

IEEE CEDA DATC RDF and METRICS2.1: Toward a Standard Platform for ML-Enabled EDA and IC Design

ASP-DAC 2022 Tutorial-1

January 17, 2022

Jinwook Jung, Andrew B. Kahng, Seungwon Kim,
and Ravi Varadarajan

Introduction

- Welcome to Tutorial 1 at ASP-DAC 2022 !
- Tutorial material: <https://github.com/ieee-ceda-datc/aspdac-2022-tutorial>
- Our speakers



Andrew B. Kahng



Seungwon Kim



Ravi Varadarajan



Jinwook Jung

Andrew B. Kahng, UC San Diego

Andrew B. Kahng is Distinguished Professor of CSE and ECE and holder of the endowed chair in high-performance computing at UC San Diego. He was visiting scientist at Cadence (1995-97) and founder/CTO at Blaze DFM (2004-06). He is coauthor of 3 books and over 500 journal and conference papers, holds 35 issued U.S. patents, and is a fellow of ACM and IEEE. He has served as general chair of DAC, ISPD and other conferences, and from 2000-2016 as international chair/co-chair of the ITRS Design and System Drivers working groups. He is currently PI of “OpenROAD” <https://theopenroadproject.org/>, a \$15M U.S. DARPA project targeting open-source, autonomous (“no human in the loop”) tools for IC implementation. He also serves as PI and director of TILOS, The Institute for Learning-enabled Optimization at Scale, which is an NSF AI Research Institute that began operations in Nov. 2021 <https://tilos.ai/>.

Some recent talks can be seen at his lab homepage:

<https://vlsicad.ucsd.edu/>



Seungwon Kim, UC San Diego

Seungwon Kim received the Ph.D. degree in electrical engineering from the Ulsan National University of Science and Technology in 2019. He is currently a Postdoctoral Researcher in the VLSI CAD lab, at the University of California San Diego. His current research interests include physical design flow optimization with machine learning. He is currently participating “OpenROAD” <https://theopenroadproject.org/>, a U.S. DARPA project targeting open-source, autonomous (“no human in the loop”) tools for IC implementation.



Ravi Varadarajan, UC San Diego

Ravi Varadarajan is currently a Ph.D. student in the ECE department at the VLSI CAD Lab at UC San Diego. He has worked in the EDA industry for over 30 years at AT&T Bell Labs, Cadence, Tera Systems Atrenta and Synopsys. His areas of interests have been in Synthesis, Place and Route, RTL and Physical Prototyping and 3D-IC design. He holds 13 US Patents. He is also a contributor to “OpenROAD”
<https://theopenroadproject.org>.



Jinwook Jung, IBM Research

Jinwook Jung is a Research Staff Member at IBM TJ Watson Research Center, Yorktown Heights, NY. At IBM, he works to advance design methodologies for AI hardware accelerators and high-performance microprocessors, leveraging machine learning and cloud computing.

He received Ph.D. from KAIST, Korea.



Tutorial Overview

- This tutorial is about IEEE CEDA DATC's efforts toward standard platform for ML-enabled EDA and IC design

• Introduction	-----	Part 1
• Challenges in ML-enabled EDA and IC design	-----	Part 2
• Toward standard platform for ML-enabled EDA and IC design	-----	
• Platform for EDA and IC Design	-----	Part 3
• DATC Robust Design Flow	-----	
• Platform for “ML-Enabled” EDA and IC Design	-----	Part 4
• Large-scale data generation in cloud-native way	-----	
• METRICS 2.1—toward standard metrics collection system	-----	Part 5
• Autonomous flow tuning framework	-----	Part 6
• Summary and future work	-----	Part 7

Agenda

- **Part 1:** Introduction
[09:00–09:10]
- **Part 2:** Challenges in ML-Enabled EDA and IC Design
[09:10–09:35]
- **Part 3:** IEEE CEDA DATC Robust Design Flow
[09:35–10:00]
- **Part 4:** Enabling Large-Scale Design Experiments with Cloud
[10:00–10:40]
- Break
[10:40–10:45]
- **Part 5:** METRICS 2.1: Toward Standard Metrics Collection System
[10:45–11:10]
- **Part 6:** AutoTuner: Autonomous RTL-to-GDS Flow Tuning
[11:10–11:40]
- **Part 7:** Summary and Future Work
[11:40–12:00]



Part 2

Challenges in ML-Enabled EDA and IC Design

Andrew B. Kahng

Challenges in ML-Enabled EDA and IC Design

- Why we care
- Generic challenges
- Infrastructure challenges
- “Measure to Improve”: METRICS2.1
- True no-human-in-the-loop: AutoTuner

Learning, Optimization, Scaling

- “Machine **Learning** (ML) is the part of AI studying how computer agents can improve their perception, knowledge, thinking, or actions **based on experience or data.**”

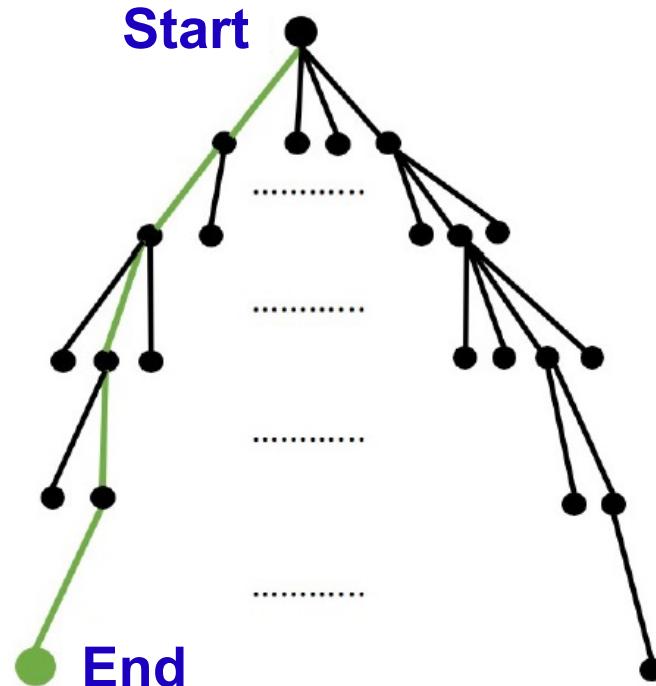
Prof. Christopher Manning, Stanford, Sept. 2020

<https://hai.stanford.edu/sites/default/files/2020-09/AI-Definitions-HAI.pdf>

- Optimization** is the universal quest to *do better*.
- Scaling** is what drives all of us.



Challenge: Optimization (IC Design) “Lives in a Box”

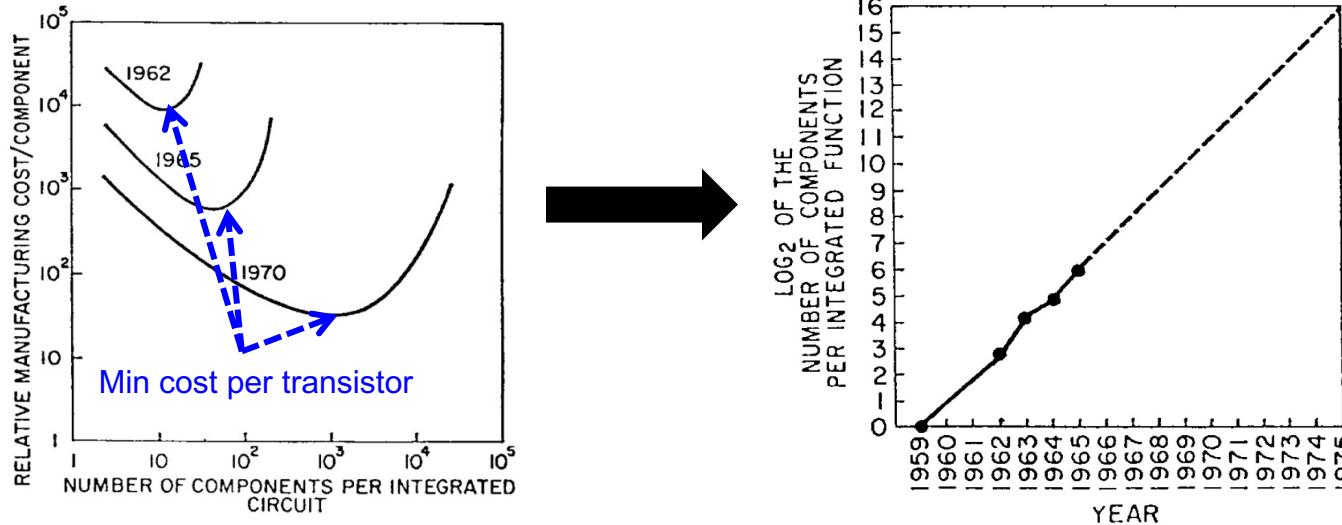


Huge space of trajectories: architecture, enablement, IPs, tools, manual fix, ...

- Start to End: expensive!
 - $O(\text{year})$ for product
 - $O(\text{weeks})$ for SP&R and Opt
- Goal: best possible End
- Constraint: stay in “Box”
 - {compute}
 - X {licenses}
 - X {people}
 - X {weeks}

Scaling: Delivers Value

- **Moore, 1965:** “The complexity for minimum component costs has increased at a rate of roughly a factor of two per year”



- **Scaling focus:** “PPAC” power, performance, area, cost
- **Moore’s Law is a law of cost reduction** 1% = 1 week
 - **Corollary:** greater reach of integration, more innovation within reach

ML for EDA and IC Design: What

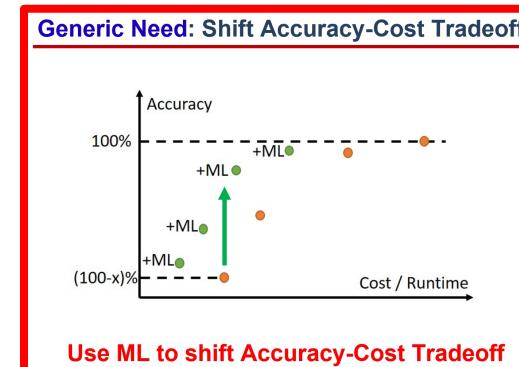
- **Predict**
 - Will RouteOpt finish with clean signoff, <1000 DRVs by tomorrow night?
- **Classify**
 - Out of these 50 floorplans + budgets, which 3 should go into trial SP&R?
- **Estimate**
 - How many hold buffers will tool eventually add into this post-CTS layout?
- **Guide / advise**
 - What P&R tool setup/script will obtain the best QOR within next 36 hours?
- More broadly: answer any question that is difficult for humans
 - Google Brain, 2020: “super-human macro placement” on arXiv
 - **Overarching: “intelligent flow”, “automated super-human expertise”**
- More directly: regressions and image classifications (LSF, litho)

ML for EDA and IC Design: Why

- A. You need models to have predictions
- B. You need predictions to leverage in exploration
- C. What you can't predict, you guardband
- D. What you don't explore, you leave on the table
- E. C and D are bad for product quality and schedule

- We are in an Era of Optimization
 - Look for ML to win quality, schedule, cost
 - E.g., reduce analysis runtime, miscorrelation

→ We hope that ML will bring Scaling



4 Aspects of ML for EDA and IC Design

1. Mechanization and Automation

Create super-human robot engineers

2. Orchestration of Search and Optimization

Optimize the use of N robot engineers

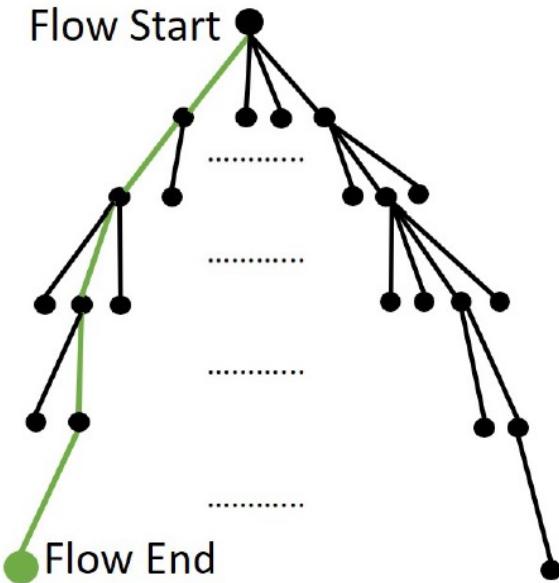
3. Pruning via Predictors, Models

Predict design-specific tool outcomes

Prune “doomed runs”

4. From Reinforcement Learning to Intelligence

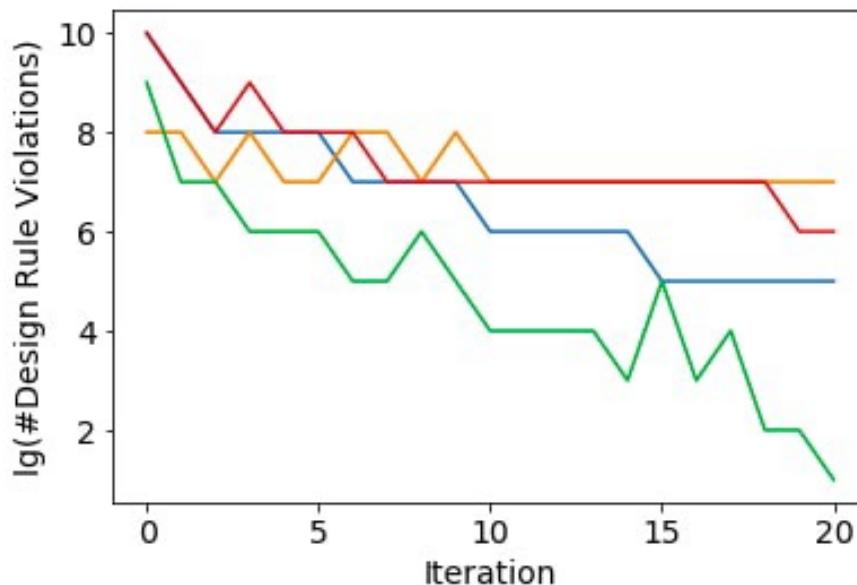
Target: “MLDA”, “self-driving tools and flows”, “superhuman”



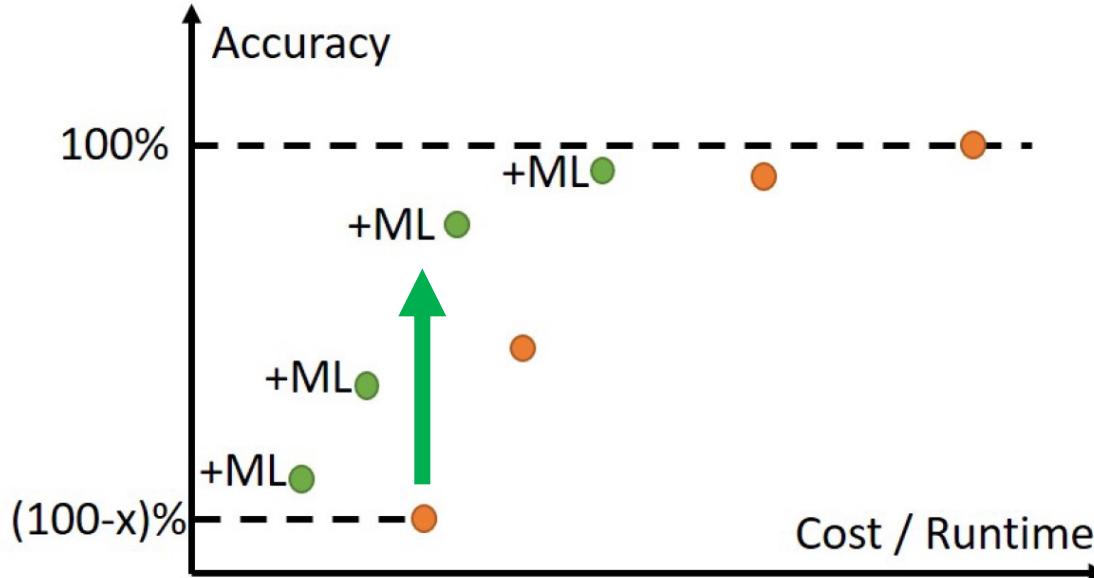
Huge space of tool, command, option trajectories through design flow

Generic Need: Predict Doomed Runs

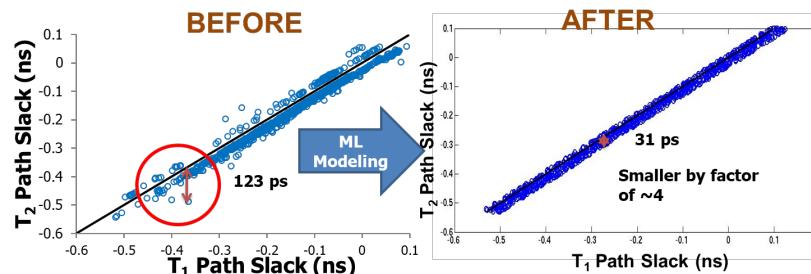
- Example: progression of DRC violations in commercial router
- Simple strategy: **track and project key metrics as time series**
- Example method: use Markov decision process (MDP): “GO” vs. “STOP” strategy card to terminate “doomed runs” early



Generic Need: Shift Accuracy-Cost Tradeoff

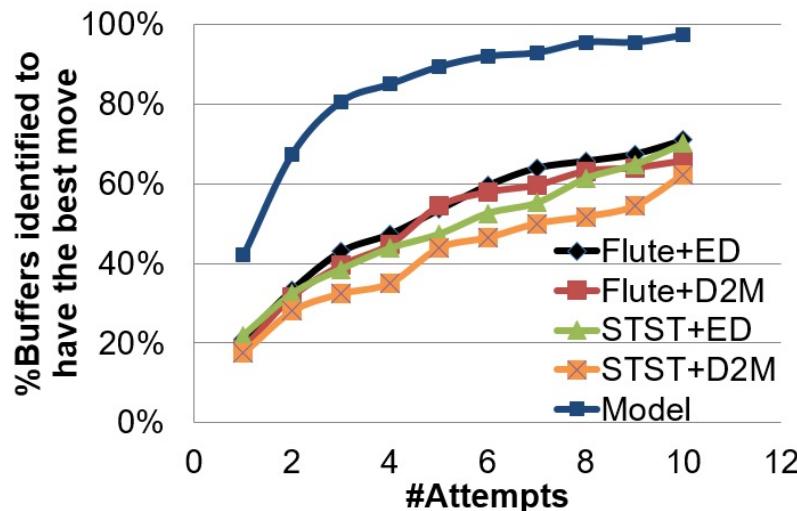
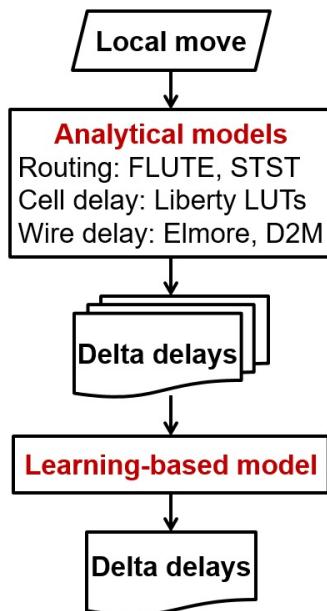
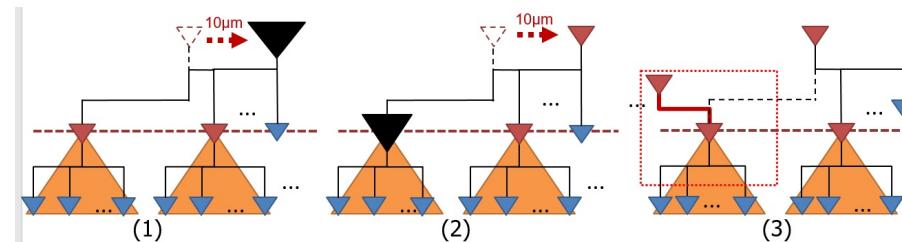


Use ML to shift Accuracy-Cost Tradeoff



Generic Need: Model-Guided Optimization

- Which CTS tweak will improve skew variation across corners?



Infrastructure Challenges and Opportunities

- **Big data and learning for IC designers**

- Model encapsulation and application
- IP preservation to enable model sharing

- **Big data and learning for EDA tools/flows**

- Data model, names, semantics **including derived data**
- Enable design-specific **modeling and prediction** of tool/flow outcomes

Recall: <http://vlsicad.ucsd.edu/GSRC/metrics>

- **Machine learning requires data**

- **Real designs, Artificial designs and “Eyecharts”**
- **Calibration data and more** (“Underwriters Lab”)
- **Develop training data together** – grid computing paradigm
 - **Year 2000:** SETI@Home **Year 2020:** Tool X on PDK Y with IP Z ?
 - + Challenges and incentives: “Kaggle for ML in IC design”
 - IEEE CEDA DA Technical Committee: “Metrics4ML” <https://github.com/ieee-ceda-datc/datc-rdf-Metrics4ML>

Infrastructure Principles

- General and extensible
 - Syntax and semantics to support future-proofing – e.g., integration and physics
 - 3D/2.5D, reliability, thermal/stress, ...
- Open, transparent and unencumbered
 - Many landmines: copyrighted command names, reporting formats, “de facto standards”
→ staying aware and clear of these helps solve the Tower of Babel problem!
 - **Permissive** open sourcing is best: BSD/MIT/Apache licenses
- No ambiguity!
 - METRICS2.1 discussed later today: each measurement concept maps to a **unique** METRICS2.1 metric name
 - Each METRICS2.1 metric name maps to a **unique** interpretation as a measurement
 - **1-1 correspondence is crucial to avoid future confusion**
- **Free, open and frictionless to everyone; agnostic to EDA vendor**

ML in IC Design Requires a Standard API!

- **Support for ML in IC design**
 - **Standards** for model encapsulation, model application, and IP preservation when models are shared
- **Standard ML platform for EDA modeling**
 - Enable design metrics collection, tool/flow model generation, design-adaptive tool/flow configuration, prediction of tool/flow outcomes
- **Datasets to support ML**
 - Real designs and artificial designs
 - Shared training data – e.g., analysis correlation, post-route DRV prediction, optimal sizing
- Challenges and incentives: “Kaggle for ML in IC design”

The screenshot shows the METRICS website with a navigation menu on the left and a main content area on the right.

Navigation Menu:

- Main
- Overview
- Infrastructure
 - Metrics List
 - Dictionary
- Publications
 - Papers
 - Presentations
 - GSRC Presentations
- Codes
- Theses
- Links
 - Conferences & Workshops
 - Softwares
 - Others

Main Content Area:

The METRICS Initiative

Recent Updates

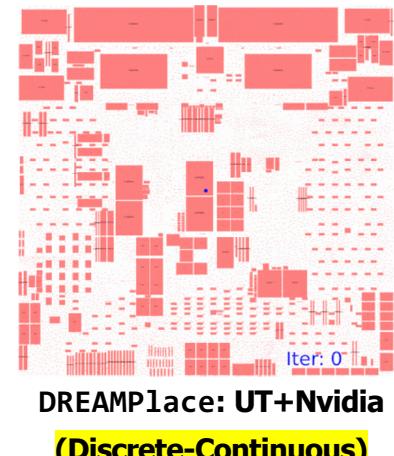
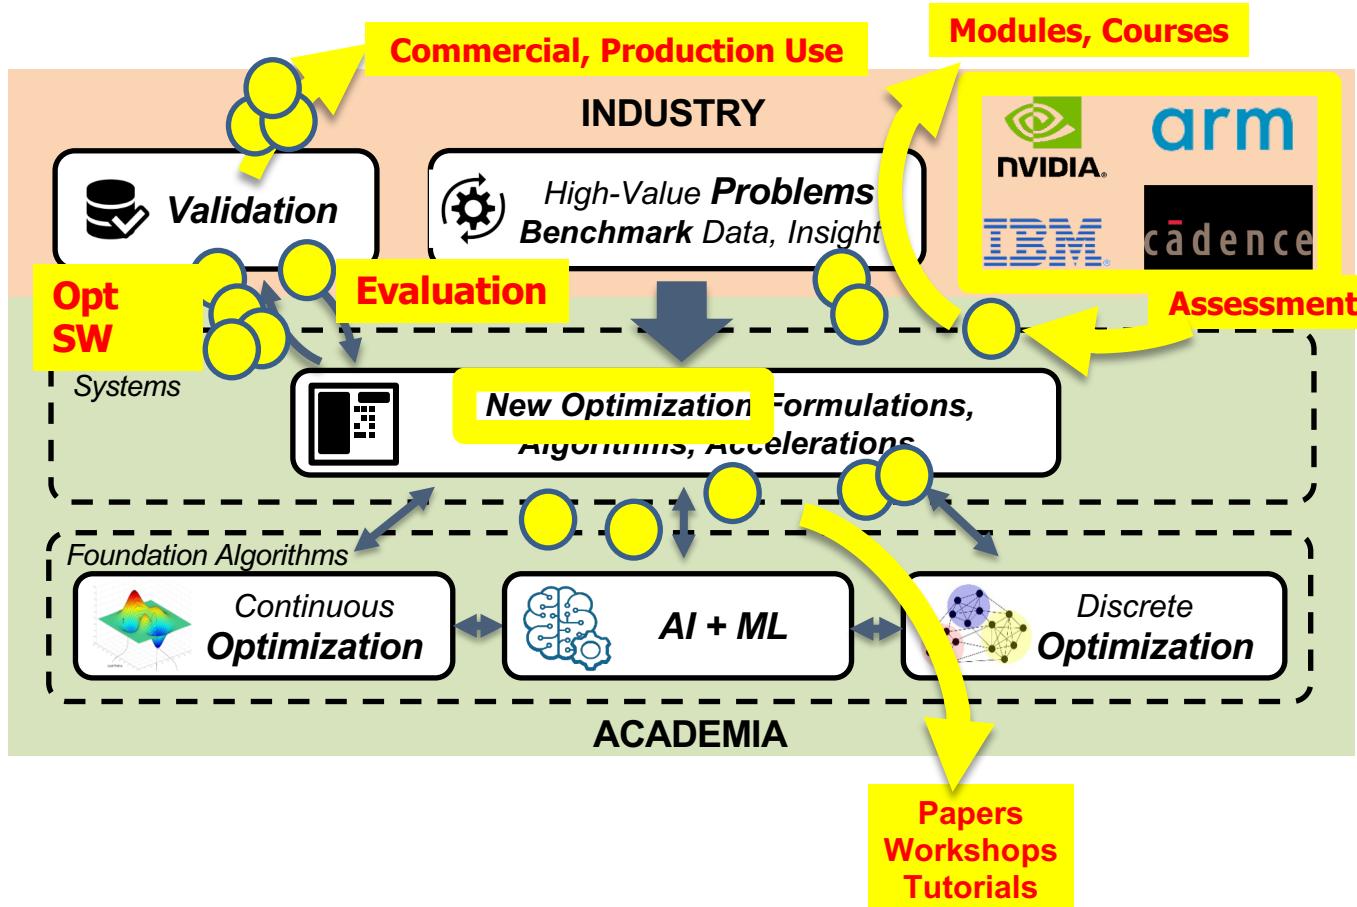
- Survey (1st draft) for design quality and productivity (the multiple-choice version)
- Reduced survey (2nd draft) for design quality and productivity that is distributed at June 2001 GSRC workshop
- Updated survey (3rd draft) for design quality and productivity that reflects the discussion at June 2001 GSRC workshop
- Workshop notes for METRICS discussion at June 2001 GSRC workshop
- List of prediction/estimator models enabled by METRICS System ***
- DAC02 Birds-of-a-Feather meeting summary (June 12, 2002) ***

- **METRICS 1.0** (1999; DAC00, ISQED01)
 - “Measure to Improve” <http://vlsicad.ucsd.edu/GSRC/metrics>
- **METRICS 2.0** ([WOSET-2018](#)) was proposed as an update of **METRICS 1.0**
- **METRICS2.1** is proposed as an extension to **METRICS 2.0**

Flow Hyperparameter AutoTuner

- A “No-Human-in-Loop” for RTL-to-GDS flow using the **OpenROAD** tool
- Automatic iterative tuning for **QoR improvement** within a given hyperparameter range space
- Interface: Python packages Ray/Tune
- Search algorithms: HyperOpt, PBT, Optuna, Nevergrad, Ax, random search
- Advantages:
 - Pre-acquired big data is not necessary for execution.
 - Less experimental trials compared to random or grid parameter sweeping
 - Powerful parallelization management → core/thread management, external server usage, web visualization
 - Suboptimality gap reduces significantly – and according to available CPU resource

Lifecycle of Learning + Optimization + Impact in Practice



THANK YOU !

- Research at UCSD is supported by NSF, DARPA, Qualcomm, Samsung, NXP, Mentor Graphics and the C-DEN center.
- Questions/Feedback: abk@eng.ucsd.edu

Part 3

IEEE CEDA DATC Robust Design Flow

Jinwook Jung

Part 3: Outline

- Introduction
 - IEEE CEDA DATC
 - DATC Robust Design Flow (RDF)
- The DATC RDF flow
 - Chisel and hardware generator designs
 - RTL obfuscation with ASSURE
 - Logic synthesis with Yosys+ABC
 - DFT insertion with Fault
 - P&R with OpenROAD or point tool-based flow
- Using DATC RDF: A short demo

Design Automation Technical Committee (DATC)

- A technical committee of IEEE Council on EDA (CEDA)
 - Provide a forum for discussing strategies and issues in design automation
- Tasks:
 1. Academic reference design flow, *DATC Robust Design (RDF)* → Part 3
 2. Standard platform for ML-enabled EDA → Parts 4-6

Committee	Affiliation	Role
Jinwook Jung	IBM Research, USA	Chair / Integration
Andrew B. Kahng	UC San Diego, USA	Co-chair / Flow
Victor N. Kravets	IBM Research, USA	Logic synthesis
Jianli Chen	Fudan Univ., China	Placement
Yih-Lang Li	National Yang Ming Chiao Tung Univ., Taiwan	Routing
Iris Hui-Ru Jiang	National Taiwan Univ., Taiwan	Previous chair

DATC Robust Design Flow (RDF)

- **Academic reference design flow**

- Initiated 2016; latest release 2021
- Design flow built upon academic tools
- Also supports complete, industrial format-based RTL-to-GDS flow with OpenROAD

- **GitHub link:**

- <https://github.com/ieee-ceda-datc/datc-rdf>

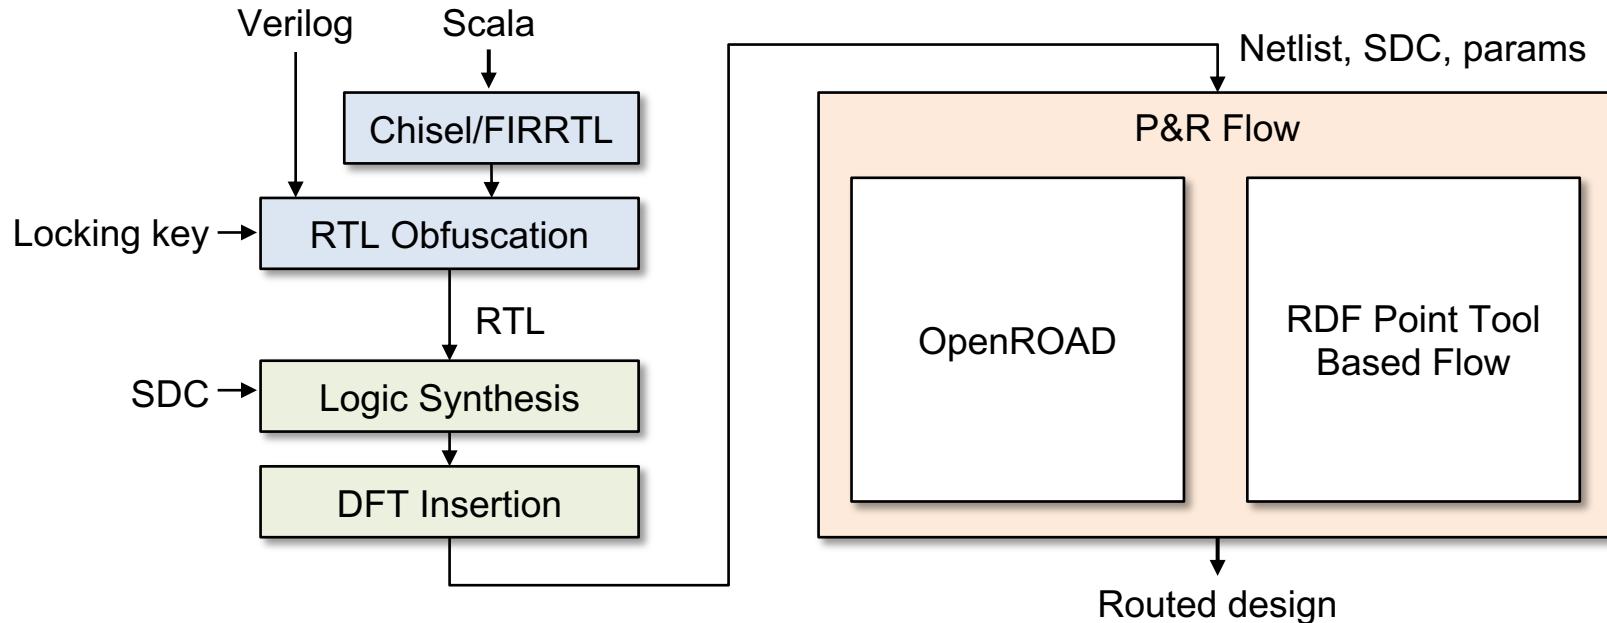
- **Goal:**

1. Preserve leading research codes, and integrate them into design flow
2. Trigger design flow-scale and cross-stage optimization research

Component	Tool
RTL generator	Chisel/FIRRTL
RTL obfuscation	ASSURE
Logic synthesis	Yosys+ABC
DFT insertion	Fault
Floorplanning	TritonFP
Global placement	RePIAce, FZUplace, NTUPlace3, ComPLx, Eh?Placer, FastPlace3-GP
Detailed placement	OpenDP, MCHL, FastPlace3-DP
Flip-flop clustering	Mean-shift, FlopTray
Clock tree synthesis	TritonCTS
Global routing	FastRoute4-lefdef, NCTUgr, CUGR
Detailed routing	TritonRoute, NCTUdr, DrCU
Layout finishing	KLayout, Magic
Gate sizing	Resizer, TritonSizer
Parasitic extraction	OpenRCX
STA	OpenSTA, iTimerC
Integrated P&R app	OpenROAD
Database	OpenDB
Libraries/PDK	NanGate45, SKY130, ASAP7, NCTUcell
Benchmark conversion	RosettaStone

DATC RDF: Flow Overview

- Starts from **pure Verilog** or **Chisel-generated Verilog** (optionally obfuscated)
- **Logic synthesis** followed by **DFT insertion**
- P&R with **OpenROAD** or **point tool-based RDF flow**



- **Chisel:** Constructing **H**ardware **I**n **S**cala **E**mbedded **L**anguage
 - Provides special classes, objects, language convention for hardware design
 - Easy to implement highly-parameterized, abstract **hardware generators**
- **Substantial interest from hardware architecture researchers**
 - RISC-V Rocket Chip Generator
 - Gemmini deep learning accelerator
 - Hwacha vector processor
- With Chisel, we can **use recent processor/accelerator designs in design flow research**

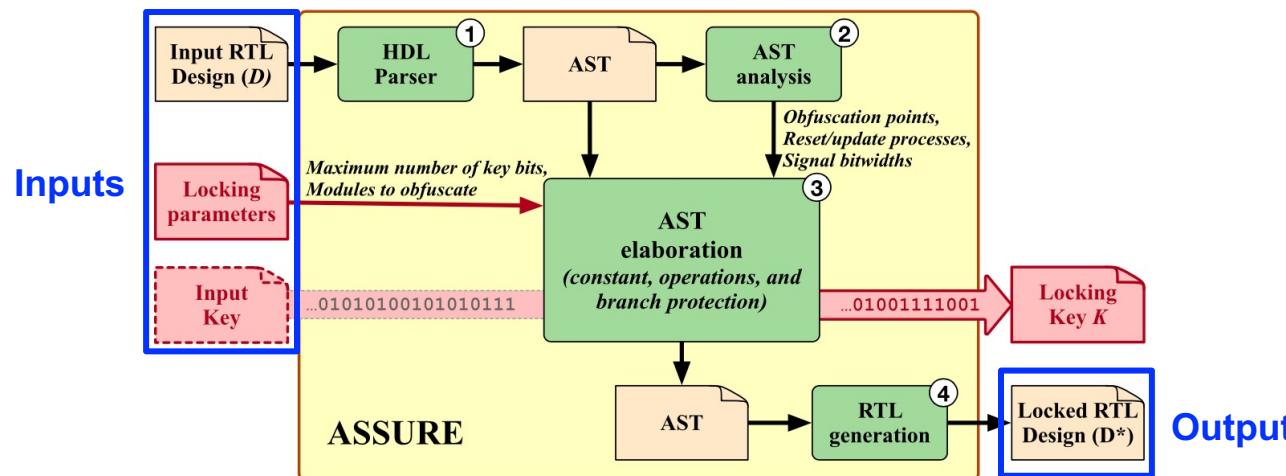
```
import chisel3._

class Filter (bitWidth: Int) extends Module {
    val io = IO(new Bundle {
        val in = Input(UInt(bitWidth.W))
        val out = Output(UInt(bitWidth.W))
    })
    val z1 = RegNext(io.in)
    val z2 = RegNext(z1)
    io.out := (io.in*1.U) + (z1*1.U) + (z2*1.U)
}

// Parameterized instantiation
val movingSum3Filter = Module(new Filter(8))
```

RTL Obfuscation with ASSURE

- **ASSURE:** RTL obfuscation framework assuming netlist-only threat model
- **3 obfuscation techniques supported:**
 - 1. Constant:** substitute constants in RTL with secret locking key
 - 2. Operation:** adds redundant operators alongside original operators in RTL
 - 3. Branch:** obfuscates branch conditions in RTL



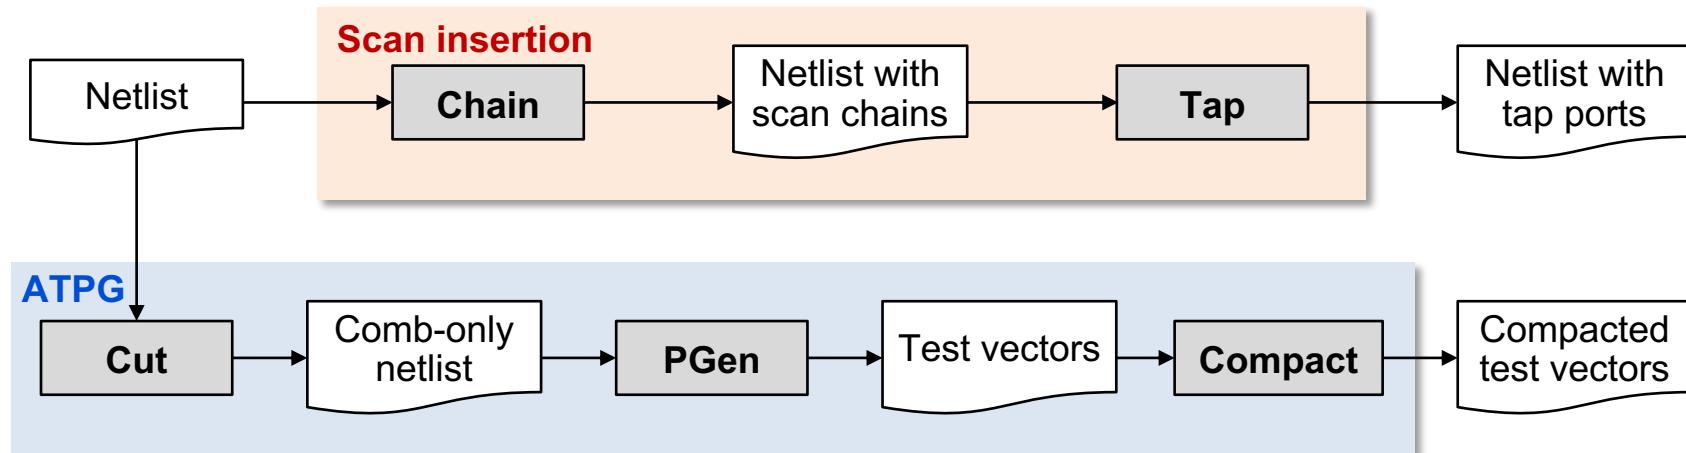
Logic Synthesis using Yosys+ABC

- Logic synthesis in RDF is done with Yosys+ABC
 - GitHub link: <https://github.com/The-OpenROAD-Project/yosys>
- Latest ABC updates incorporated in DATC RDF
 - Tight integration of SAT solver with internal design representation
 - Improved scalability and new options to improve technology mapping quality
- LazyMan synthesis scripts:
 - Reuse logic rewriting choices empirically discovered in earlier runs
 - Reduce runtime while improving quality

```
abc 01> &cec -x nn.170.aig ← New option “-x”
cnn.187.aig Networks are equivalent. Time = 1238.07 sec
abc 02> &cec cnn.170.aig cnn.187.aig
Networks are equivalent. Time = 12058.28 sec
```

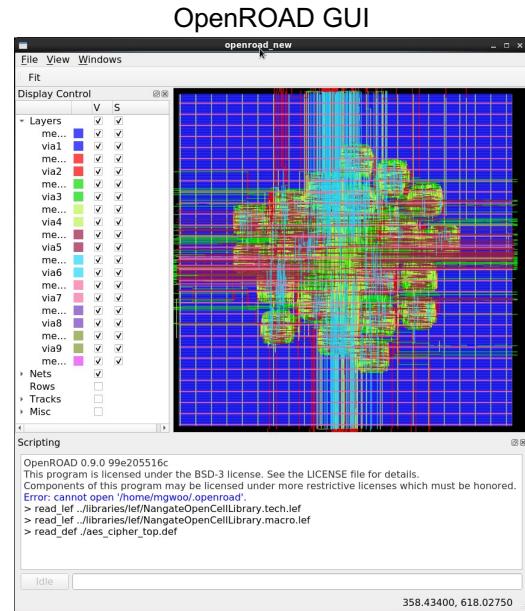
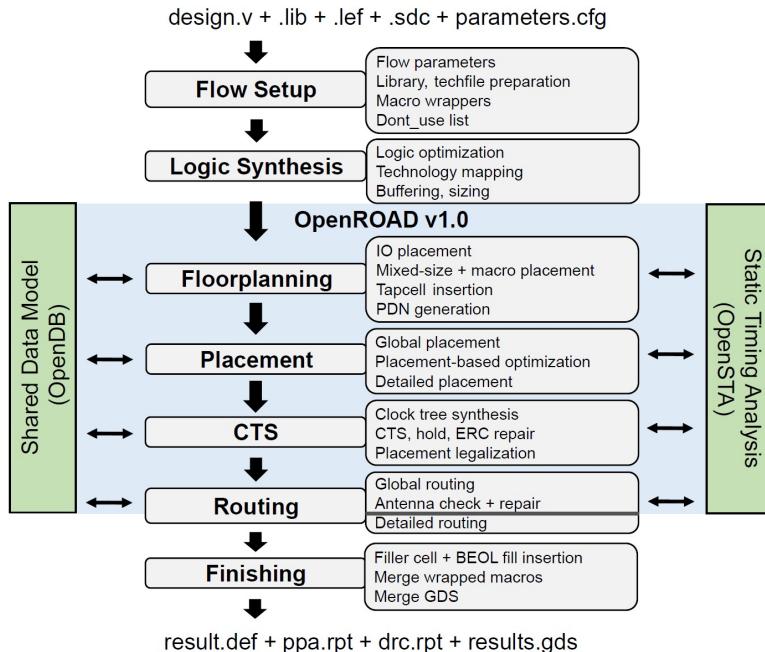
DFT Support with Fault

- **Fault:** open-source DFT infrastructure
 - GitHub link: <https://github.com/Cloud-V/Fault>
 - **Scan chain insertion:** add test signals, scan chains, and JTAG interface
 - **ATPG:** generate compressed test vectors



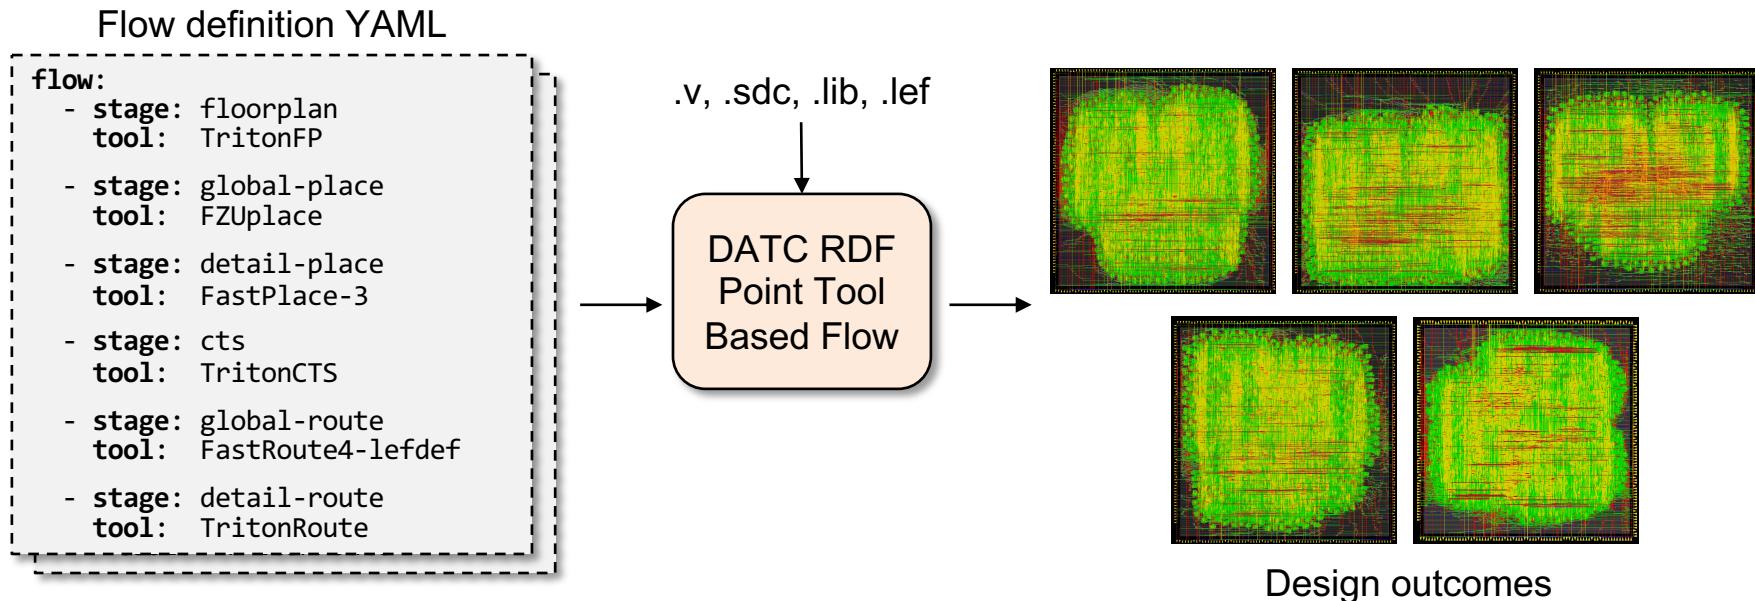
P&R Flow: OpenROAD

- OpenROAD: famous open-source tool for complete RTL-to-GDSII flow
 - GitHub link: <https://github.com/The-OpenROAD-Project>
 - Enable robust handling of industry standard design enablements in RDF



P&R Flow: Point-Tool Based Flow

- Composed with academic point tools, especially contest-winning tools
 - GitHub link: <https://github.com/ieee-ceda-datc/RDF-2020>
 - Preserve leading research codes and integrate them into complete design flow



Using DATC RDF

- **DATC RDF consists of flow collateral and tool binaries**
 - Flow collateral in GitHub: <https://github.com/ieee-ceda-datc/datc-rdf>
 - Tool binaries provided in Docker image: <https://hub.docker.com/r/jinwookjung/rdf>
- **Installing and launching RDF:**

```
# Clone DATC RDF repository
git clone https://github.com/ieee-ceda-datc/datc-rdf.git
cd datc-rdf
git submodule update --init

# Launch DATC RDF
docker run --rm -it \
-v $(pwd)/submodules/OpenROAD-flow-scripts/flow:/openroad-flow \
-v $(pwd)/example:/root/example \
jinwookjung/rdf
```

Demonstration of DATC RDF

1. RTL generation with Chisel
2. RTL obfuscation with ASSURE
3. Logic synthesis with Yosys
4. DFT insertion with Fault
5. P&R with OpenROAD

Summary

- **DATC RDF:** Academic reference **RTL-to-GDS** design flow
 1. Preserve/integrate leading research codes in complete design flow
 2. Trigger design flow and cross-stage optimization research
- **The DATC RDF Flow**
 - Starts from pure Verilog or Chisel-generated Verilog (optionally obfuscated)
 - Logic synthesis followed by DFT insertion
 - P&R done with OpenROAD integrated app or point tool-based RDF flow
- **GitHub link:** <https://github.com/ieee-ceda-datc/datc-rdf>

Part 4

Enabling Large-Scale Design Experiments with Cloud

Jinwook Jung

Outline

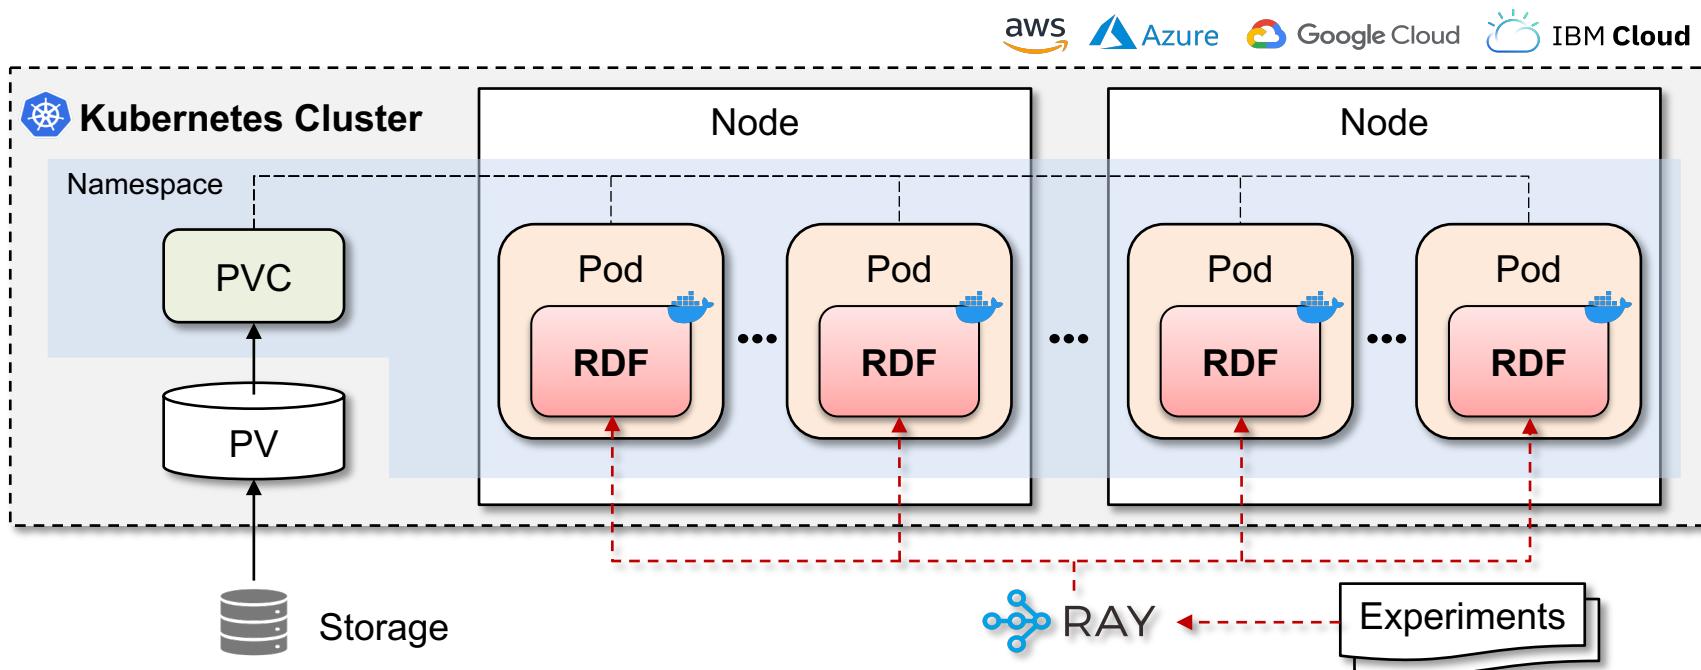
- Introduction and overview
- Docker and Kubernetes
- Kubernetes cluster architecture for DATC RDF
- Scaling-out design experiments
 - Multi-container design experiments using Ray
 - Multi-node design experiments using Ray and public cloud services
- Summary

Introduction

- “***Machine learning requires data***”
 - Generating data for ML-enabled EDA needs compute resource and time
 - Efficient way to generate large amount data for ML-enabled EDA?
- **Public cloud provides compute resources on-demand**
 - How to best-utilize public cloud for ML-CAD?
 - Obtain cloud instance, install libraries, compile tools, ... → Not scale well
 - Need to establish ***generic*** and ***cloud-native*** way for enabling large-scale design experiments on public cloud services

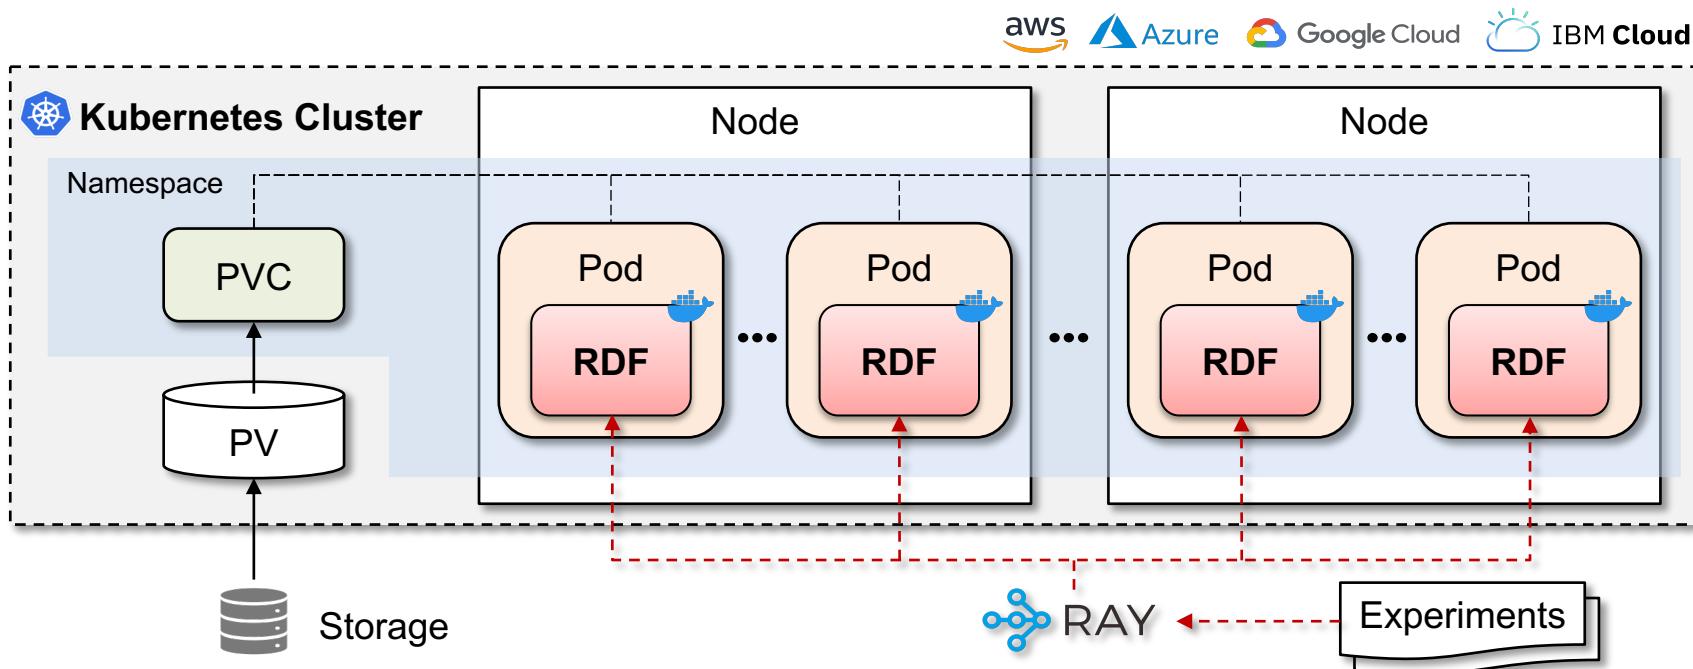
Goal of Part 4

- We will learn how to deploy large-scale design experiments on cloud, utilizing *Docker*, *Kubernetes*, *Ray*, and *public cloud services*



Overview of Part 4

1. Basics of Docker and Kubernetes
2. Minimal Kubernetes cluster for design experiments
3. Scale-out to multi-container, multi-node cluster using public cloud services

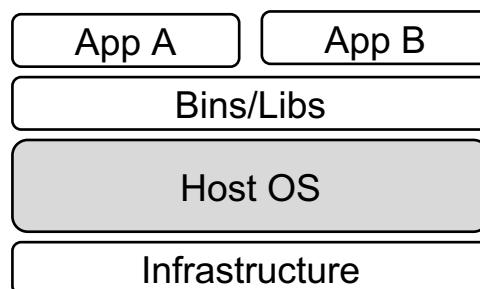


Outline

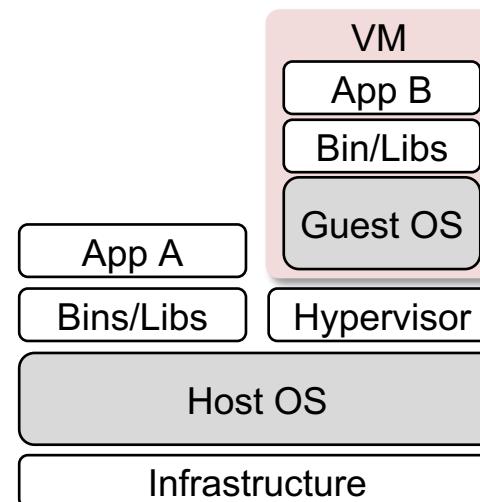
- Introduction and overview
- Docker and Kubernetes
- Kubernetes cluster architecture for DATC RDF
- Scaling-out design experiments
 - Multi-container design experiments using Ray
 - Multi-node design experiments using Ray and public cloud services
- Summary

Docker and Container

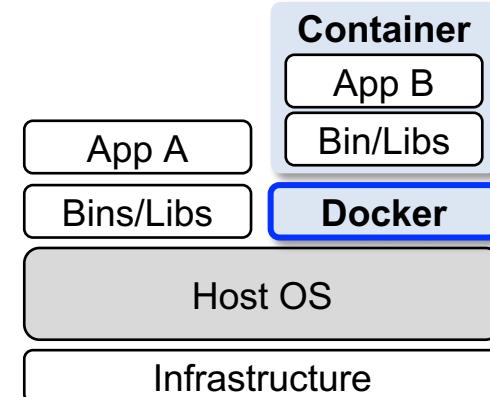
- **Docker:** creates application into *container*
 - **Container** = application + dependencies (libs/bins)
 - Isolate applications from each other, while sharing same OS kernel
- Comparison to different types of applications



Traditional app



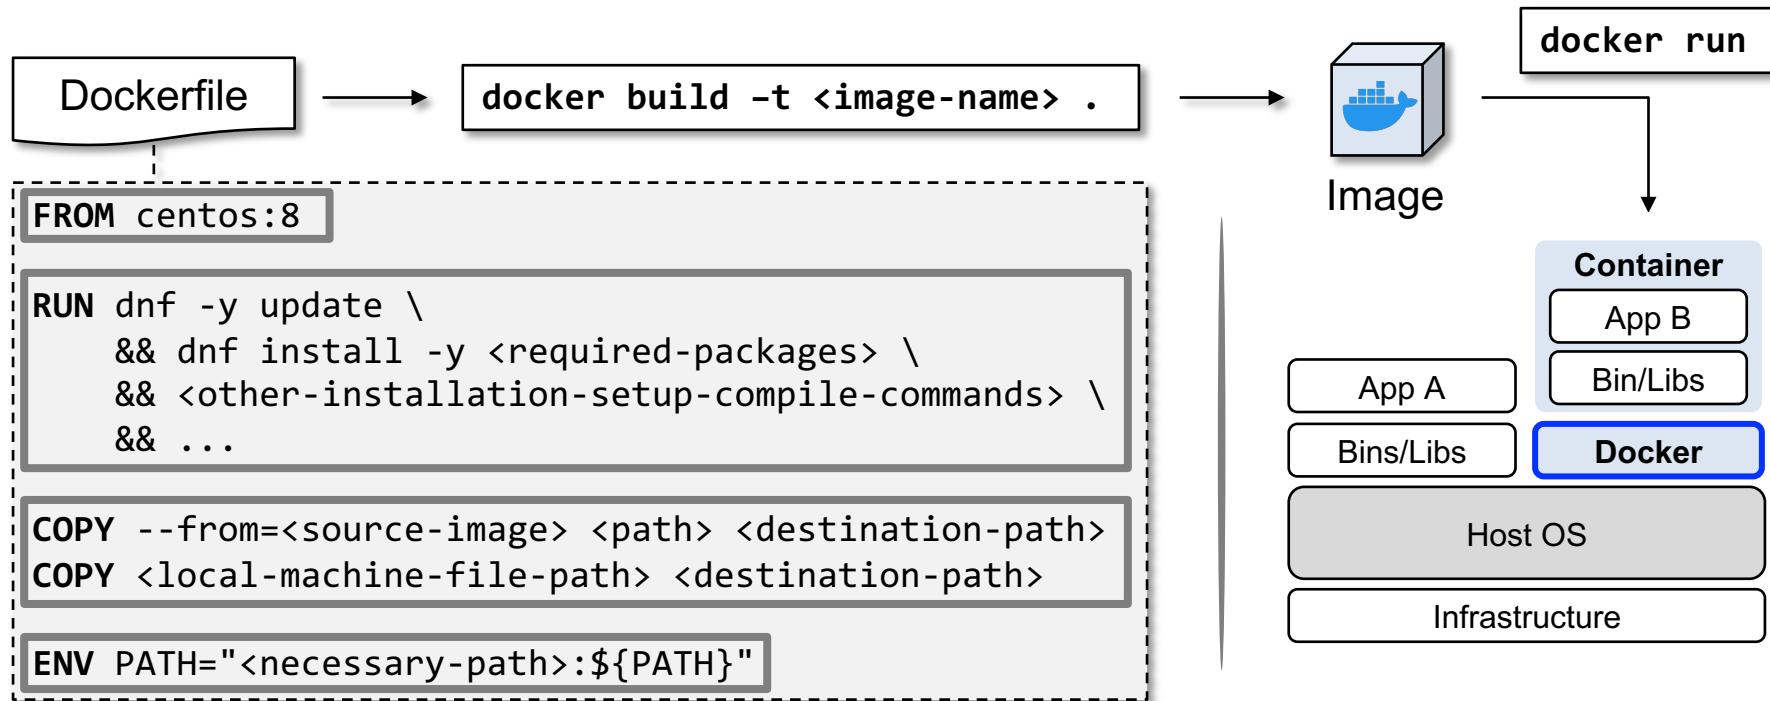
VM-based app



Containerized app

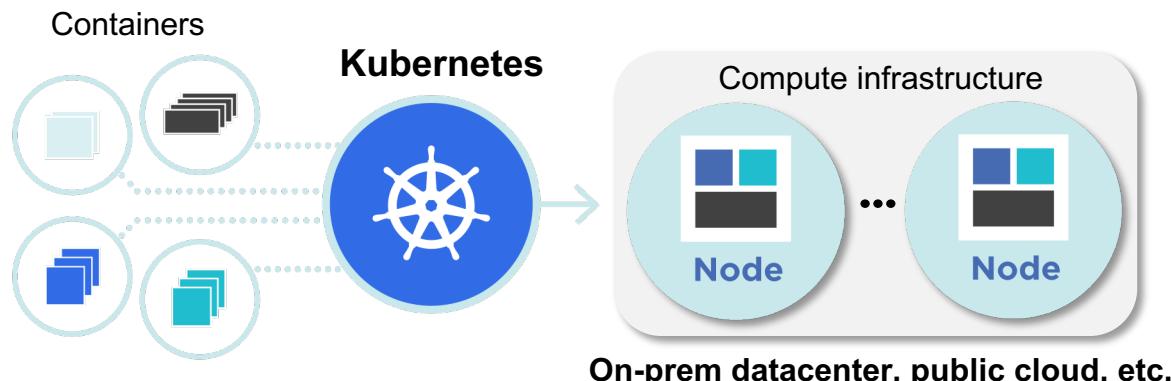
Docker Image and Runtime

- **Image:** application packaged by Docker and deployable as container
- **Docker runtime:** executes container from image



Kubernetes and Container Orchestration

- **Docker:** build and run containers
- **Kubernetes:** provide *orchestration services*
 - Deploy, scale, and manage containers on large-scale compute infrastructure
 - Abstracts underlying compute infrastructure (on-premise datacenter machines, local servers, and public cloud services)
 - Infrastructure agnostic → standard way to deploy containerized applications



Basic Kubernetes Terminologies

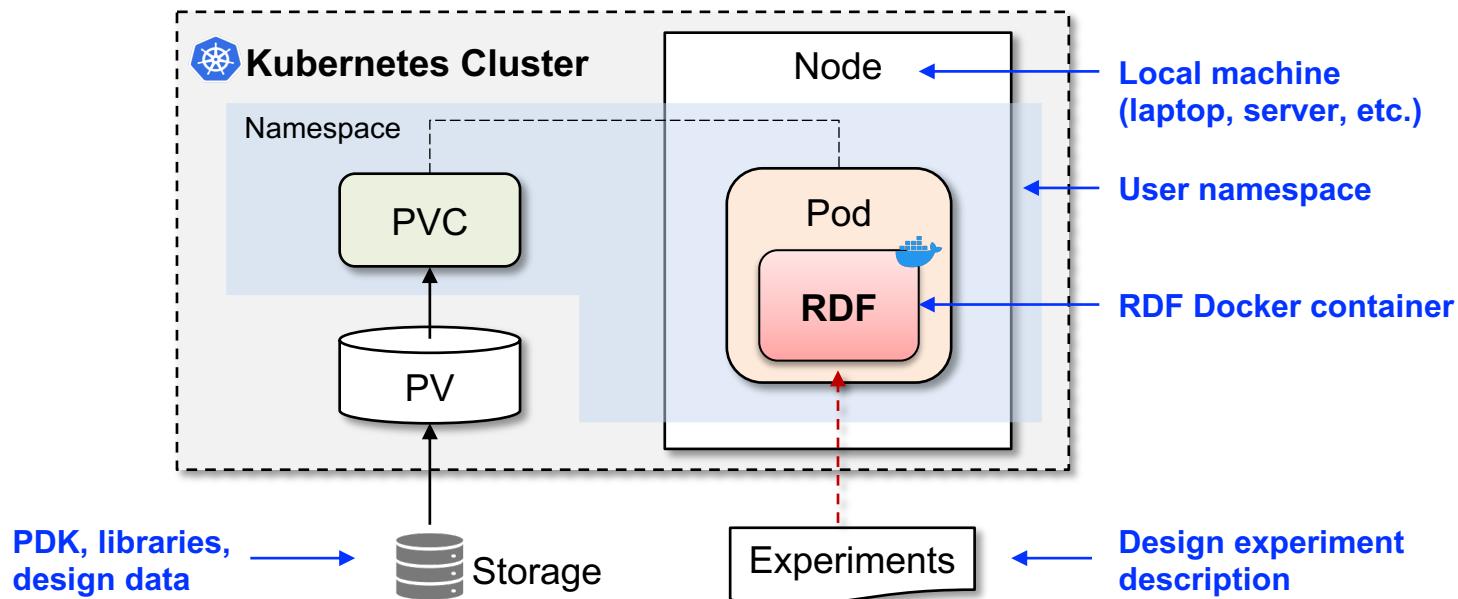
- **Node:** Compute machine, e.g., physical server, VM, cloud instance, etc.
- **Cluster:** A set of nodes
- **Pod:** Construct for running containers in Node
- **Persistent Volume (PV):** External storages made available in Cluster, such as local directory, NFS and cloud file storage
- **Persistent Volume claim (PVC):** “Ticket” authorizing Pod to mount PV

Outline

- Introduction and overview
- Docker and Kubernetes
- **Kubernetes cluster architecture for DATC RDF**
- Scaling-out design experiments
 - Multi-container design experiments using Ray
 - Multi-node design experiments using Ray and public cloud services
- Summary

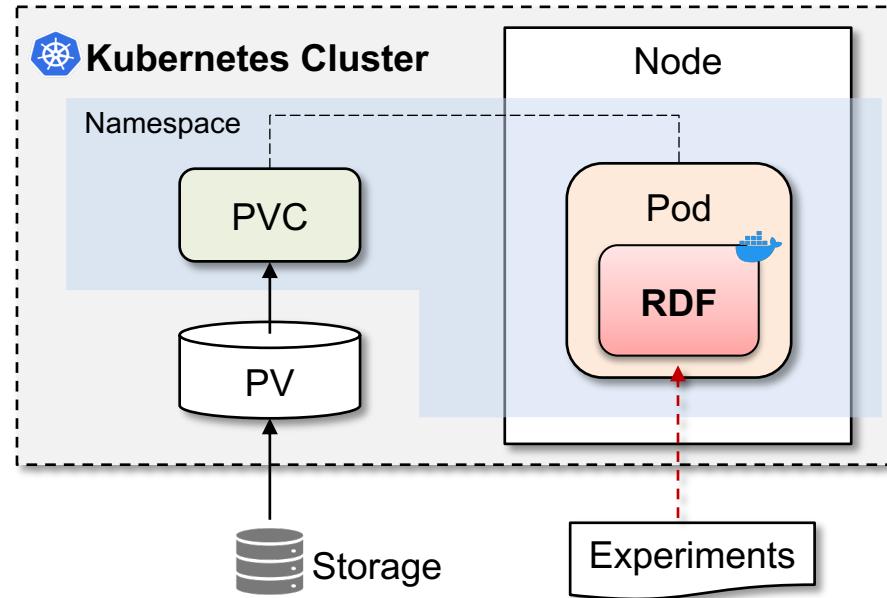
“Minimal” Kubernetes Cluster Architecture for DATC RDF

- Single node
- Single volume (PV/PVC)
- Single pod and container



Demonstration

- How to create the minimal Kubernetes cluster and run DATC RDF there

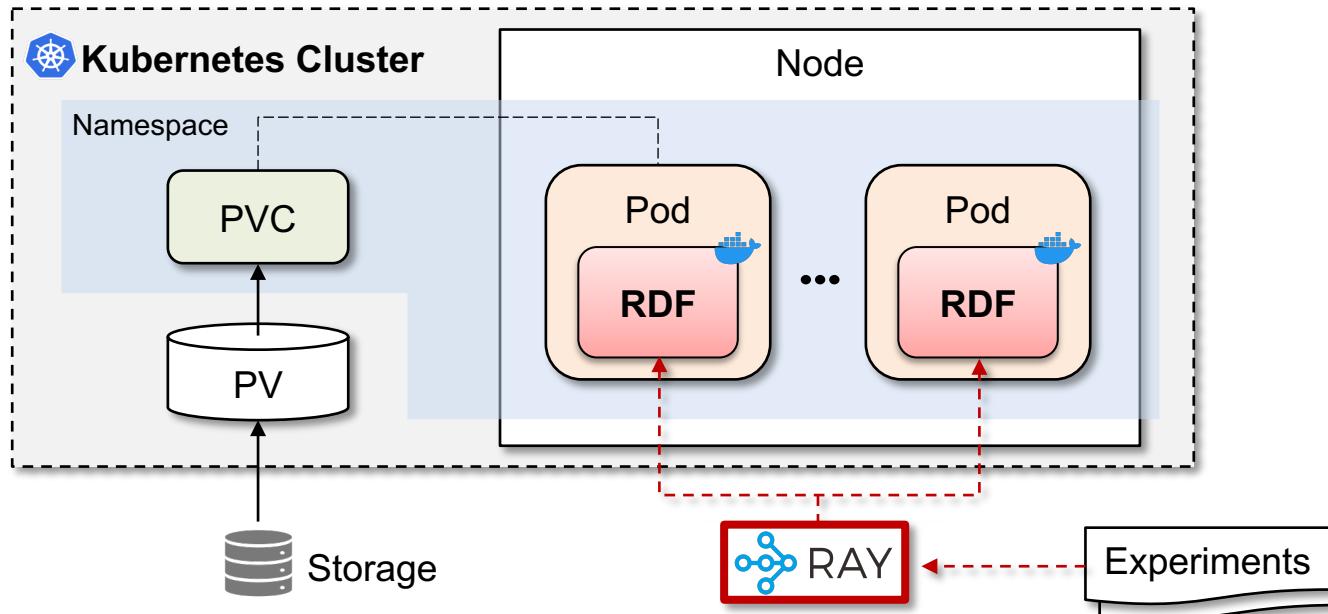


Outline

- Introduction and overview
- Docker and Kubernetes
- Kubernetes cluster architecture for DATC RDF
- Scaling-out design experiments
 - Multi-container design experiments using Ray
 - Multi-node design experiments using Ray and public cloud services
- Summary

Scaling Out DATC RDF Containers

- Single-node/single-pod cluster → limited compute resource
- For large-scale design experiments, we want larger cluster
 - One way to achieve it is to scale-out containers/pods in Kubernetes cluster
 - How to efficiently **distribute experiments?** → We introduce *Ray*



Ray

- **Distributed execution framework** developed by UC Berkeley
 - GitHub Link: <https://github.com/ray-project/ray>
 - Provides a simple Python API to build distributed applications
 - Works with Kubernetes and other cloud-based clusters

- **Example:**

Original program

```
def f(x):
    return x+1

[f(i) for i in range(100)]
```

Distributed version with Ray

```
import ray
ray.init(...)

@ray.remote
def f(x):
    return x * x

[f.remote(i) for i in range(100)]
```

→

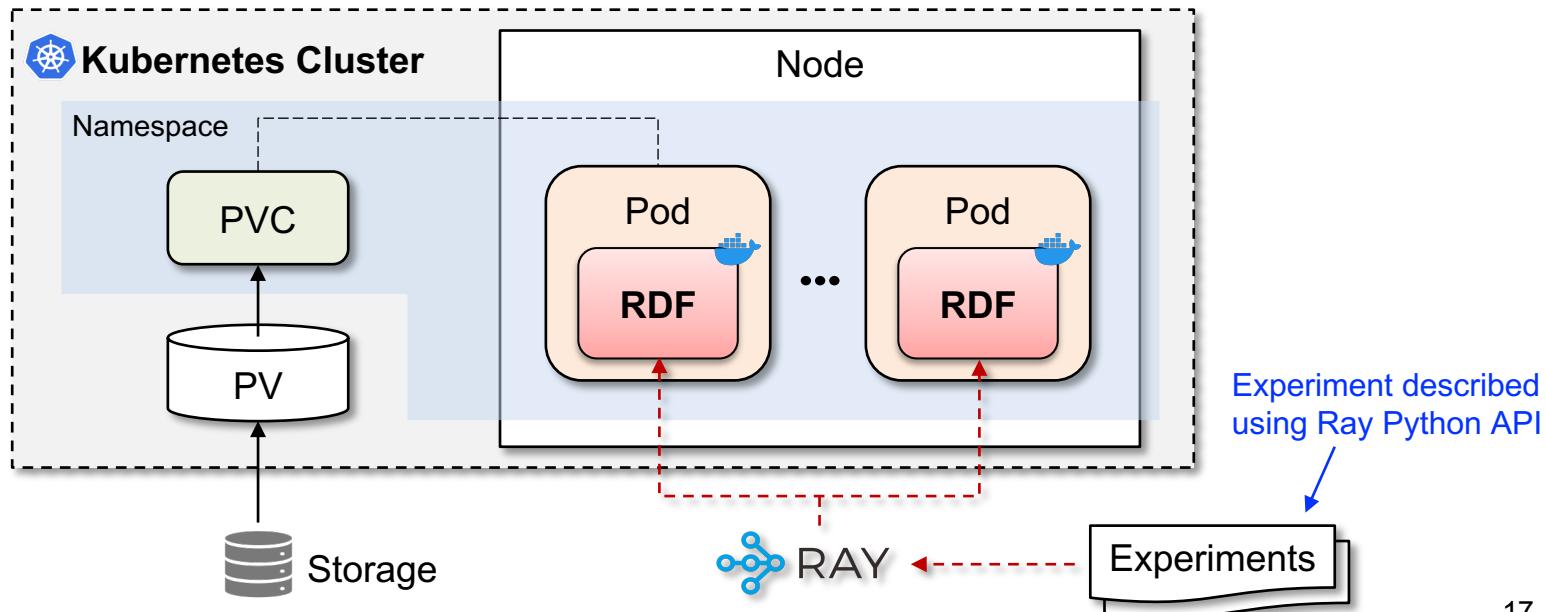
Initialize Ray

Decorate function



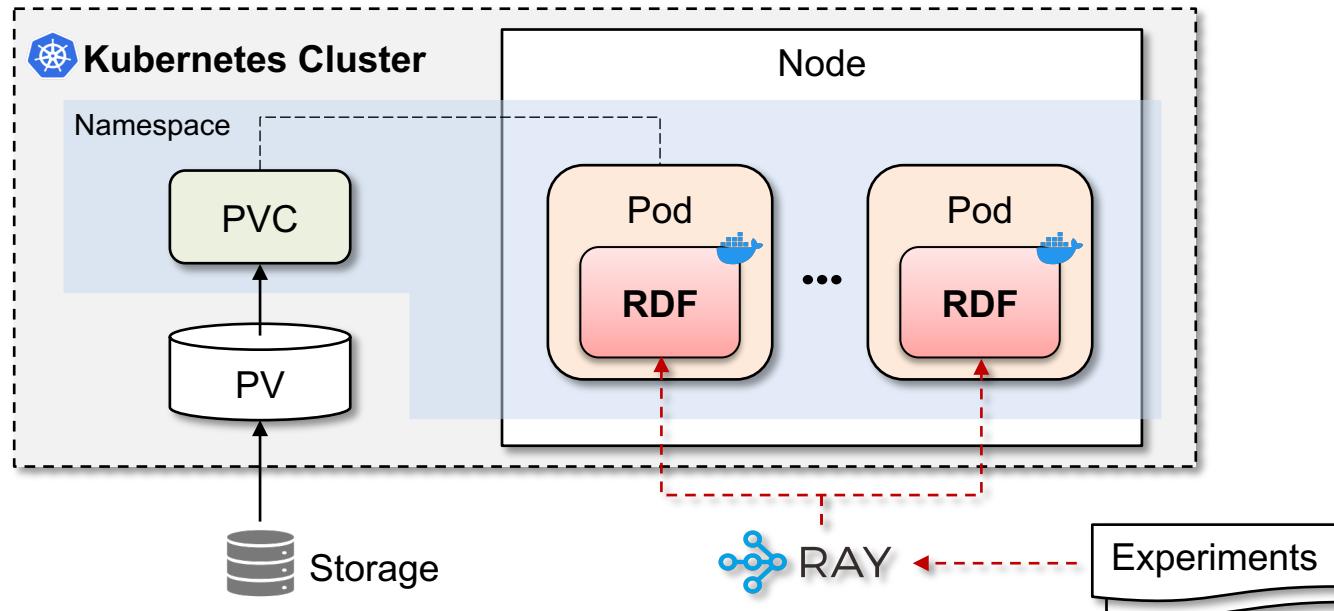
Distributed Design Experiments with Ray and Kubernetes

1. Create multi-container Kubernetes cluster
2. Describe large-scale design experiments using Ray Python API
3. Deploy experiments on Kubernetes cluster via Ray



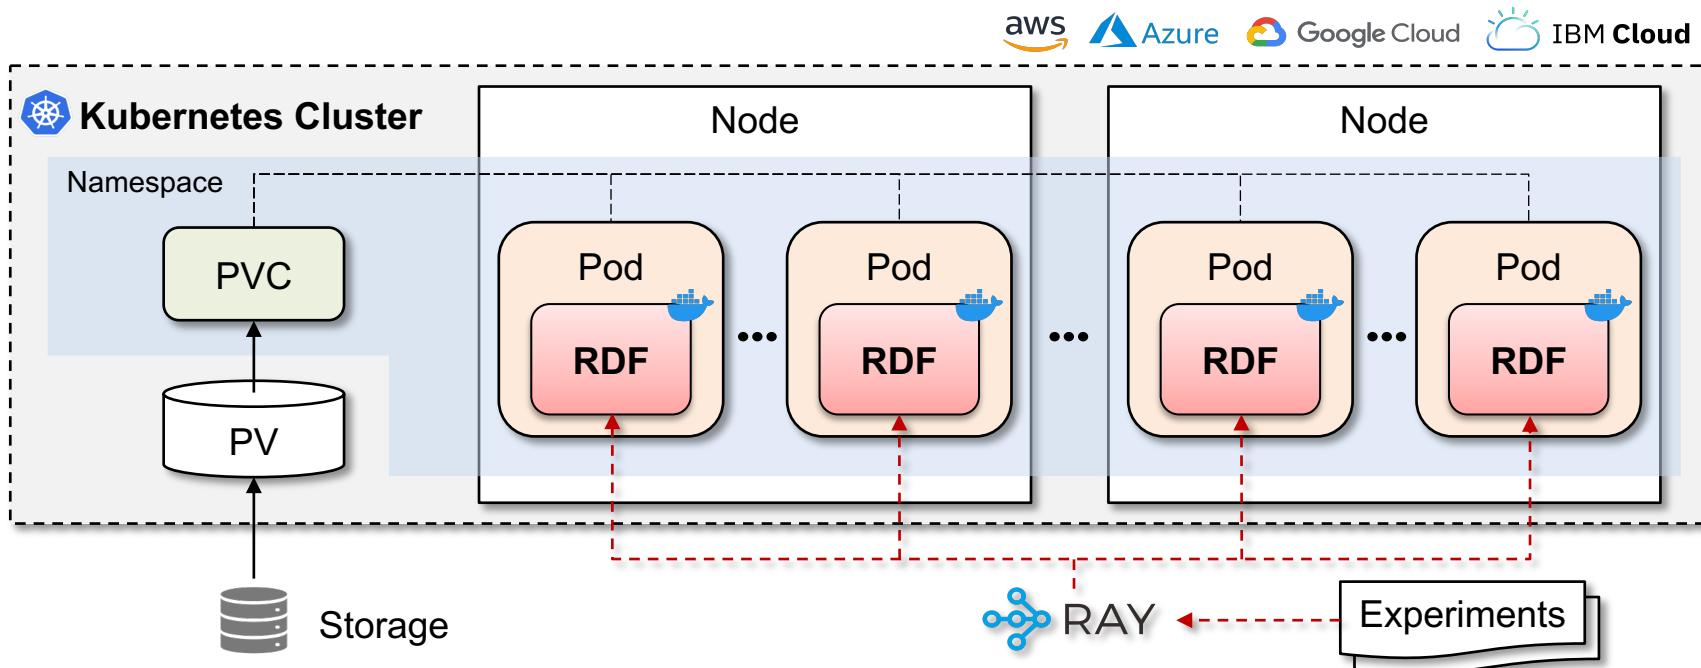
Demonstration

- Consider design experiment to find maximum achievable utilization
 - Given design, PDK, libraries
 - Try many floorplan utilizations, and check whether flow fails or not
- Use Ray and Kubernetes to distribute such experiments efficiently



Scale Out DATC RDF Containers using Public Cloud

- Multi-container Kubernetes cluster still uses *single node*
- Now, we compose **multi-node Kubernetes cluster** using public cloud service to enable even larger-scale design experiments

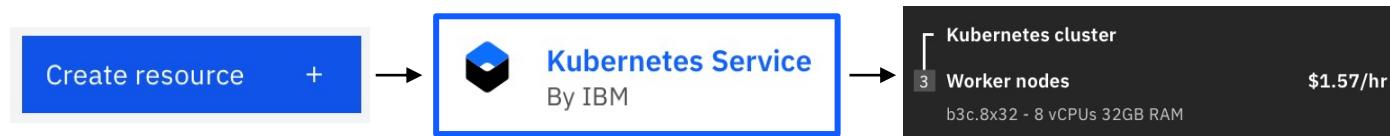


Kubernetes Available on Public Cloud

- Major public cloud providers all offer Kubernetes service
 - **AWS**: Amazon Elastic Kubernetes Service (EKS)
 - **Microsoft Azure**: Azure Kubernetes Service (AKS)
 - **Google Cloud Platform**: Google Kubernetes Engine (GKE)
 - **IBM Cloud**: IBM Cloud Kubernetes Service (IKS)
- Using public cloud, we can create multi-node Kubernetes cluster in minutes with a few clicks

Example: Kubernetes Cluster on IBM Cloud

1. Go to <https://cloud.ibm.com> and create account
2. Create “Kubernetes Service”
3. Choose desired node type, i.e., CPUs and RAM, and create cluster
4. The cluster will be ready in a few minutes



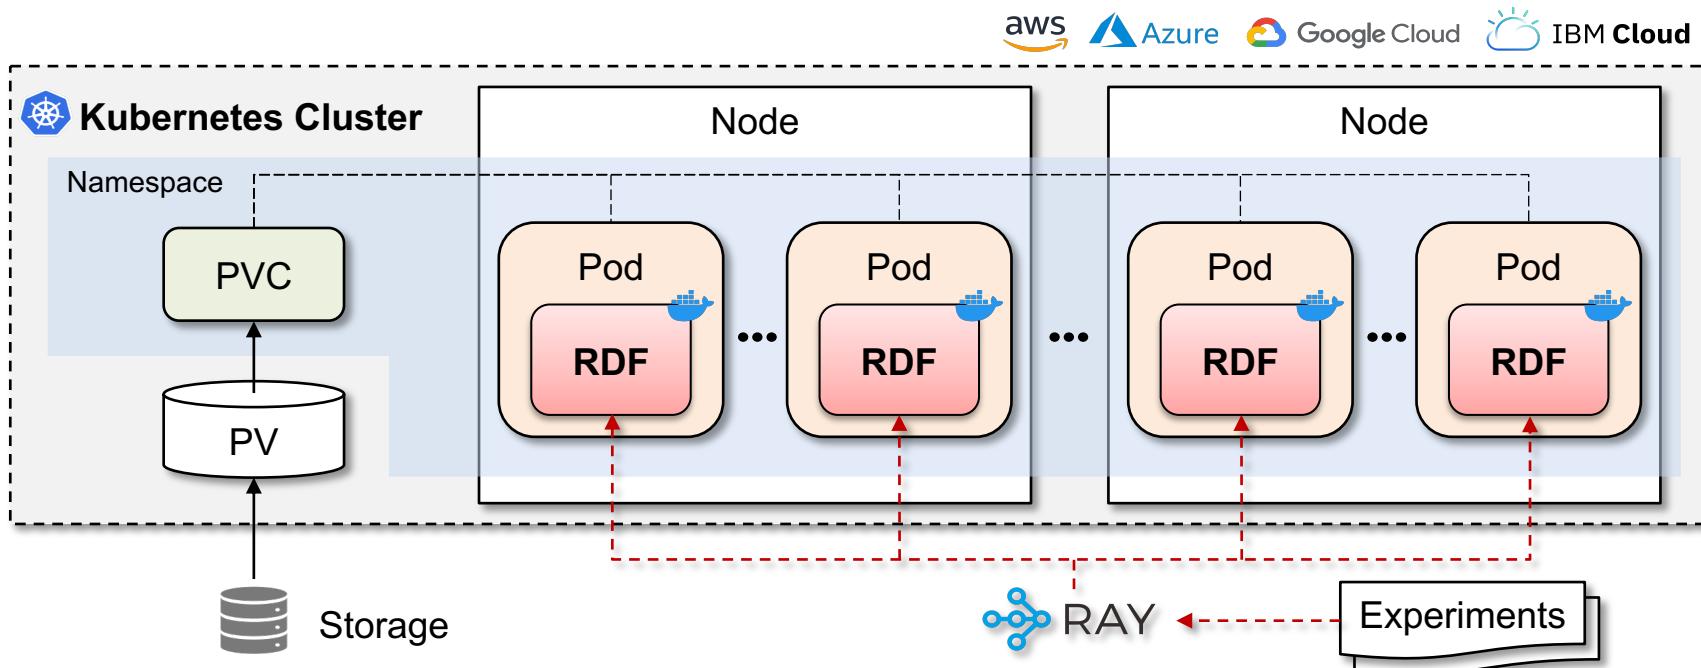
Clusters / rdf-k8s-cluster

Help Kubernetes dashboard

Overview	Node status	Add-on status	Master status
Worker nodes Normal	3 of 3 Normal	0 of 0 Normal	Normal
Worker pools			
DevOps New	Details ↓	Details ↓	Docs ↗

Demonstration

- Once created, Kubernetes cluster on public cloud can be used in same way
- Next demo shows how to deploy distributed design experiments on public cloud using Ray and Kubernetes



Outline

- Introduction and overview
- Docker and Kubernetes
- Kubernetes cluster architecture for design experiments with RDF
- Scaling-out design experiments
 - Multi-container design experiments using Ray
 - Multi-node design experiments using Ray and public cloud services
- **Summary**

Summary

- “**Machine learning requires data**” → Time and compute resources
 - We learned how to utilize public cloud to generate data for ML-CAD
 - We learned how to deploy large-scale experiments on public cloud using Ray and Kubernetes, in *generic/cloud-native way*
- **Code, notebooks:** <https://github.com/ieee-ceda-datc/aspdac-2022-tutorial>
- In this part, metrics collection from data were done in ad-hoc manners
 - “*Can we collect key metrics in a common, standard form?*”
 - **METRICS 2.1:** standardized metrics collection/design process recording → **Part 5**

Part 5

METRICS 2.1: Toward Standard Metrics Collection System

Ravi Varadarajan

ML in EDA and IC Design Requires Data and Standards

- **Standard ML platform(s) for EDA and design process modeling**
 - Enable design metrics collection, tool/flow model generation, design-adaptive tool/flow configuration, prediction of tool/flow outcomes
- **Datasets to support ML**
 - Real and artificial designs
 - Shared training
 - Challenges and incentives: “Kaggle for ML in IC design”

The screenshot shows the METRICS website homepage. The header features the word "METRICS" in large blue letters. Below it is a navigation menu with the following items:

- Main
- Overview
- Infrastructure
 - Metrics List
 - Dictionary
- Publications
 - Papers
 - Presentations
 - GSRC Presentations
- Codes
- Theses
- Links
 - Conferences & Workshops
 - Softwares
 - Others

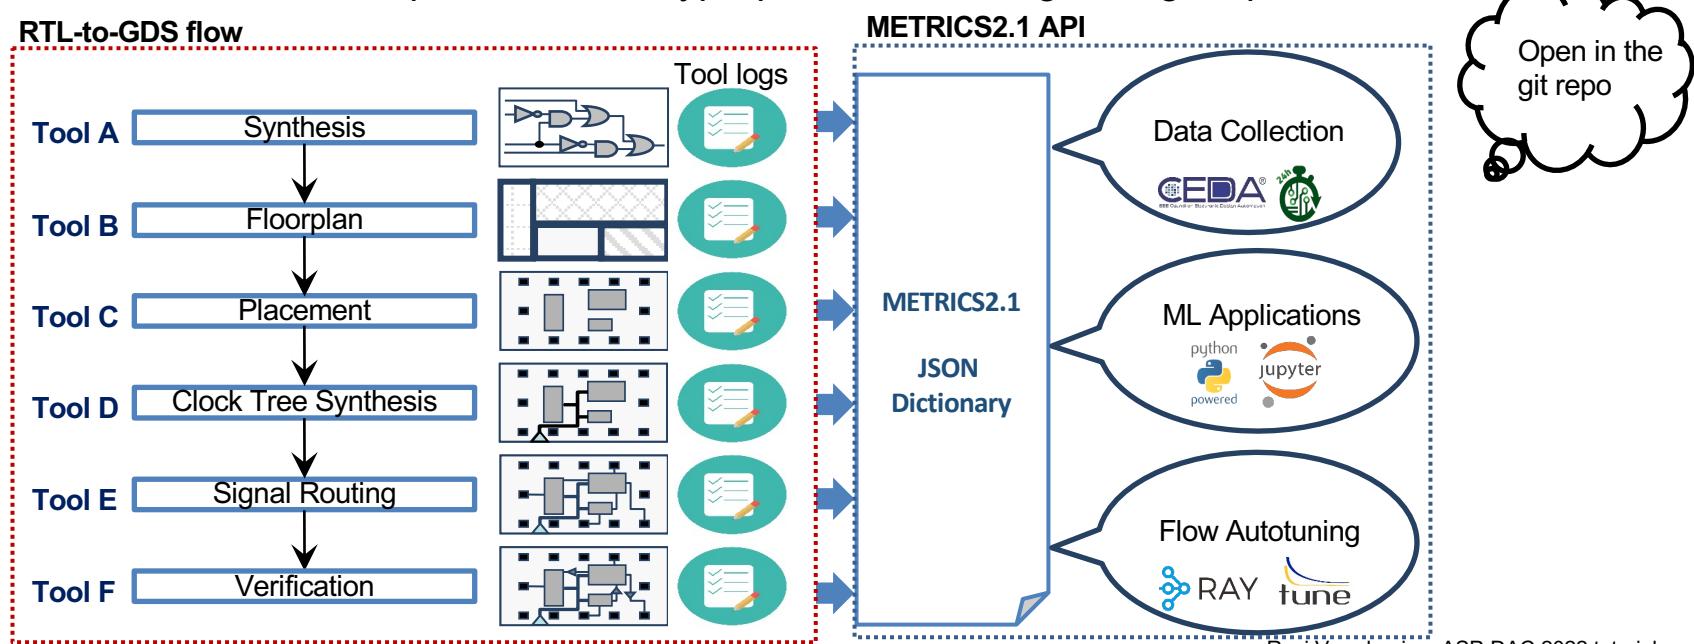
The main content area has a title "The METRICS Initiative" and a section titled "Recent Updates" with the following list:

- [Survey \(1st draft\)](#) for design quality and productivity ([the multiple-choice version](#))
- [Reduced survey \(2nd draft\)](#) for design quality and productivity that is distributed at June 2001 GSRC workshop
- [Updated survey \(3rd draft\)](#) for design quality and productivity that reflects the discussion at June 2001 GSRC workshop
- [Workshop notes](#) for METRICS discussion at June 2001 GSRC workshop
- [List of prediction/estimator models](#) enabled by METRICS System
- [DAC02 Birds-of-a-Feather meeting summary](#) (June 12, 2002)

- **METRICS 1.0** (1999; DAC00, ISQED01)
 - “Measure to Improve” <http://vlsicad.ucsd.edu/GSRC/metrics>
- **METRICS 2.0** ([WOSET-2018](#)) was proposed as an update of **METRICS 1.0**
- **METRICS2.1** is proposed as an extension to **METRICS 2.0**

METRICS2.1 Overview

- An initiative of the **IEEE CEDA Design Automation Technical Committee (DATC)**
- Collection and sharing of design of experiment (DoE) public datasets using the **OpenROAD** tool, with intelligent setting of constraints
- Sample Jupyter notebooks on the datasets to visualize and build simple prediction models
- True “No-Human-in-Loop”: Automatic hyperparameter tuning of engine parameters



METRICS2.1 Principles

- General and extensible
 - Syntax and semantics to support future addition of new metrics
- No ambiguity!
 - Any desired measurement must map to a unique METRICS2.1 metric
 - Every METRICS2.1 metric must map to a unique interpretation as a measurement
- This two-way mapping is crucial to avoid future confusion
- Ability to capture the same metric at different stages of the design flow
- Free, open and frictionless to everyone; agnostic to any EDA vendor

METRICS2.1 Organization

- METRICS2.1 is organized as a hierarchical JSON object
- Top level: predefined *stage* or a user-defined *snapshot*
 - Default stages are *init*, *synth*, *floorplan*, *globalplace*, *detailedplace*, *cts*, *globalroute*, *detailedroute* and *finish*
- Inside each stage: metrics are further organized as a predefined metrics *category*, a metrics *name*, an optional metrics *name modifier*, and optional metrics *classification modifiers*
 - Predefined metrics categories are *flow*, *design*, *timing*, *power* and *route*
- A metric name and its optional *name modifier* uniquely define the metric and its units, e.g., *design_instance_count*
- A metric *classification modifier* subdivides the metric into a specific type (stdcells, macros) or a specific structure (clock domain, analysis view etc.)
- Precedence order is part of the METRICS2.1 specification
- METRICS2.1 naming convention: <https://github.com/ieee-ceda-datc/datc-rdf-Metrics4ML#metrics-naming-convention>

Metrics Stage

Metrics Category

Metrics Name

Metrics Name Modifiers

Metrics Classification
Modifiers

METRICS2.1 Examples

- Sample metrics

Metric	Description
<i>timing_setup_wns</i>	Setup worst negative slack in the design
<i>timing_setup_wns_clock:clk_a</i>	Setup worst negative slack for clock "clk_a" in the design
<i>timing_setup_wns_analysis_view:slow</i>	Setup worst negative slack for the analysis view "slow"
<i>power_total</i>	Total power consumption
<i>power_leakage</i>	Total leakage power
<i>power_leakage_clock</i>	Total leakage power in the clock network

Metrics2.1 Use Model in OpenROAD

- Metrics2.1 used in OpenROAD
 - Monitor and gate the PRs into the OpenROAD repo
 - Create the daily dashboards for OpenROAD and OpenROAD flow
 - <https://github.com/The-OpenROAD-Project/OpenROAD>
 - <https://github.com/The-OpenROAD-Project/OpenROAD-flow-scripts>
 - Quantify improvements to engines and flows
- Metrics2.1 used to collect data from DoE runs
 - Metrics collection of DoE runs to support AutoTuner
 - Metrics collection of DoE runs for visualization and ML applications

OpenROAD Public Dashboards

- OpenROAD public dashboard based on Metrics2.1 is updated daily at jenkins.openroad.tools
- Multiple pipelines are run daily in cloud

The screenshot shows the Jenkins public dashboard interface. On the left, there's a sidebar with links like 'Blue Ocean', 'Buildable Build Status', 'History' (selected), 'Builds...', 'Jan 14, 2022, 7:59 AM', 'Jan 14, 2022, 2:56 AM', 'Jan 13, 2022, 9:59 PM', 'Build Jan 13, 2022, 2:59 PM' (selected), 'No build', 'Jan 13, 2022, 9:19 AM', 'Jan 13, 2022, 3:34 AM', 'Build Jan 12, 2022, 9:59 PM' (disabled), 'Jan 12, 2022, 2:59 PM', 'Build Jan 12, 2022, 9:19 AM' (disabled), 'Jan 12, 2022, 3:28 AM', '1 Idle', 'Jan 12, 2022, 10:01 PM', and 'Jan 11, 2022, 4:43 PM'. The main area displays a grid of pipeline runs. The columns represent different stages: Declarative: Checkout SCM (master branch), Build, Test, Docker, asap7 aes, asap7 ethmac, asap7 gcd, asap7 ibex, asap7 jpeg, asap7 sha3, and asap7 uart. Each row contains a pipeline ID (#1122, #1121, #1120, #1119, #1118, #1117) and its status (e.g., 'No Changes', '8 commits', '1 commit', '9 commits'). The grid cells show various execution times.

	Declarative: Checkout SCM master branch	Build	Test	Docker	asap7 aes	asap7 ethmac	asap7 gcd	asap7 ibex	asap7 jpeg	asap7 sha3	asap7 uart			
#1122	1min 57s	12s	10min 16s	589ms	17min 15s	15min 14s	33min 28s	4min 26s	26min 22s	41min 23s	11min 49s	13min 20s		
#1121	27min	1min 40s	12s	9min 8s	596ms	15min 46s	13min 22s	31min 46s	3min 18s	25min 37s	40min 23s	11min 36s	14min 22s	
#1120	Jan 13 23:59	No Changes	2min 0s	11s	10min 17s	588ms	35min 36s	26min 34s	49min 54s	3min 11s	20min 50s	42min 36s	11min 42s	15min 0s
#1119	Jan 13 13:59	8 commits	2min 6s	13s	14min 35s	605ms	18min 57s	14min 47s	34min 32s	4min 52s	29min 24s	42min 0s	18min 10s	16min 14s
#1118	Jan 13 06:59	1 commit	1min 52s	12s	9min 31s	599ms	15min 41s	13min 19s	31min 31s	4min 45s	22min 37s	43min 4s	11min 11s	14min 31s
#1117	Jan 13 01:19	9 commits	1min 55s	12s	9min 49s	629ms	8min 51s	15min 0s	33min 47s	4min 58s	22min 3s	44min 26s	11min 57s	14min 5s

OpenROAD flow scripts “rules.json”

- OpenROAD flow scripts uses “rules.json” to determine pass or fail of designs based on different metrics.
- Currently 14 metrics are checked for each design
- rules.json can be automatically created from metrics data
- Sample rules.json for sky130HD “ibex_core” is shown

```
{  
    "field": "cts_timing_setup_ws_pre_repair",  
    "value": -2.75,  
    "compare": ">="  
},  
{  
    "field": "cts_timing_setup_ws_post_repair",  
    "value": -2.75,  
    "compare": ">="  
},  
{  
    "field": "globalroute_timing_clock_slack",  
    "value": -2.12,  
    "compare": ">="  
},  
{  
    "field": "globalroute_timing_setup_ws",  
    "value": -2.12,  
    "compare": ">="  
},  
{  
    "field": "detailedroute_route_wirelength",  
    "value": 915751,  
    "compare": "<="  
},  
{  
    "field": "detailedroute_route_drc_errors",  
    "value": 0,  
    "compare": "<="  
},  
{  
    "field": "finish_timing_setup_ws",  
    "value": -2.75,  
    "compare": ">="  
},  
{  
    "field": "finish_design_instance_area",  
    "value": 231165,  
    "compare": "<="  
}
```

Sample Collection of DoE Runs

https://github.com/ieee-ceda-datc/datc-rdf-Metrics4ML/tree/main/experiments/sky130hd_ibex_core_sample



Sample METRICS2.1 Application: ibex, SKY130HD

- **ibex** = lowRISC RISC-V core [<https://github.com/lowRISC/ibex>] (15,708 instances)
- Design run with OpenROAD flow scripts
- Parameters swept:
 - Core utilization of 20% to 40% in increments of 2%
 - Layer track resource adjustment for M1 to M4 from 0.1 to 0.2 in steps of .025
- Metrics of interest:
 - Successful or Doomed (failure in Global or Detailed Route)
 - Route runtime, DRC errors
 - Routed wirelength
- Build prediction models for Success/Failure and estimated wirelength for new parameter settings

Sample METRICS2.1 Application

- https://github.com/ieee-ceda-datc/datc-rdf-Metrics4ML/tree/main/experiments/sky130hd_ibex_core_sample_fine_grained contains the data and the Jupyter Notebook
- Jupyter Notebook:
 - Read METRICS data
 - Data engineering to pull, organize relevant data
 - Visualize trends, statistics
- Leverages python packages json, pandas, matplotlib, plotly, sklearn
- Use sklearn
 - Logistic regression to predict successful/doomed runs
 - Linear regression to estimate wirelength, runtime

nbviewer
data-rdf-Metrics4ML / experiments / sky130hd_ibex_core_sample_fine_grained

```
In [1]: import numpy as np
import pandas as pd
import json
from pandas.io.json import json_normalize
import matplotlib.pyplot as plt
import seaborn as sns
import glob, os

from datetime import date, timedelta

import plotly
import plotly.graph_objects as go
import plotly.express as px
from plotly.subplots import make_subplots

%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
pd.set_option('display.max_colwidth', -1)
```

Design : ibex

Platform: sky130hd

A sample dataset with multiple runs varying the design utilization and the layer_adjust parameters for the routing layers. All design metrics from the runs are collected for analysis.

Metrics data is represented as json files in the METRICS2.1 format. Each experiment in the run is a separate json file. All of the files are read into a DataFrame 'json_df'.

- Rows of the DataFrame represent an experiment.
- Columns of the DataFrame represent the metrics.

```
In [2]: path = './fine-grain-metrics_new'
json_df = pd.DataFrame()
for filename in glob.glob(os.path.join(path, '*.json')):
    with open(os.path.join(os.getcwd(), filename), 'r') as f:
        data = f.read()
        data_json = json_normalize(json.loads(data))
        json_df = json_df.append(data_json)
```

```
In [3]: print(f'Number of runs in the dataset: {json_df.shape[0]})')
print(f'Number of metrics in each data: {json_df.shape[1]})')
```

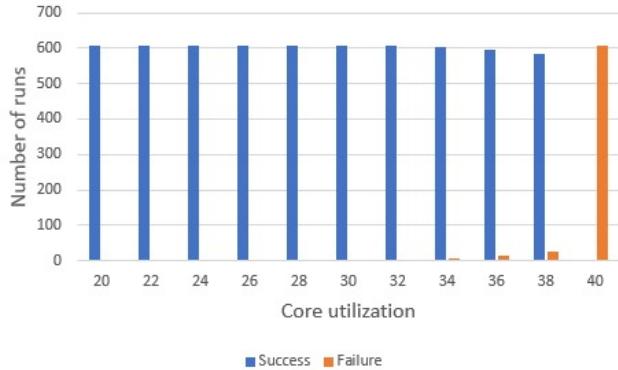
Number of runs in the dataset: 6699
Number of metrics in each data: 87

- The columns in the DataFrame, which are the metrics collected for each run.

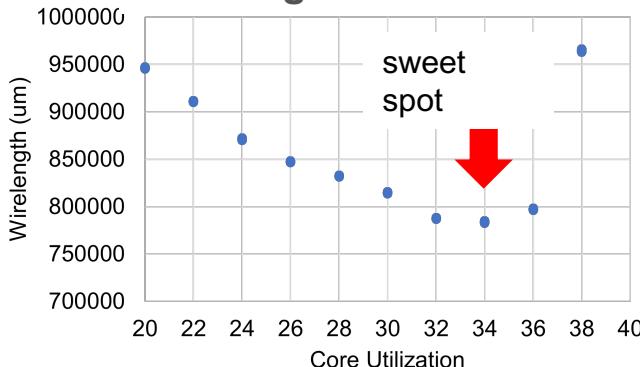
```
In [4]: json_df.columns
Out[4]: Index(['run_flow_generate_date', 'run_flow_metrics_version',
   'run_flow_openroad_version', 'run_flow_openroad_commit',
   'run_flow_scripts_commit', 'run_flow_uid', 'run_flow_design',
   'run_flow_platform', 'run_flow_platform_commit', 'run_flow_variant',
   'synth.design_instance_stcell_count',
   'synth.design_instance_stcell_area', 'synth.runtime_total',
   'runth_rms_total', 'runth_max_mask'])
```

Sample Plots from Jupyter Notebook

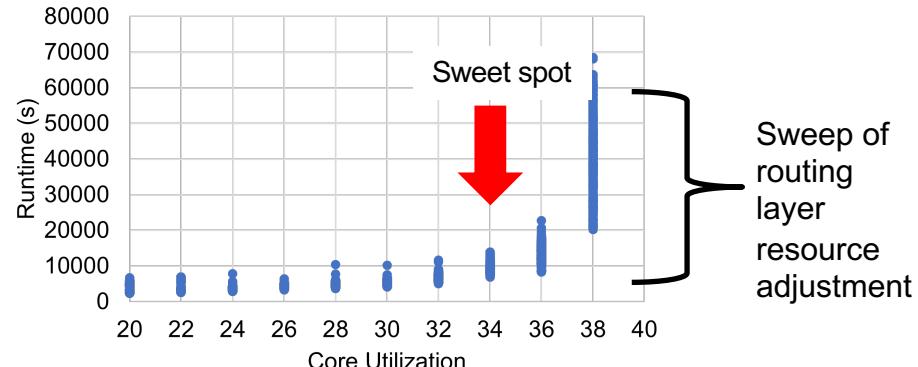
Successes vs. Failures at Different Utilizations



Wirelength vs. Core Utilization

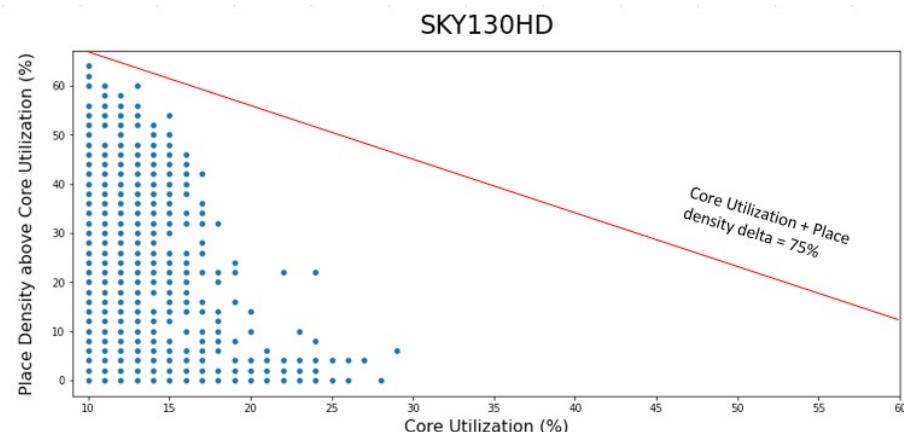
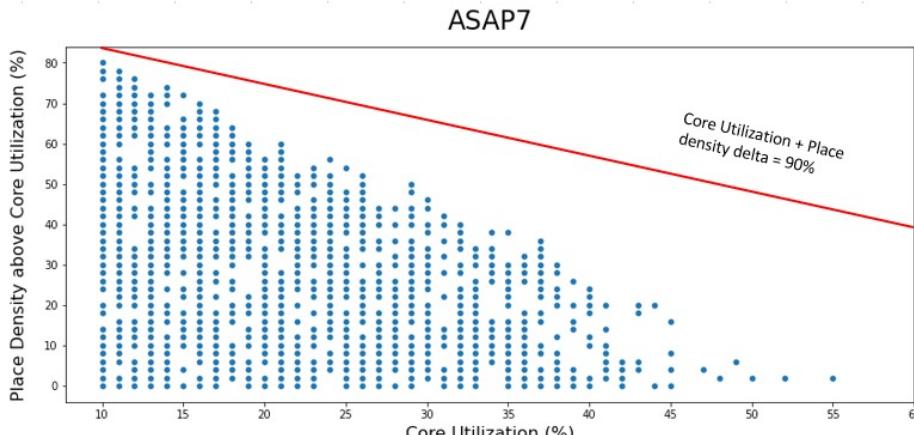


Runtime vs. Core Utilization



Another Sample METRICS2.1 Application

- Design: ibex on ASAP7 and SKY130HD platforms
- Parameters swept
 - Core Utilization: 10% to 60% in increments of 1%
 - Place Density (RePIAce) delta above Core Utilization: 0% to 90% in increments of 2%
 - 2295 data entries for each platform
- Metrics of interest
 - Successful or Doomed
 - Routed Wirelength
- Shown in Plots
 - 728 of 2295 ASAP7 runs are successful
 - 252 of 2295 SKY130HD runs are successful



Summary

- **METRICS2.1 infrastructure**
 - An initiative of the IEEE CEDA Design Automation Technical Committee (DATC)
 - Enables standardized metrics collection and design process recording
 - Used in the OpenROAD project for CI and track QoR
 - We provide archives, reproducible config files and ML applications in the form of Jupyter notebooks
- **We invite you for collaboration!**
 - METRICS2.1
 - <https://github.com/ieee-ceda-datc/datc-rdf-Metrics4ML>

THANK YOU ! (Questions / Follow ups?)

We thank the OpenROAD team and the IEEE CEDA Design Automation Technical Committee for many helpful interactions.

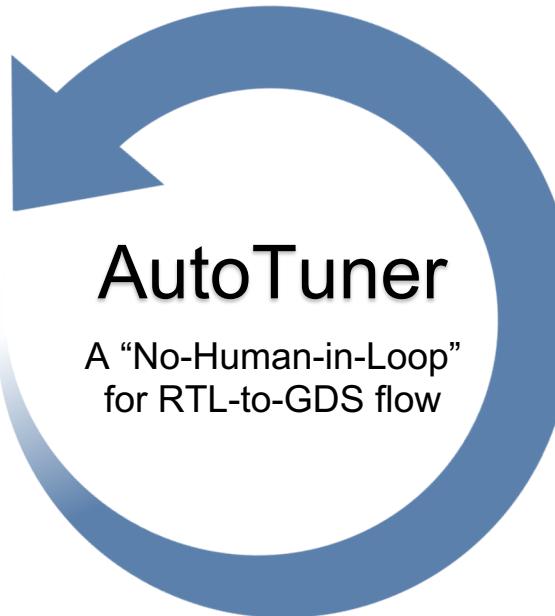
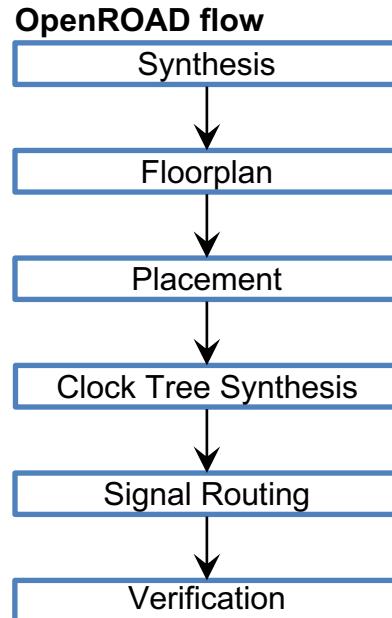
This work is supported in part by the U.S. National Science Foundation (CCF-2112665) and the U.S. Defense Advanced Research Projects Agency (HR0011-18-2-0032).
<https://theopenroadproject.org> <https://tilos.ai> <https://vlsicad.ucsd.edu>

Part 6

AutoTuner: Autonomous RTL-to-GDS Flow Tuning

Seungwon Kim

AutoTuner: True No-Human-in-the-Loop Operation



Best human-tuned flow
vs.
Auto-tuned flow
(No Human In Loop)

23% Faster!
45% Less Power Consumption!

Tested in SKY130HD
lowRISC RISC-V core (ibex)

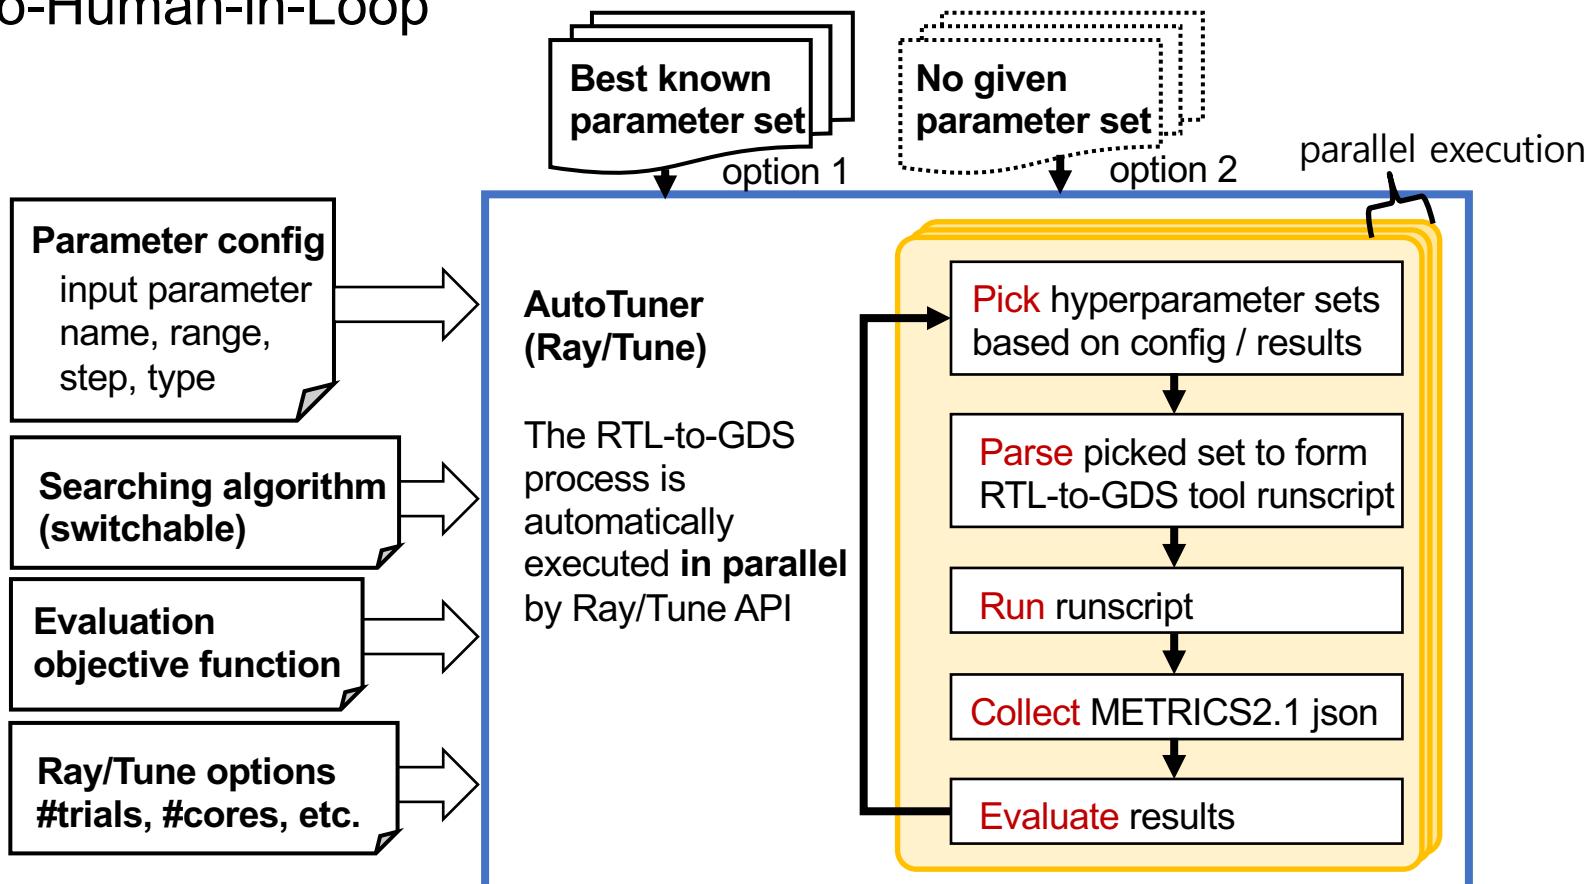
Note: AutoTuner being deployed on GCP for
public enablements (SKY130*, ASAP7).

Flow Parameter AutoTuner – Overview

- A “No-Human-in-Loop” for RTL-to-GDS flow using the **OpenROAD** tool
- Automatic iterative tuning for **QoR improvement** within a given hyperparameter range space
- Interface: Python packages Ray/Tune
- Search algorithms: HyperOpt, PBT, Optuna, Nevergrad, Ax, random search
- Advantages:
 - Pre-existing big data is not necessary for execution
 - Fewer trials compared to random or grid parameter sweeping
 - Powerful parallelization management: core/thread management, external server usage, web visualization
 - Suboptimality gap reduces with compute resource, schedule budget

Flow Parameter AutoTuner – Architecture

- No-Human-in-Loop



Case Study: SkyWater 130HD, ibex (1/3)

- Tech: Skywater 130nm HD
- Design: ibex
- Tested Algorithms
 - HyperOpt
- Parameter config

Assuming full factorial combinations,
1,058,298,150
= #possible combinations!

```
1 {  
2     "GP_PAD": {"type": "int", "minmax": [0,4], "step": 1 },  
3     "DP_PAD": {"type": "int", "minmax": [0,4], "step": 1 },  
4     "LAYER_ADJUST": {"type": "float", "minmax": [0.1,0.7], "step": 0 },  
5     "PLACE_DENSITY": {"type": "float", "minmax": [0.1,1.0], "step": 0 },  
6     "FLATTEN": {"type": "int", "minmax": [0,1], "step": 1 },  
7     "PINS_DISTANCE": {"type": "int", "minmax": [1,3], "step": 1 },  
8     "CTS_CLUSTER_SIZE": {"type": "int", "minmax": [10,40], "step": 1 },  
9     "CTS_CLUSTER_DIAMETER": {"type": "int", "minmax": [80,120], "step": 1 },  
10    "GR_OVERFLOW": {"type": "int", "minmax": [1,1], "step": 0 }  
11 }
```

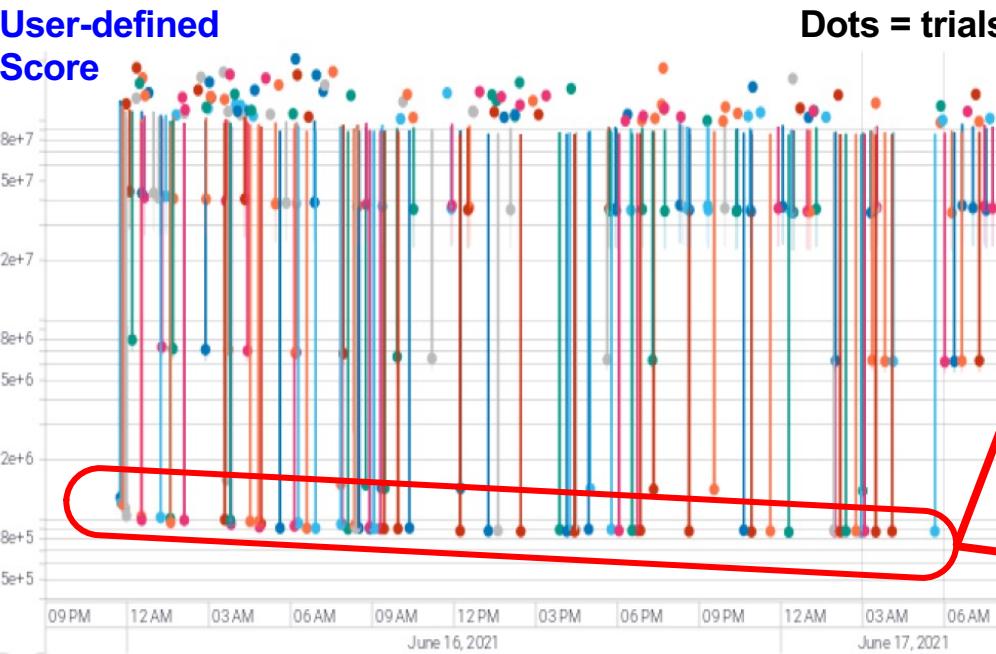
- HyperParameter name, type, minmax, step.
- When type is int and step = 0, it means constant value
- When type is float and step = 0, it means continuous range

Case Study: SkyWater 130HD, ibex (2/3)

- GUI feature is supported with **TensorBoard**
- Results over Wall Time

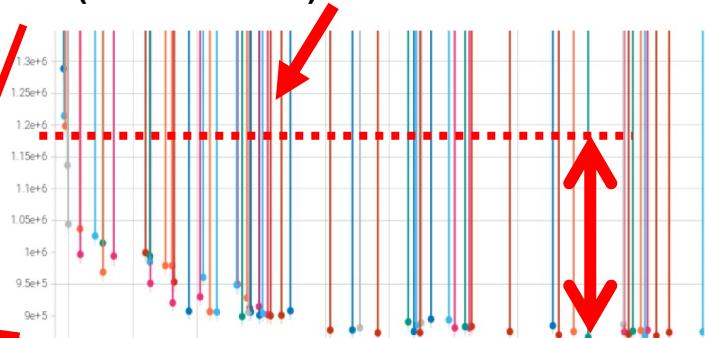
User-defined

Score



Dots = trials

Default flow score = 1,174,346
Our Best Score = 855,373
(370 trials in total 500 #trials)
(less is better)



Improvement

WL 1003801um → 843258um (-16%)
Effective CP 20.935ns → 16.185 ns (-23%)
Total power 0.024 W → 0.0133 W (-45%)

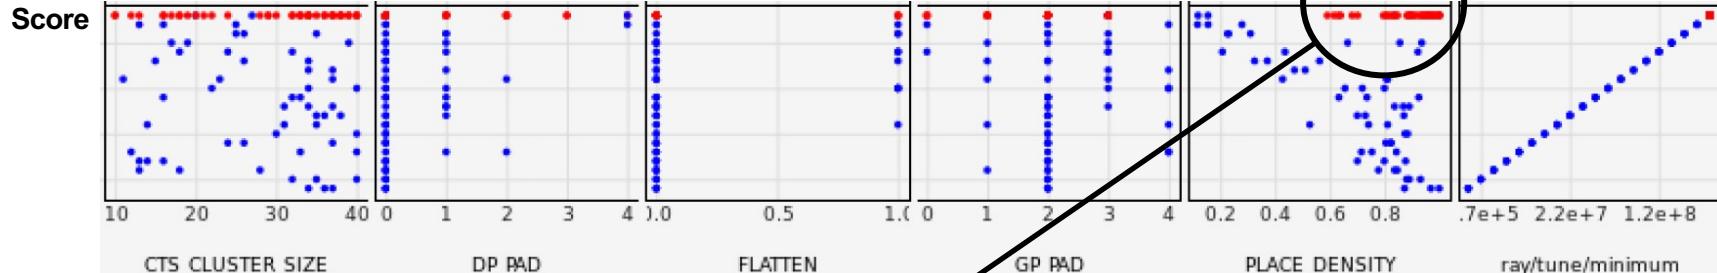


Results over Wall Time

Time

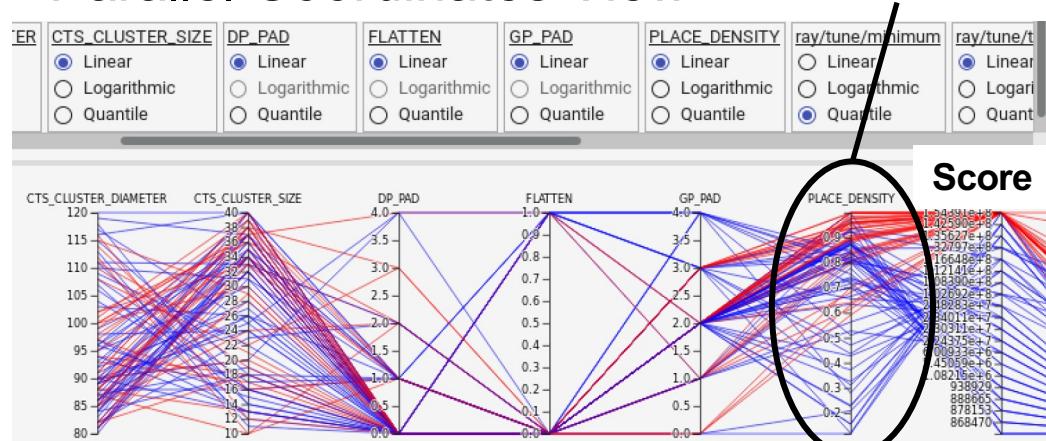
Case Study: SkyWater 130HD, ibex (3/3)

- Hyperparameter Scatter Plot Matrix View



- Parallel Coordinates View

high risk, high return



Parameters	Human	AutoTuner
Flatten (0 = hier)	1	0
Global placement padding	4	2
Detailed placement padding	2	0
IO pin distance (um)	2	2
CTS clustering size (um)	30	37
CTS clustering diameter (um)	100	95
Layer resource adjustment	0.5	0.42
Target GP place density	0.6	0.99

Search Algorithms

- Types of search algorithms currently supported by AutoTuner

Types	Name	Year	Author	Affiliation
Random/Grid search	Random /Grid search			
Population Based Training	PBT	2017	M. Jaderberg	Google Deepmind
Tree Parzen Estimator (TPE)	HyperOpt	2013	J. Bergstra	Rowland Institute at Harvard
Bayesian + Multi-Armed Bandit	AxSearch	2019	B. Letham	Facebook
TPE + CMA-ES[†]	Optuna	2019	T. Akiba	Preferred Network, Inc.
Evolutionary Algorithm (Gradient-free optimization)	Nevergrad	2018	J. Rapin	Facebook

[†]CMA-ES: Covariance matrix adaptation evolution strategy

- Each algorithm has its pros and cons depending on:

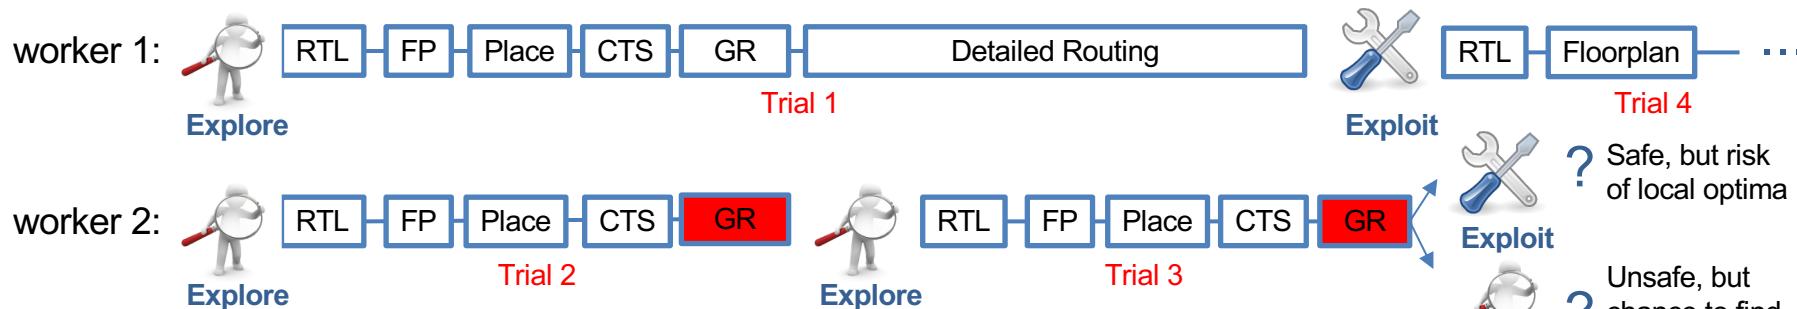
- Difficulty of the target black box problem
- Number of hyperparameters
- Type of hyperparameters (continuous, discrete, mixed)
- Parallelization, and whether synchronous / asynchronous

Search Algorithm Usage Strategy: Considerations (1)

- Considerations for **RTL-to-GDS** domain
 - **Problem size:** Big
 - Long runtime for each trial
 - **Difficult:** Many infeasible input parameter combinations for hyperparameters
 - Active exploration is required to find feasible starting points
 - **Mixed types:** continuous and discrete parameters
 - Some Bayesian optimization algorithms do not support discrete parameters
 - **CAD Tool noise:** influence on the final QoR from very small variations in input values

Search Algorithm Usage Strategy: Considerations (2)

- Considerations for algorithm selection
 - Most time-consuming stage is detailed routing
 - The runtime required per trial varies greatly
 - Each algorithm has different strategy of exploration and exploitation



- Significant difference in total walltime with **same #trials**
- Many infeasible hyperparameter sets

- Results should be compared by the **same walltime constraint**
- Active exploration with **asynchronous scheduling** helps

? Safe, but risk of local optima
? Unsafe, but chance to find better solution
Depends on the search algorithm

Objective Function – Power, Performance and Area (PPA)

- IC design is a multi-objective optimization (e.g., power, performance and area)
- We integrate three PPA metrics into one objective function
 - To increase the general availability
 - To use algorithms that do not support multi-objective functions
- User-settable coefficient values ($C_{\text{power}} \cdot C_{\text{perform}} \cdot C_{\text{area}}$) of three objectives to set the direction of tuning

Minimize $\text{Score} = (\text{PPA}_{\text{imp}}^{\text{UB}} - \text{PPA}_{\text{imp}}) \cdot (1 + \alpha \cdot N_{\text{DRV}})$ α : penalty factor (default: 0.1)

s.t.

$$\begin{aligned}\text{PPA}_{\text{imp}} &= (C_{\text{power}} \cdot \text{Power}_{\text{imp}} + C_{\text{perform}} \cdot \text{Perform}_{\text{imp}} + C_{\text{area}} \cdot \text{Area}_{\text{imp}}) \\ \text{PPA}_{\text{imp}}^{\text{UB}} &= (C_{\text{power}} + C_{\text{perform}} + C_{\text{area}}) \\ \text{Power}_{\text{imp}} &= (P_{\text{ref}} - P)/P_{\text{ref}} \\ \text{Perform}_{\text{imp}} &= (\text{effCP}_{\text{ref}} - \text{effCP})/\text{effCP}_{\text{ref}} \\ \text{Area}_{\text{imp}} &= (A_{\text{ref}} - A)/(A_{\text{ref}})\end{aligned}$$

P : Total power
effCP: Effective clock period
A: Cell area
ref: Reference metric

Demonstration of AutoTuner

- AutoTuner provides two main functionalities as follows
 - **Automatic hyperparameter tuning framework** for OpenROAD-flow-scripts
 - **Parametric sweeping experiments** for OpenROAD-flow-scripts
- Installation of **OpenROAD-flow-scripts (ORFS)** is required
 - <https://github.com/The-OpenROAD-Project/OpenROAD-flow-scripts>
- In this demonstration, the following processes are included
 - Tunable / sweepable parameters
 - How to define inputs
 - Input JSON structure (autotuner.json)
 - User-settable coefficient values (distributed.py)
 - List of input arguments
 - How to run in shell by AutoTuner in ORFS
 - How to find the logs and results
 - How to use GUI
- Instructions for AutoTuner
 - <https://github.com/The-OpenROAD-Project/OpenROAD/blob/master/docs/user/InstructionsForAutoTuner.md>

Demonstration – Tunable / Sweepable Parameters (1)

- **Any variable** that can be set from the command line can be used
 - Technology platform config variables (with "?=" type definition)
 - Design config variables
 - The other environmental variables used in ORFS
- **Technology platform variables** (e.g., ASAP 7nm)

- Location of technology platform variables

config file: OpenROAD-flow-scripts/flow/platforms/asap7/config.mk

- Example variables defined by type "?="

```
# Cell padding in SITE widths to ease rout-ability. Applied to both sides
export CELL_PAD_IN_SITES_GLOBAL_PLACEMENT ?= 2
export CELL_PAD_IN_SITES_DETAIL_PLACEMENT ?= 1
```

- The reason why we need to define "?=" type for platform config
 - ORFS config reading order: design config → AutoTuner→ platform config
 - "?=" only has an effect if the variable is not yet defined

Demonstration – Tunable / Sweepable Parameters (2)

- **Design config variables** (e.g., aes design, ASAP 7nm)

- Location of design config variables

config file: OpenROAD-flow-scripts/flow/designs/asap7/aes/config.mk

- Example variables

```
export CORE_UTILIZATION      = 30
export CORE_ASPECT_RATIO     = 1
export CORE_MARGIN           = 2
```

- Duplicated design config variables will be overwritten by AutoTuner

Demonstration – Tunable / Sweepable Parameters (3)

- The other environmental variables (1)

- Location of flow scripts in ORFS

scripts path: OpenROAD-flow-scripts/flow/scripts

- Example variable in flow script (detailed_route.tcl):

```
if { [info exists ::env(OR_SEED)] } {  
    append additional_args " -or_seed $::env(OR_SEED)"
```

- As above example, any env variable in flow scripts is Tunable / Sweepable

Demonstration – Tunable / Sweepable Parameters (4)

- The other environmental variables (2)

- For SDC:

<code>_SDC_FILE_PATH</code>	Path relative to the current JSON file to the SDC file.
<code>_SDC_CLK_PERIOD</code>	Design clock period. This will create a copy of <code>_SDC_FILE_PATH</code> and modify the clock period.
<code>_SDC_UNCERTAINTY</code>	Clock uncertainty. This will create a copy of <code>_SDC_FILE_PATH</code> and modify the clock uncertainty.
<code>_SDC_IO_DELAY</code>	I/O delay. This will create a copy of <code>_SDC_FILE_PATH</code> and modify the I/O delay.

- For FastRoute (global router):

<code>_FR_FILE_PATH</code>	Path relative to the current JSON file to the fastroute.tcl file.
<code>_FR_LAYER_ADJUST</code>	Layer adjustment. This will create a copy of <code>_FR_FILE_PATH</code> and modify the layer adjustment for all routable layers, i.e., from <code>\$MIN_ROUTING_LAYER</code> to <code>\$MAX_ROUTING_LAYER</code> .
<code>_FR_LAYER_ADJUST_NAME</code>	Layer adjustment for layer NAME. This will create a copy of <code>_FR_FILE_PATH</code> and modify the layer adjustment only for the layer NAME.
<code>_FR_GR_SEED</code>	Global route random seed. This will create a copy of <code>_FR_FILE_PATH</code> and modify the global route random seed.

Demonstration – How to define inputs (1)

- **Input JSON structure (autotuner.json)**

- A part of sample JSON file for ASAP 7nm aes design:

```
{  
    "_SDC_FILE_PATH": "constraint.sdc",  
    "_S  
        "type": "PERIOD",  
        "minmax": [  
            100,  
            600  
        ],  
        "step": 0  
,  
    "CORE_MARGIN": [  
        "type": "int",  
        "minmax": [  
            2,  
            2  
        ],  
        "step": 0  
,  
    ...  
}
```

- Parameter names for sweeping/tuning
- Parameter type ("float" or "int") for sweeping/tuning
- Min-to-max range for sweeping/tuning. The unit follows the default value of each technology std cell library
- Parameter step within the minmax range.
Step 0 for type "float" means continuous step.
Step 0 for type "int" means a constant parameter

path: OpenROAD-flow-scripts/flow/designs/asap7/aes/autotuner.json

Demonstration – How to define inputs (2)

- User-settable coefficient values (distributed.py)

- A part of distributed.py :

```
class PPAImprov(AutoTunerBase):  
    ...  
  
    PPAImprov  
    ...  
  
    @classmethod  
    def get_ppa(cls, metrics):  
        ...  
  
        Compute PPA term for evaluate.  
        ...  
  
        coeff_perform, coeff_power, coeff_area = 10000, 100, 100
```

- Function “get_ppa” in class “PPAImprove” contains coefficient variables
- User can adjust three coefficient to set the PPA optimization direction of tuning

path: OpenROAD-flow-scripts/flow/util/distributed.py

Demonstration – List of input arguments (1)

- **Input arguments are classified into five categories**
 - Target Design
 - Experiment Setup
 - Git Setup
 - AutoTuner Specific Arguments
 - Workload
- For arguments with default settings, refer to distributed.py

Demonstration – List of input arguments (2)

- **Target Design**

- **--design**
 - Name of the design for autotuning
- **--platform**
 - Name of the platform for autotuning

- **Experiment setup**

- **--config**
 - Configuration file that sets which knobs to use for autotuning
- **--experiment**
 - Experiment name. This parameter is used to prefix the **FLOW_VARIANT** and to set the Ray log destination

Demonstration – List of input arguments (3)

- **Git setup**

- **--git-clean**
 - Clean binaries and build files
- **--git-clone**
 - Force new git clone
- **--git-clone-args**
 - Additional git clone arguments
- **--git-latest**
 - Use the latest version of OR app
- **--git-or-branch**
 - OR app branch to use
- **--git-orfs-branch**
 - ORFS branch to use
- **--git-url**
 - ORFS repo URL to use
- **--build-args**
 - Additional arguments given to
./build_openroad.sh

Demonstration – List of input arguments (4)

- AutoTuner Specific Arguments

- **--algorithm**
 - Search algorithm to use for autotuning
- **--eval**
 - Evaluate function to use with search algorithm
- **--samples**
 - Number of samples for autotuning
- **--iterations**
 - Number of iterations for autotuning
- **--resume**
 - Resume previous run
- **--reference**
 - Reference METRICS2.1 json for "PPAImprov" evaluation
 - **If not defined**, "AutoTunerBase" will be used (performance opt only)
- **--perturbation**
 - Perturbation interval for PopulationBasedTraining
- **--seed**
 - Random seed for parameter selection during autotuning

Demonstration – List of input arguments (5)

- **Workload**

- **--jobs**
 - Max number of concurrent jobs
- **--openroad-threads**
 - Max number of threads OR app can use
- **--server**
 - The address of Ray server to connect
- **--port**
 - The port of Ray server to connect
- **-v, --verbose**
 - Verbosity level
 - 0: only print Ray status
 - 1: also print training stderr
 - 2: also print training stdout

Demonstration – How to run in the shell (1)

- Run in Shell
 - **distributed.py** scripts handles sweeping and tuning of ORFS parameters
 - AutoTuner run example (SKY130nm HD ibex design):

```
shell) python3 distributed.py --design ibex --platform sky130hd \
    --config ../designs/sky130hd/ibex/autotuner.json \
    --jobs 40 --openroad-threads 4 \
    tune --samples 1000
```
- The order of the parameters matter
 - Arguments **--design**, **--platform** and **--config** are always required and should precede
 - **AutoTuner specific arguments** must be typed after “**tune**”

launch path: OpenROAD-flow-scripts/flow/util

Demonstration – How to run in the shell (2)

```
shell) python3 distributed.py --design ibex --platform sky130hd \
    --config ../designs/sky130hd/ibex/autotuner.json \
    --jobs 40 --openroad-threads 4 \
    tune --samples 1000
```

Launch

```
(swkim38) [swkim@autotuner-1 util]$ python3 distributed.py --design ibex --platform sky130hd      --config ../designs/sky130hd/ibex/autotuner.j
son      --jobs 40 --openroad-threads 4  tune --samples 1000
2022-01-04 13:44:22,736 INFO services.py:1272 -- View the Ray dashboard at http://127.0.0.1:8265
== Status ==
Memory usage on this node: 120.9/3783.4 GiB
Using AsyncHyperBand: num_stopped=0
Bracket: Iter 64.000: None | Iter 16.000: None | Iter 4.000: None | Iter 1.000: None
Resources requested: 0/160 CPUs, 0/0 GPUs, 0.0/3478.44 GiB heap, 0.0/186.26 GiB objects
Result logdir: /home/swkim/01_OpenROAD/daily_runs/OpenROAD_daily_2021-12-28_00-07/OpenROAD-flow-scripts/flow/logs/sky130hd/ibex/test-tune-2022-01-04-13-44-21
Number of trials: 1/1000 (1 PENDING)
+-----+-----+-----+-----+-----+
|-----+-----+-----+-----+-----+
| Trial name          | status   | loc     | CELL_PAD_IN_SITES_DETAIL_PLACEMENT | CELL_PAD_IN_SITES_GLOBAL_PLACEMENT | CORE_ASPECT_RATIO |
| CORE_MARGIN         | CORE_UTILIZATION | CTS_CLUSTER_DIAMETER | CTS_CLUSTER_SIZE | PLACE_DENSITY_LB_ADDON | _FR_GR_OVERFLOW | _FR_LAYER_AD
JUST | _PINS_DISTANCE | _SDC_CLK_PERIOD | _SYNTH_FLATTEN |           |           |           |
|-----+-----+-----+-----+-----+
|-----+-----+-----+-----+-----+
| AutoTunerBase_6e9e8a08 | PENDING |          |          |          |          |          |
| 2 |          98 |          |          |          |          |          |
4004 |          2 |          10.1019 |          1 |          20 |          0.763522 |          4 |          1 |          0.860483 |
|-----+-----+-----+-----+-----+
|-----+-----+-----+-----+-----+
|-----+-----+-----+-----+-----+
```

status of server and tunning

status of initial trial

Demonstration – How to run in the shell (2)

```
shell) python3 distributed.py --design ibex --platform sky130hd \
    --config ..../designs/sky130hd/ibex \
    --jobs 40 --openroad-threads 4 | tee log.txt
```

In progress of autotuning

Current best trial: 6e9e8a08 with minimum=9999999999900.0 and parameters={'_SDC_CLK_PERIOD': 10.101861658418082, 'CORE_UTILIZATION': 98, 'CORE_ASPECT_RATIO': 0.8604828491607834, 'CORE_MARGIN': 2, 'CELL_PAD_IN_SITES_GLOBAL_PLACEMENT': 4, 'CELL_PAD_IN_SITES_DETAIL_PLACEMENT': 1, '_FR_LAYER_ADJUST': 0.6240038814181154, 'PLACE_DENSITY_LB_ADDON': 0.763522369775535, '_SYNTH_FLATTEN': 1, '_PINS_DISTANCE': 2, 'CTS_CLUSTER_SIZE': 20, 'CTS_CLUSTER_DIAMETER': 116, '_FR_GR_OVERFLOW': 1}

Number of trials: 47/1000 (39 RUNNING, 8 TERMINATED) | current progress | current best

NUMBER OF TRIALS: 49/1000 (39 RUNNING, 8 TERMINATED)

current progress

current best

Trial name	status	loc	CELL_PAD_IN_SITES_DETAIL_PLACEMENT	CELL_PAD_IN_SITES_GLOBAL_PLACEMENT	CORE_ASPECT_RATIO	
CORE_MARGIN	CORE_UTILIZATION	CTS_CLUSTER_DIAMETER	CTS_CLUSTER_SIZE	PLACE_DENSITY_LB_ADDON	_FR_GR_OVERFLOW	_FR_LAYER
ADJUST	_PINS_DISTANCE	_SDC_CLK_PERIOD	_SYNTH_FLATTEN	iter	total time (s)	minimum
AutoTunerBase_6ec6d760	RUNNING			4		
398468	2	57	83	36	0.9	0.571708
AutoTunerBase_6eda5bfa	RUNNING			2		0.
381646	2	53	112	20	0.825331	1.55892
AutoTunerBase_6f140e04	RUNNING			3		0.
636041	2	37	85	26	0.879843	1.45467
AutoTunerBase_6f277958	RUNNING			1		0.
567149	2	76	112	15	0.0445132	0.486949
AutoTunerBase_6f619dcc	RUNNING			3		0.
31722	2	20	108	24	0.869995	0.866958
		5.49068				0.

Demonstration – How to run in the shell (3)

```
shell) python3 distributed.py --design ibex --platform sky130hd \
    --config ../designs/sky130hd/ibex/autotuner.json \
    --jobs 40 --openroad-threads 4 \
    tune --samples 1000
```

best parameters

```
2022-01-05 13:58:10,963 INFO tune.py:549 -- Total run time: 87225.26 seconds (87224.14 seconds for the tuning loop).
[INFO TUN-0002] Best parameters found: {'_SDC_CLK_PERIOD': 8.64722781707361, 'CORE_UTILIZATION': 20, 'CORE_ASPECT_RATIO': 1.5685476008329589, 'CORE_MARGIN': 2, 'CELL_PAD_IN_SITES_GLOBAL_PLACEMENT': 3, 'CELL_PAD_IN_SITES_DETAIL_PLACEMENT': 0, '_FR_LAYER_ADJUST': 0.10066503378702243, 'PLACE_DENSITY_LB_ADDON': 0.46158507922619585, '_SYNTH_FLATTEN': 0, '_PINS_DISTANCE': 3, 'CTS_CLUSTER_SIZE': 31, 'CTS_CLUSTER_DIAMETER': 117, '_FR_GR_OVERFLOW': 1}
(pid=162978) [INFO TUN-0003] Best parameters written to logs/sky130hd/ibex/test-tune-2022-01-04-13-44-21/autotuner-best.json
(lswkim38) lswkim@autotuner-1 util$
```

Demonstration – How to find the logs and results

- Logs

logs path: OpenROAD-flow-

scripts/flow/logs/sky130hd/ibex/test-tune-2022-01-04-12-09-50

Platform	Design	Exp name	Mode	Date
AutoTunerBase_fe2336fe_102_CELL_PAD_IN_SITES_DETAIL_PLACEMENT=2,CELL_PAD_IN_SITES_GLOBAL_PLACEMENT=2,CORE_ASPECT_RATIO=1.3982,CORE_2022-01-04_12-15-23				
AutoTunerBase_fe937b3e_163_CELL_PAD_IN_SITES_DETAIL_PLACEMENT=2,CELL_PAD_IN_SITES_GLOBAL_PLACEMENT=2,CORE_ASPECT_RATIO=0.64249,COR_2022-01-04_12-22-33				
AutoTunerBase_fed9edd0_164_CELL_PAD_IN_SITES_DETAIL_PLACEMENT=2,CELL_PAD_IN_SITES_GLOBAL_PLACEMENT=1,CORE_ASPECT_RATIO=0.49114,COR_2022-01-04_12-22-34				
experiment_state-2022-01-04_12-09-56.json				
searcher-state-2022-01-04_12-09-56.pkl				
search_gen_state-2022-01-04_12-09-56.json				
variant-003555d328664202f711c14f5982bc65				
variant-0180c5ec69abdd0e7dd6a1435594668b				
variant-01e4d6e7456106435e4c45ac17fe7e67				
variant-02b25c4237f54a9866aa3827b28d7271				
variant-03e1f1f2e479ac640c2ba0afc02d6282				
variant-04a5c61eb3263921bbfe330e75456451				

AutoTuner logs for each trial

Current state of tuning for checkpoint and logging

OR logs for each trial
(hash can be found
in **variant_hash.txt**
in AutoTuner logs)

Demonstration – How to find the logs and results

- Results

```
results path: OpenROAD-flow-
scripts/flow/results/sky130hd/ibex/test-tune-2022-01-04-12-
09-50
{results path}/variant-3603bf82cd6b0e40993341690550d8af
```

hash

```
[swkim@autotuner-1 variant-3603bf82cd6b0e40993341690550d8af]$ ls
1_1_yosys.v          2_floorplan.def      4_1_cts.def
1_synth.sdc          2_floorplan.sdc      4_2_cts_fillcell.def
1_synth.v            2_floorplan.v        4_cts.def
2_1_floorplan.def   3_1_place_gp.def    4_cts.sdc
2_2_floorplan_io.def 3_2_place_iop.def  4_cts.v
2_3_floorplan_tdms.def 3_3_place_resized.def output_guide.mod
2_4_floorplan_macro.def 3_4_place_dp.def route.guide
2_5_floorplan_tapcell.def 3_place.def     updated_clks.sdc
2_6_floorplan_pdn.def 3_place.sdc
```

Intermediate/final DEFs, SDCs, GR guides, and Verilog files

Demonstration – How to use GUI (Tensorflow GUI) (1)

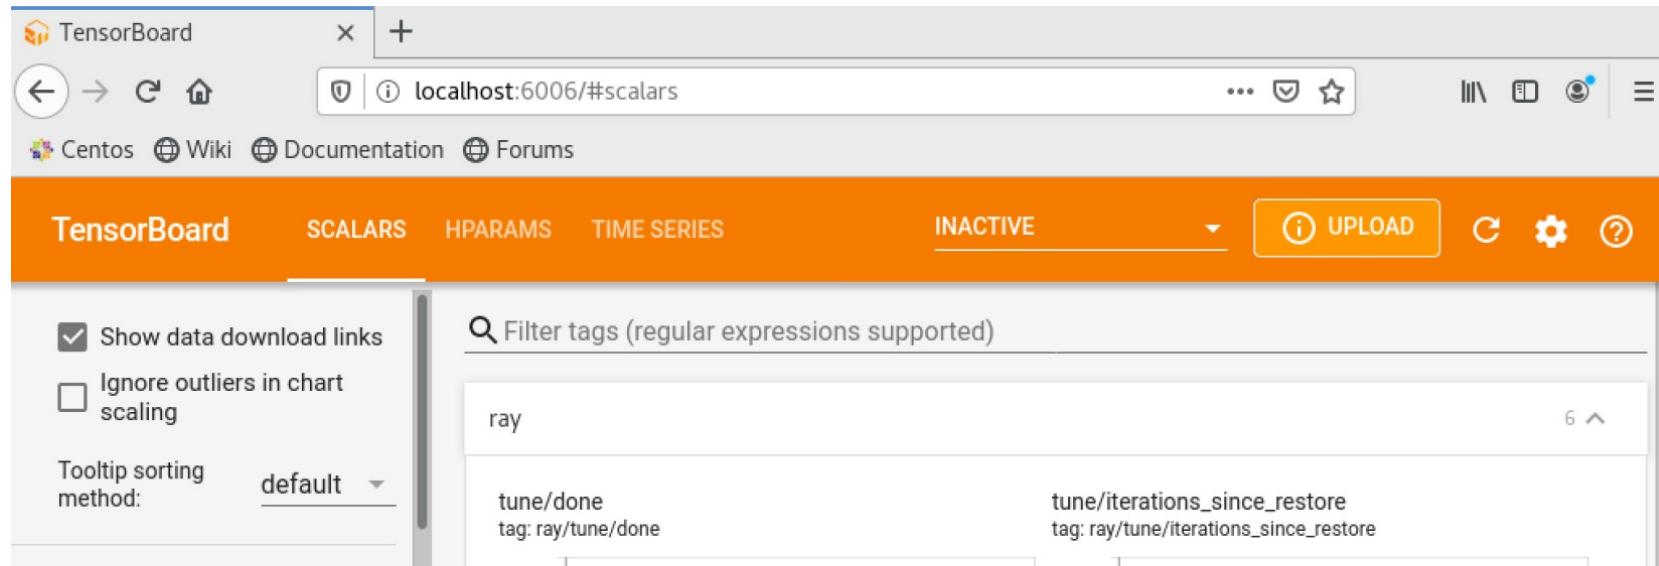
- How to run Tensorflow GUI

- Run Tensorboard

```
shell) cd {log paths}
```

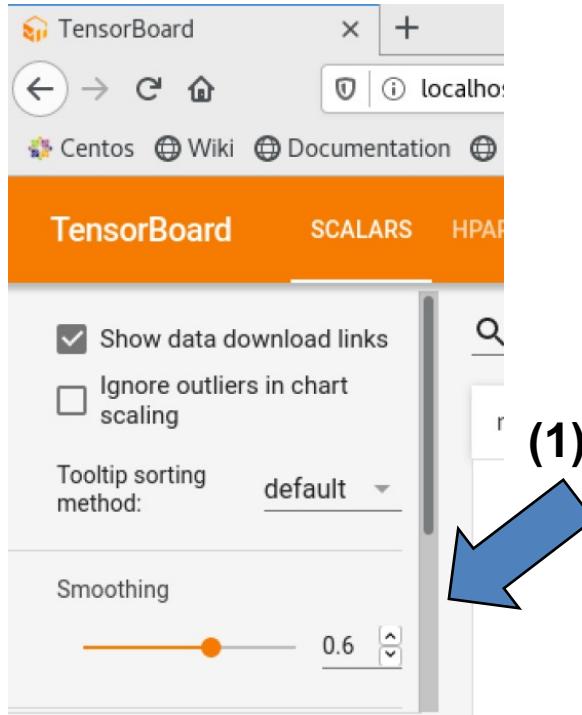
```
shell) tensorboard --logdir=./
```

- Open the local webpage (<http://localhost:6006>)

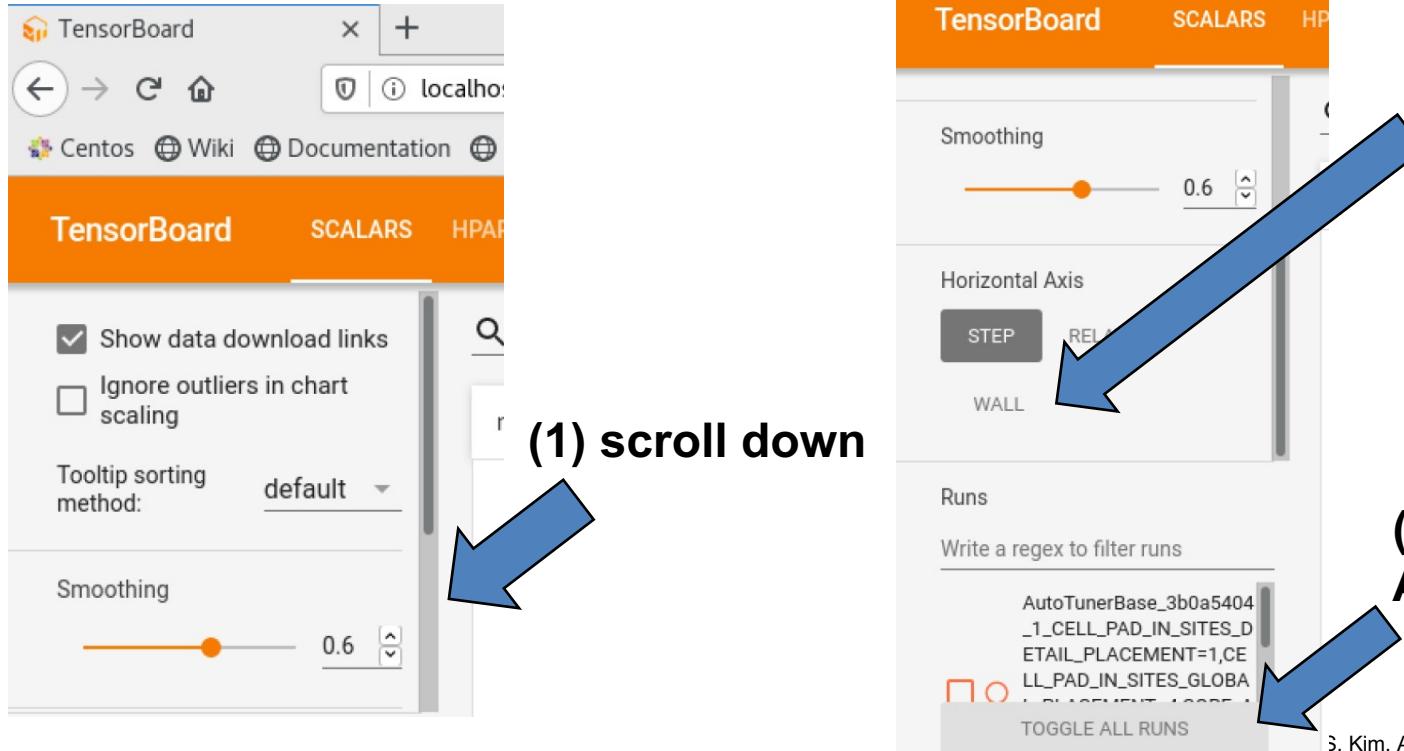


Demonstration – How to use GUI (Tensorflow GUI) (2)

- How to see the results
- Results over walltime



(1) scroll down

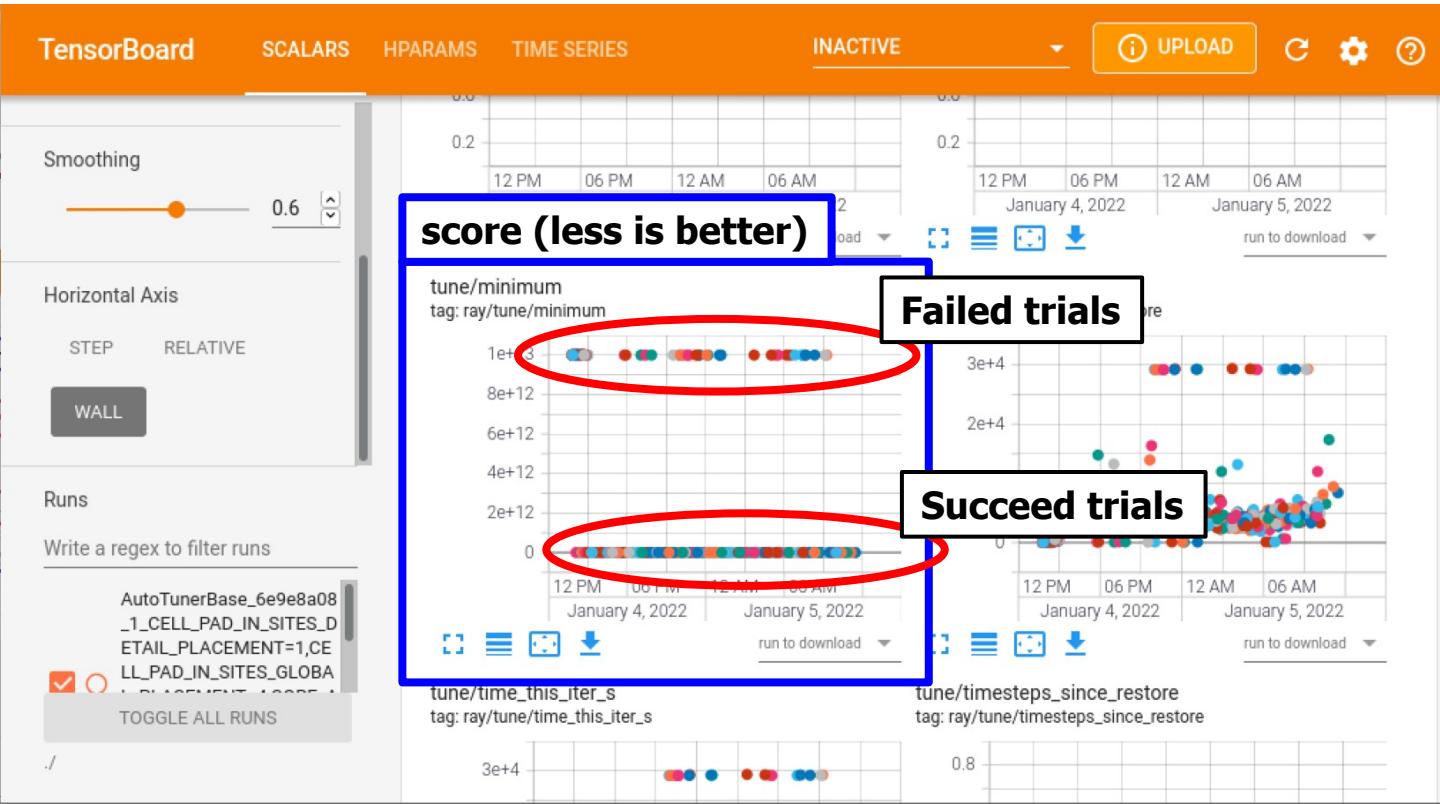


(2) click WALL

(3) click TOGGLE ALL RUNS

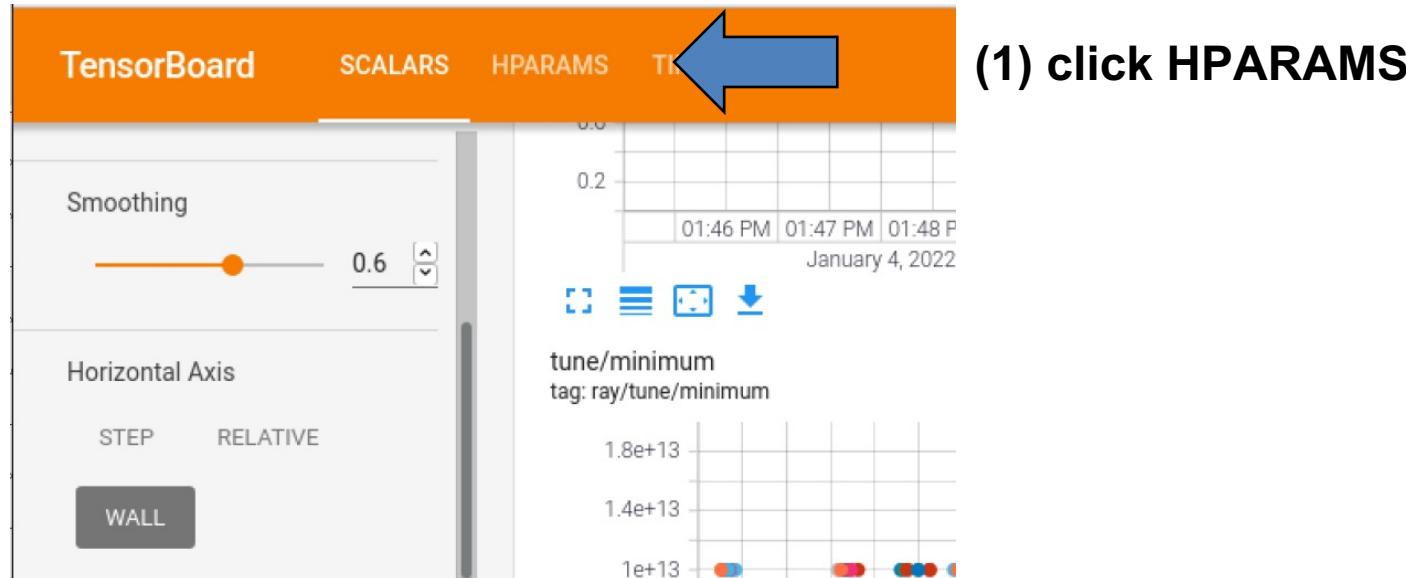
Demonstration – How to use GUI (Tensorflow GUI) (3)

- How to see the results
 - Results over walltime



Demonstration – How to use GUI (Tensorflow GUI) (4)

- How to see the results
- Parallel Coordinates View



Demonstration – How to use GUI (Tensorflow GUI) (5)

- How to see the results

- Parallel Coordinates View

(2) click PARALLEL COORDINATES VIEW

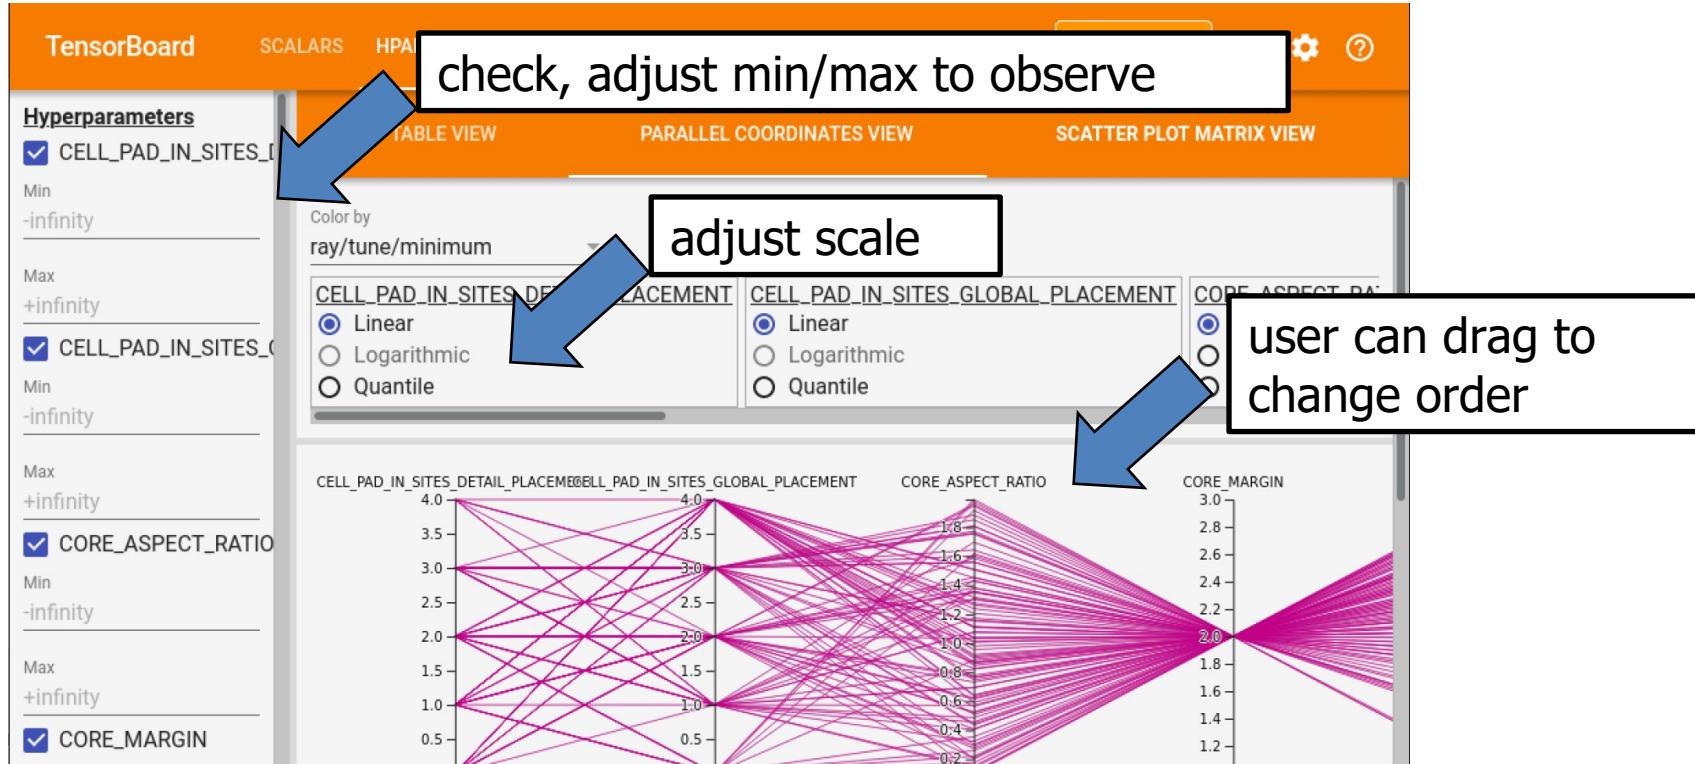
The screenshot shows the TensorBoard interface with the following details:

- Hyperparameters:** A sidebar on the left lists hyperparameters with checkboxes:
 - CELL_PAD_IN_SITES_D (checked)
 - CELL_PAD_IN_SITES_O (unchecked)
 - CORE_ASPECT_RATIO (checked)
 - CORE_MARGIN (checked)
- Filtering:** Below each checked hyperparameter are "Min" and "Max" fields set to "-infinity".
- View Options:** An orange bar at the top right contains three tabs:
 - TABLE VIEW
 - PARALLEL COORDINATES VIEW (highlighted with a blue arrow)
 - SCATTER PLOT MATRIX VIEW
- Data Table:** The main area displays a table of trial data. The columns are:
 - Trial ID
 - Show Metrics
 - CELL_PAD_IN_SITES_DETAIL_PLACEMENT
 - CELL_PAD_IN_SITES_GLOBAL_PLACEMENT
 - CORE_ASPECT_RATIO
 - CORE_JThe table has 10 rows of data.

Trial ID	Show Metrics	CELL_PAD_IN_SITES_DETAIL_PLACEMENT	CELL_PAD_IN_SITES_GLOBAL_PLACEMENT	CORE_ASPECT_RATIO	CORE_J
AutoTunerBase_...	<input type="checkbox"/>	4.0000	3.0000	1.0132	2.0000
AutoTunerBase_...	<input type="checkbox"/>	3.0000	3.0000	1.6440	2.0000
AutoTunerBase_...	<input type="checkbox"/>	4.0000	1.0000	0.69444	2.0000
AutoTunerBase_...	<input type="checkbox"/>	4.0000	2.0000	1.5878	2.0000
AutoTunerBase_...	<input type="checkbox"/>	3.0000	2.0000	1.3567	2.0000
AutoTunerBase_...	<input type="checkbox"/>	2.0000	0.0000	0.88496	2.0000
AutoTunerBase_...	<input type="checkbox"/>	4.0000	4.0000	1.1097	2.0000
AutoTunerBase_...	<input type="checkbox"/>	3.0000	4.0000	0.91041	2.0000
AutoTunerBase_...	<input type="checkbox"/>	4.0000	1.0000	0.14049	2.0000
AutoTunerBase_...	<input type="checkbox"/>	3.0000	1.0000	1.3074	2.0000

Demonstration – How to use GUI (Tensorflow GUI) (6)

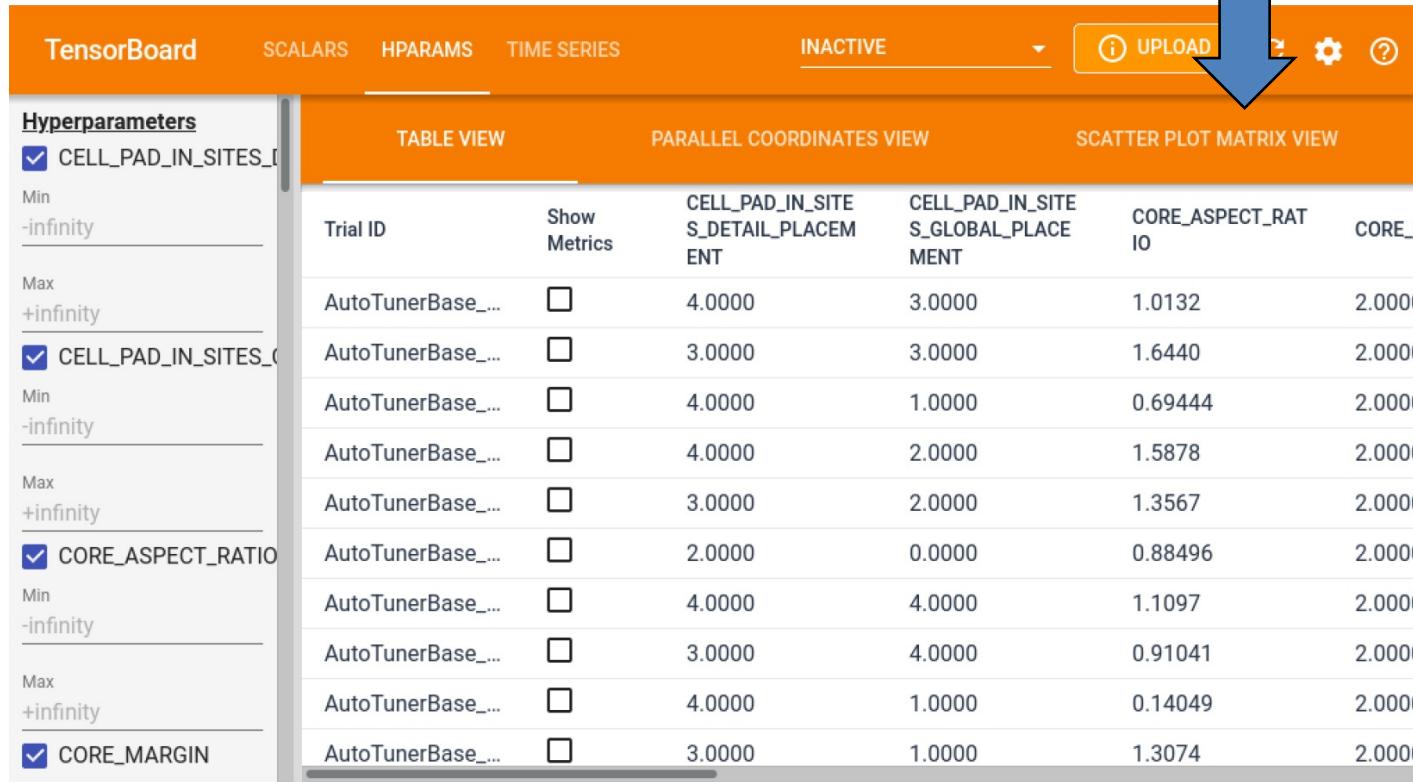
- How to see the results
 - Parallel Coordinates View



Demonstration – How to use GUI (Tensorflow GUI) (7)

- How to see the results
 - Scatter plot matrix

(1) click SCATTER PLOT MATRIX VIEW



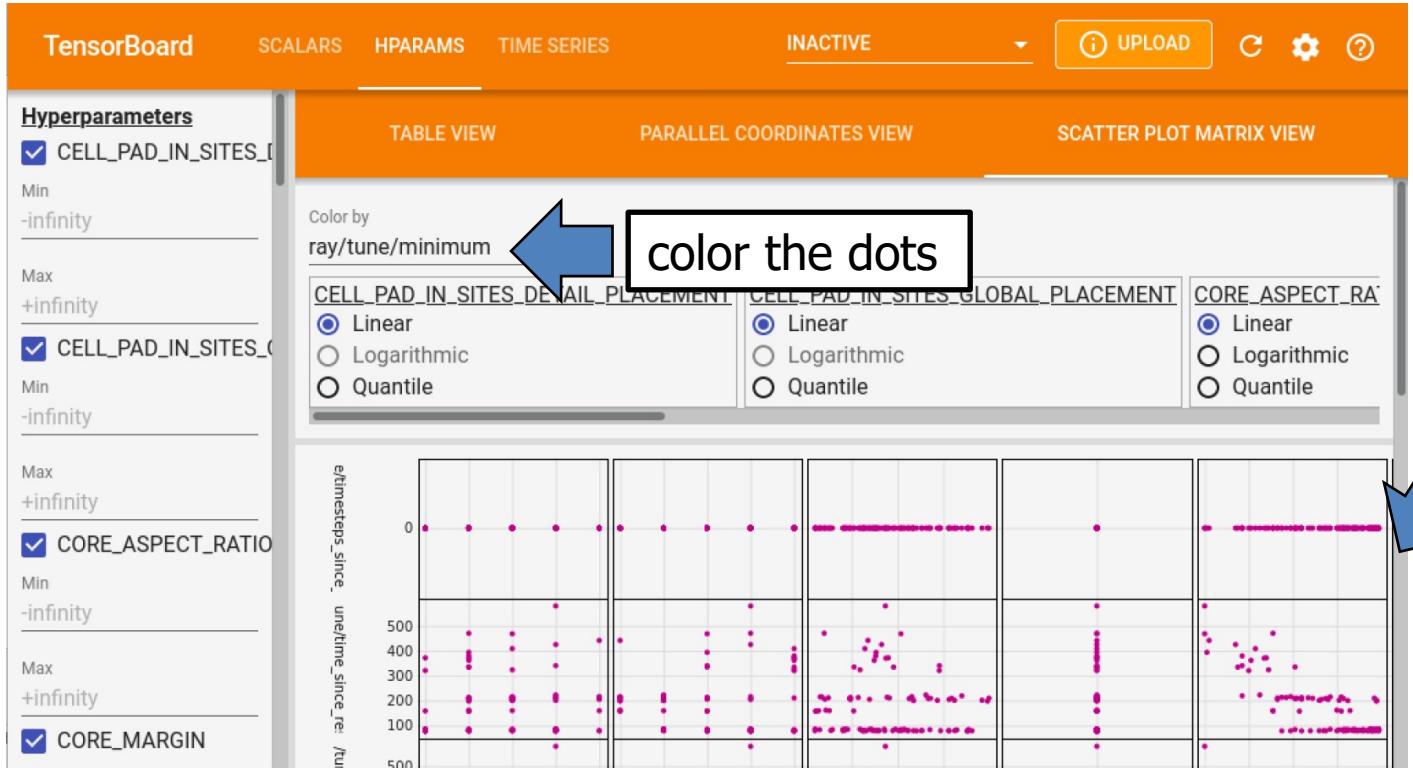
The screenshot shows the TensorBoard interface with the following details:

- Hyperparameters:** A sidebar on the left lists hyperparameters with checkboxes:
 - CELL_PAD_IN_SITES_D (checked)
 - CELL_PAD_IN_SITES_O (unchecked)
 - CORE_ASPECT_RATIO (checked)
 - CORE_MARGIN (checked)
- Metrics:** A table view displays metrics for multiple trials. The columns include Trial ID, Show Metrics, CELL_PAD_IN_SITE_S_DETAIL_PLACEMENT, CELL_PAD_IN_SITE_S_GLOBAL_PLACEMENT, CORE_ASPECT_RATIO, and CORE_J. Each row contains a checkbox next to the Trial ID.
- Top Bar:** Includes tabs for TensorBoard, SCALARS, HPARAMS, and TIME SERIES. The SCALARS tab is selected. It also features an INACTIVE dropdown, an UPLOAD button, and other icons.
- Bottom Buttons:** Three buttons are visible: TABLE VIEW, PARALLEL COORDINATES VIEW, and SCATTER PLOT MATRIX VIEW. The SCATTER PLOT MATRIX VIEW button is highlighted with a large blue arrow pointing to it.

	Trial ID	Show Metrics	CELL_PAD_IN_SITE_S_DETAIL_PLACEMENT	CELL_PAD_IN_SITE_S_GLOBAL_PLACEMENT	CORE_ASPECT_RATIO	CORE_J
1	AutoTunerBase_...	<input type="checkbox"/>	4.0000	3.0000	1.0132	2.0000
2	AutoTunerBase_...	<input type="checkbox"/>	3.0000	3.0000	1.6440	2.0000
3	AutoTunerBase_...	<input type="checkbox"/>	4.0000	1.0000	0.69444	2.0000
4	AutoTunerBase_...	<input type="checkbox"/>	4.0000	2.0000	1.5878	2.0000
5	AutoTunerBase_...	<input type="checkbox"/>	3.0000	2.0000	1.3567	2.0000
6	AutoTunerBase_...	<input type="checkbox"/>	2.0000	0.0000	0.88496	2.0000
7	AutoTunerBase_...	<input type="checkbox"/>	4.0000	4.0000	1.1097	2.0000
8	AutoTunerBase_...	<input type="checkbox"/>	3.0000	4.0000	0.91041	2.0000
9	AutoTunerBase_...	<input type="checkbox"/>	4.0000	1.0000	0.14049	2.0000
10	AutoTunerBase_...	<input type="checkbox"/>	3.0000	1.0000	1.3074	2.0000

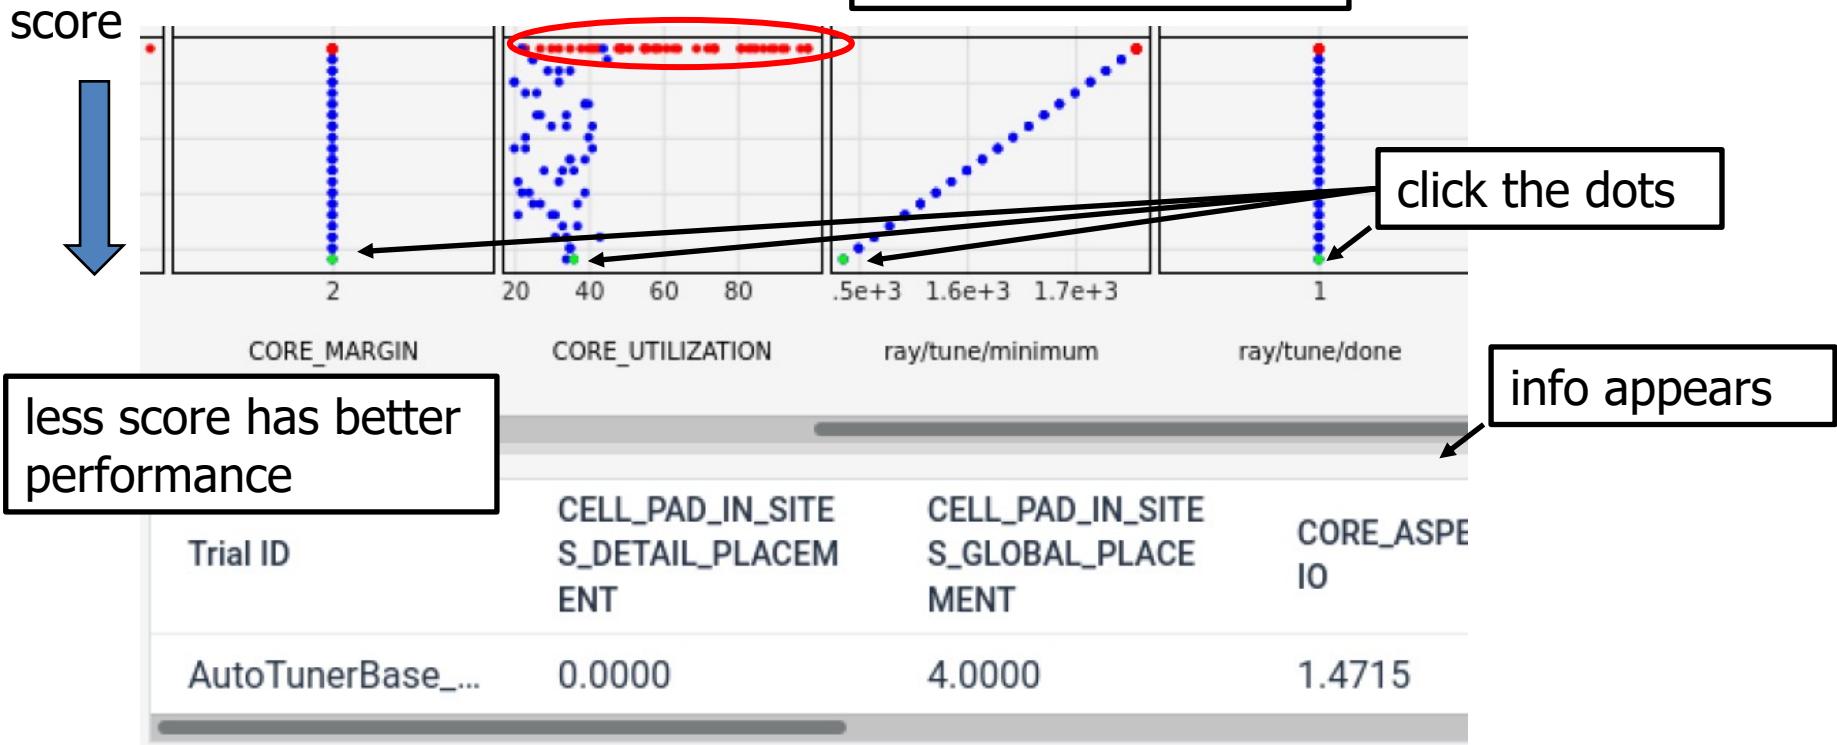
Demonstration – How to use GUI (Tensorflow GUI) (8)

- How to see the results
 - Scatter plot matrix



Demonstration – How to use GUI (Tensorflow GUI) (9)

- How to see the results
 - Scatter plot matrix



Demonstration – How to use GUI (Tensorflow GUI) (10)

- How to see the results
- Download CSV, JSON, LaTex

TensorBoard SCALARS HPARAMS TIME SERIES INACTIVE UPLOAD C G ?

Status
 Unknown
 Success Failure
 Running

Sorting
Sort by Direction

Paging
Number of matching session groups: 500
Page... Max # of session ...
1 / 5 100

Download data as [CSV](#) [JSON](#) [LaTeX](#)

TABLE VIEW PARALLEL COORDINATES VIEW SCATTER PLOT MATRIX VIEW

ray/tune/minimum	ray/tune/done	ray/tune /time_this_iter_s	ray/tune /time_since_restore	ray/tune /timesteps_since_restore
1.0000e+13	1.0000	82.423	82.423	0.0000
1.0000e+13	1.0000	555.50	555.50	0.0000
1.0000e+13	1.0000	87.077	87.077	0.0000
1.0000e+13	1.0000	217.22	217.22	0.0000
1.0000e+13	1.0000	79.131	79.131	0.0000
1.0000e+13	1.0000	2323.5	2323.5	0.0000
1.0000e+13	1.0000	79.782	79.782	0.0000
1.0000e+13	1.0000	228.83	228.83	0.0000
1.0000e+13	1.0000	79.137	79.137	0.0000
1537.1	1.0000	2986.1	2986.1	0.0000

(1) scroll down and click to download



Experiments and Results

- Experimental Setup
 - Framework is written in Python 3.8
 - Machine: Google Cloud Platform m1-ultramem-160. 160 2.2GHz vCPUs, 3.8TB memory
 - Platform: **ASAP 7nm, Skywater 130nm HD**
 - Design: **aes, ibex, jpeg**
 - Coefficient Setting (C_{power} , C_{perform} , C_{area})
 - $\text{power}_{\text{opt}}$: (10000, 100, 100)
 - $\text{perform}_{\text{opt}}$: (100, 10000, 100)
 - area_{opt} : (100, 100, 10000)
 - $\text{uniform}_{\text{opt}}$: (100, 100, 100)
 - All studies are parallelized to 40 concurrent OpenROAD jobs, with each job being assigned to 4 cores.

Experiments and Results

- Experimental Setup
 - Tunable tool and design parameters

Parameters	Description	Type	Range
<code>_SDC_CLK_PERIOD</code>	Target clock period (ns)	float	[T_{\min} , T_{\max}]
<code>CORE_UTILIZATION</code>	Target core utilization (%)	int	[20, 99]
<code>CORE_ASPECT_RATIO</code>	Floorplan aspect ratio	float	[0.1, 2.0]
<code>CELL_PAD_IN_SITES_GLOBAL_PLACEMENT</code>	Cell padding for global placement (site)	int	[0, 4]
<code>CELL_PAD_IN_SITES_DETAIL_PLACEMENT</code>	Cell padding for detailed placement (site)	int	[0, 4]
<code>_FR_LAYER_ADJUST</code>	Layer resource adjustment for global routing (%)	float	[0.1, 0.7]
<code>PLACE_DENSITY_LB_ADDON</code>	Additional lower bound increase of the target local global placement density (%)	float	[0.00, 0.99]
<code>_SYNTH_FLATTEN</code>	Design hierarchy flattening	int	[0, 1]
<code>_PINS_DISTANCE</code>	Minimum IO pad distance (#tracks)	int	[1, 3]
<code>CTS_CLUSTER_SIZE</code>	Target CTS sink cluster size	int	[10, 40]
<code>CTS_CLUSTER_DIAMETER</code>	Target CTS sink cluster diameter (um)	int	[80, 12]

Experiments and Results – Search Algorithm Comparison

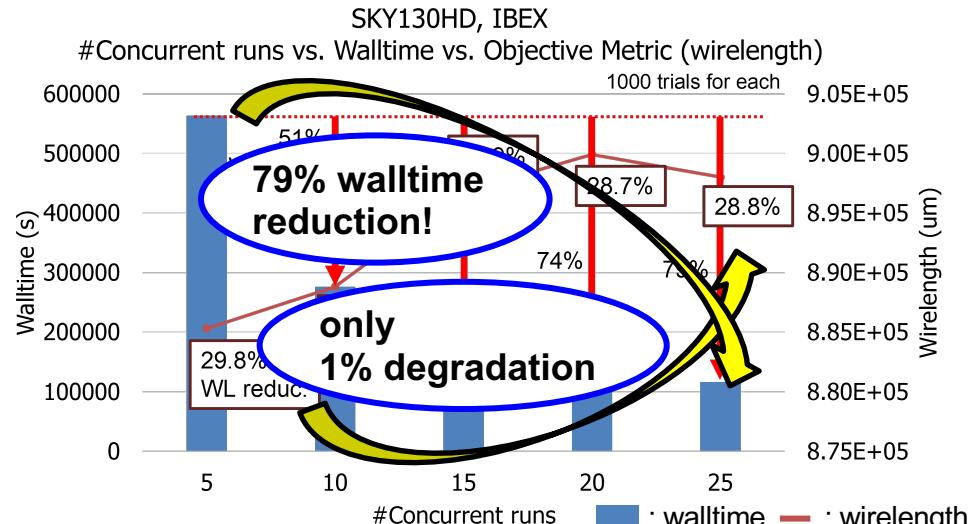
- Coefficient setting: **perform_{opt}**
- Walltime constraint = 63966s (~17.8 hrs)
- Tuning scheduler: asynchronous scheduler

Algorithm	#trials	effCP (ns)	Improv.	Util (%)	Power	Improv.	Score
Reference		17.775		20	0.0523		1020000
PBT	1000	15.1621	15%	37	0.0143	73%	863611
HyperOpt	742	14.9405	16%	40	0.0286	45%	853503
Ax	478	16.0395	10%	26	0.0172	67%	914902
Optuna	1493	14.8693	16%	42	0.0278	47%	849094
Nevergrad	2925	15.2311	14%	18	0.0194	63%	870843

- **Nevergrad** have numerous trials by actively exploring loose hyperparameters or infeasible sets.
- **HyperOpt** and **Optuna** shows better (i.e., smaller) final scores
- However, appropriate search algorithm may vary depending on the design difficulty, experimental setup, CAD tools and parallelization!

Experiments and Results – Exploration vs. Exploitation

- High parallelization may cause performance degradation!
 - For example, if we use N concurrent jobs for N trials, walltime is minimized but this is the same as random (“multi-start”) search (no iteration, no exploitation)
- Objective function: wirelength minimization
- #trials = 1000
- By parallelization,
 - Reduces the total walltime (left y-axis) up to 79%
 - Worse performance (wirelength) (right y-axis), but only by 1% degradation
- Why? → With asynchronous scheduler:
 - Walltime is proportional to $\frac{1}{\#concurrent_jobs'}$
 - Performance degradation is almost proportional to $\frac{\#concurrent_jobs}{\#trials}$



If the number of trials is large enough, active parallelization significantly reduces the total latency of autotuning

Experiments and Results – Automatic PPA Push

- 1000 trials for each setting. The table shows two of the three designs tested.
- For {SKY130, ASAP7} respectively, we achieve on average:
- {41%, 32%} power consumption reduction with power setting
- {21%, 18%} performance improvement with performance setting
- {68%, 58%} area reduction with area setting
- Comparison to OpenROAD-flow-scripts design-platform default in master branch
- Walltime: all < 2 days

Platform	Design	Setting	effCP (ns)	Improv.	Util (%)	Improv.	Power (W)	Improv.
SKY130	aes	ref	4.554	0%	19	0%	0.0838	0%
		power	4.373	4%	23	17%	0.0569	32%
		perf.	3.491	23%	22	14%	0.0942	-12%
		area	3.763	17%	42	55%	0.0558	33%
	jpeg	uniform	3.38	26%	35	46%	0.0612	27%
		ref	9.277	0%	20	0%	0.0005	0%
		power	7.781	16%	37	46%	0.0711	21%
		perf.	7.259	22%	31	35%	0.135	-49%
		area	7.781	-1%	80	78%	0.101	-12%
		uniform	7.88	15%	64	76%	0.08	12%
Platform	Design	Setting	effCP (ps)	Improv.	Util (%)	Improv.	Power (W)	Improv.
ASAP7	aes	ref	497.630	0%	29	0%	0.0170	0%
		power	445.810	10%	64	55%	0.0098	42%
		perf.	408.012	18%	52	44%	0.0325	-91%
		area	462.413	7%	70	59%	0.0154	9%
	jpeg	uniform	443.146	11%	60	52%	0.0102	40%
		ref	1039.200	0%	30	0%	0.0348	0%
		power	1021.997	2%	58	48%	0.0340	2%
		perf.	838.097	19%	59	49%	0.0520	-49%
		area	1066.411	-3%	68	56%	0.0366	-5%
		uniform	945.232	9%	64	53%	0.0379	-9%

Summary

- **METRICS2.1 infrastructure (Tutorial Part 5)**
 - An initiative of the IEEE CEDA Design Automation Technical Committee (DATC)
 - Enables standardized metrics collection and design process recording
 - We provide archives, reproducible config files and ML applications in the form of Jupyter notebooks
- **Flow Tuning (AutoTuner) (Tutorial Part 6)**
 - Open-source design flow autotuning framework
 - Supports synchronous/asynchronous parallelization, various search algorithms, and metrics collection using METRICS2.1
- **We invite you for collaboration!**
 - METRICS2.1
 - <https://github.com/ieee-ceda-datc/datc-rdf-Metrics4ML>
 - Flow Tuning (AutoTuner)
 - <https://github.com/ieee-ceda-datc/datc-rdf-flow-tuner>

THANK YOU !

We thank the OpenROAD team and the IEEE CEDA Design Automation Technical Committee for many helpful interactions.

This work is supported in part by the U.S. National Science Foundation (CCF-2112665) and the U.S. Defense Advanced Research Projects Agency (HR0011-18-2-0032).

<https://theopenroadproject.org> <https://tilos.ai> <https://vlsicad.ucsd.edu>

Part 7

Summary and Future Work

Jinwook Jung and Andrew B. Kahng

Summary and Future Work

- Coming full circle: EDA = Optimization
- Futures in Learning + Optimization

EDA is Optimization

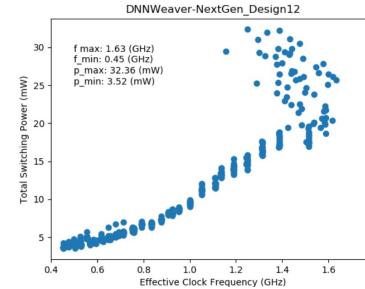
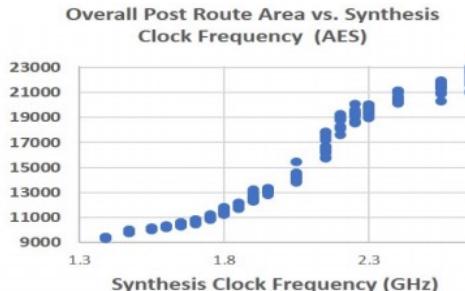
- **EDA is about optimizations and algorithms**
 - High stakes: performance, power, design closure
- **Discrete, combinatorial formulations at huge scale**
 - **Optimizations:** ILP, MCF, QAP, SAT/SMT, LR, ...
 - **Algorithms:** min-cost flow, high-dim DP, ...
- **But need an answer overnight**
- Reality under the hood: **metaheuristics**
 - Annealing, multi-start, PSO, ripup-reroute, greed, ...
 - Convenient objectives, customer-/tech-specific tuning, ...
 - *... which comes at a cost*

Never Enough Time: Heuristics and Chaos

- “CAD tools are chaos machines”

-- Ward Vercruyse, Sun UltraSPARC III CAD manager, **Physical Design Workshop 1996**

- Push harder on a tool that is made up of heuristics stacked on top of heuristics → result becomes less predictable
- Change initial conditions slightly → outcome can change a lot



- Recurring theme in my group...
 - ISQED02 (**Noise**), ISQED10 (**Chaos**)
 - DAC18 WIP (**Multi-Armed Bandits**)

Measurement of Inherent Noise in EDA Tools *

Andrew B. Kahng [†] and Stefanus Mantik

[†] UCSD CSE and ECE Departments, La Jolla, CA 92093-0114
UCLA Computer Science Department, Los Angeles, CA 90095-1596
abk@ucsd.edu, stefanus@cs.ucla.edu

Methodology From Chaos in IC Implementation

Kwangok Jeong¹ and Andrew B. Kahng^{1,2}
¹ECE and ²CSE Departments, University of California at San Diego, La Jolla, CA, USA
kjeong@vlsicad.ucsd.edu, abk@cs.uscd.edu

A No-Human-in-the-Loop Methodology Toward
Optimal Utilization of EDA Tools and Flows

Andrew B. Kahng, Shriram Kumar and Tushar Shah, UCSD



IC Design = “Multi-Armed Bandit” Problem

- Multi Armed Bandit Problem (MAB): Given a slot machine with N arms, maximize total reward obtained using T “pulls” (iterations)
 - Involves an explore-exploit tradeoff
 - Draw samples to learn model parameters (e.g., distributions of outcomes)
 - Simultaneously maximize reward

IC Design Problem

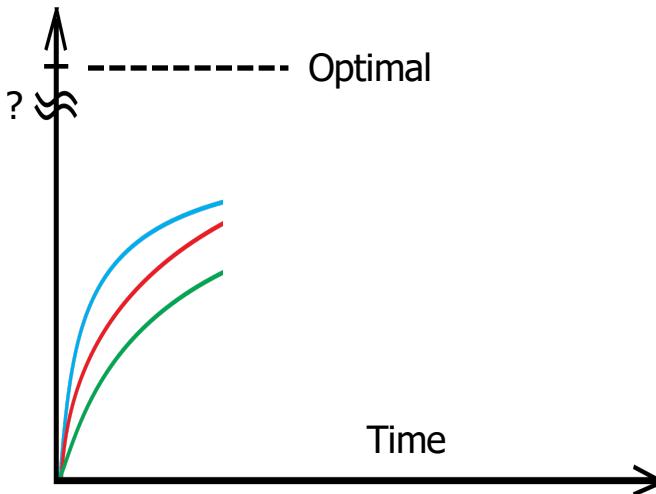
- Each “arm” = a target frequency.
- Each “pull” = a run of the tool flow

Three well-known sampling strategies

- Thompson Sampling
- ϵ -Greedy Sampling
- Softmax Sampling
- (+ Naive: uniformly sample from all available arms)

Have We Lost Sight of the Suboptimality Gap?

- Reality of optimization
 - Better, faster, cheaper – pick any two
- IC EDA: want all three at once
 - “Unfortunately, the runtime of ...”
- But the world has changed
 - Automation, cloud, ...



Question: If you give your SP&R flow 3 extra days of runtime, would it know how to use this extra time?

Question: If you could run 10,000 copies of your P&R tool at the same time, what QOR improvement **should** you expect?

Generic Need: Re-Focus on Suboptimality (to get closer to Optimality!)

Suboptimality ... in what sense?

- Need proper aiming points (which AI can help us **learn**)

Benchmarking

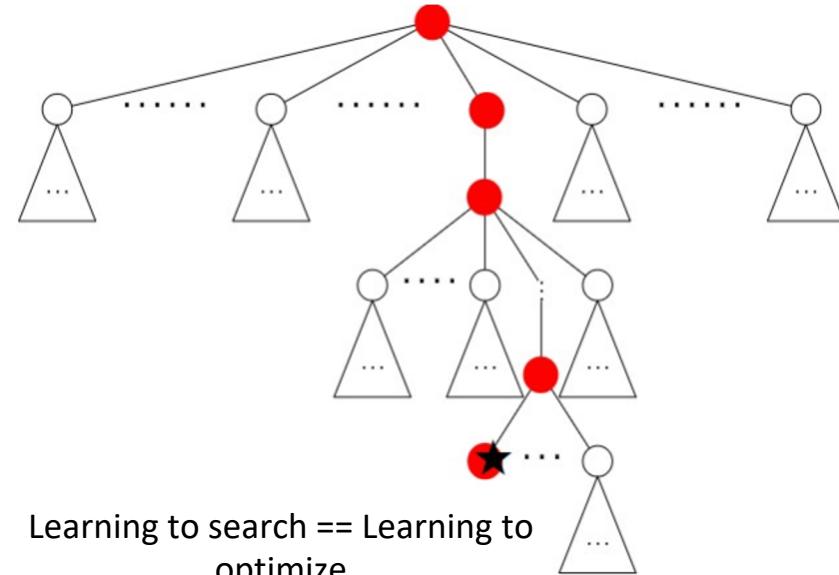
- “Real” benchmarks in EDA have been obfuscated, incomplete, non-vertical, old...
- “Artificial” benchmarks have tiptoed between realism, known optimal solution quality, scalability ...
- Enough (30+ years of this) is enough → find a next level
 - E.g., “Underwriters Laboratories for IC Design Tools”

Comprehending modern compute resources cloud, GPU, accelerators, ...

- EDA Optimization + Learning **naturally** live in the cloud
- **Most of today’s EDA optimization algorithms are “EDA1.0” from the 1980s**
 - → Develop new, cloud-scalable “EDA2.0”

Optimization Challenges

- **Why is optimization difficult?**
 - Combinatorial and high-dimensional
 - Many optimizers/solvers and tuning knobs
 - Greedy, tree-search, branch and bound
 - Gradient descent
- **EDA optimization adds further constraints**
 - Technology, libraries, PVT, RC
 - Tool (per stage)
 - Layout, timing constraints
 - Design rules
 - Tool parameter settings
 - Many more



How to automate efficient search?
Can we formalize learning problem?

Learning to Optimize

Minimize $f(x_1, x_2, x_3, \dots)$

Subject to $g(x_1, x_2, x_3, \dots)$



State: x_1, x_2, \dots, x_{i-1}

Action: x_i

Reward: $f(x_1, x_2, x_3, \dots), g(x_1, x_2, x_3, \dots)$
or $f'(x_1, x_2, \dots, x_i), g'(x_1, x_2, \dots, x_i)$

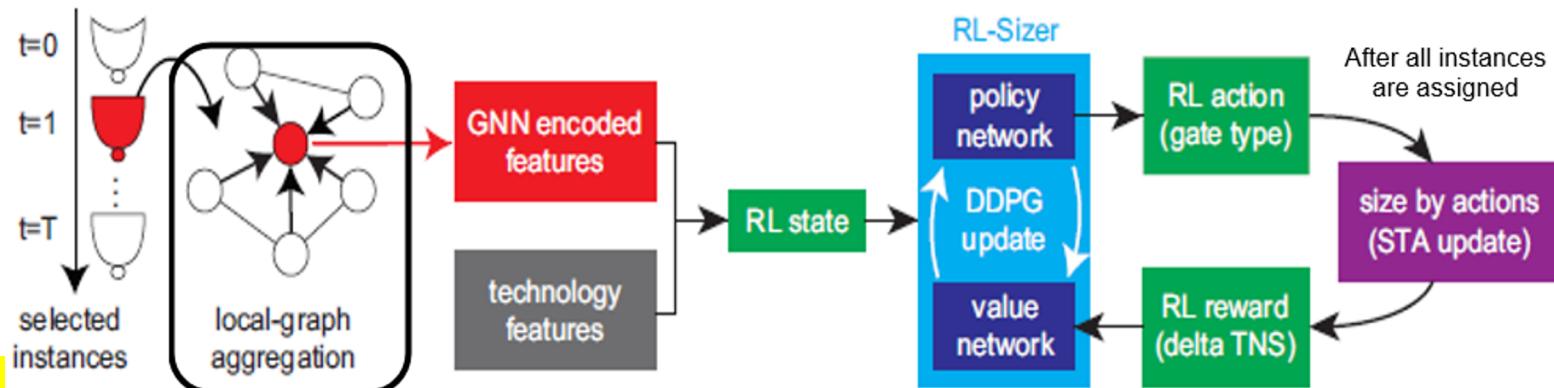
- Key questions

- What are we learning?
- Which algorithm / model are we learning?
- How are we learning?

- Reinforcement learning is promising for sequential decision-making

- Sequential approach to select the best action with feedback (reward)

- Recent RL-based gate sizing (DAC21) achieves up to 30% TNS improvement w.r.t. commercial tool's sizer



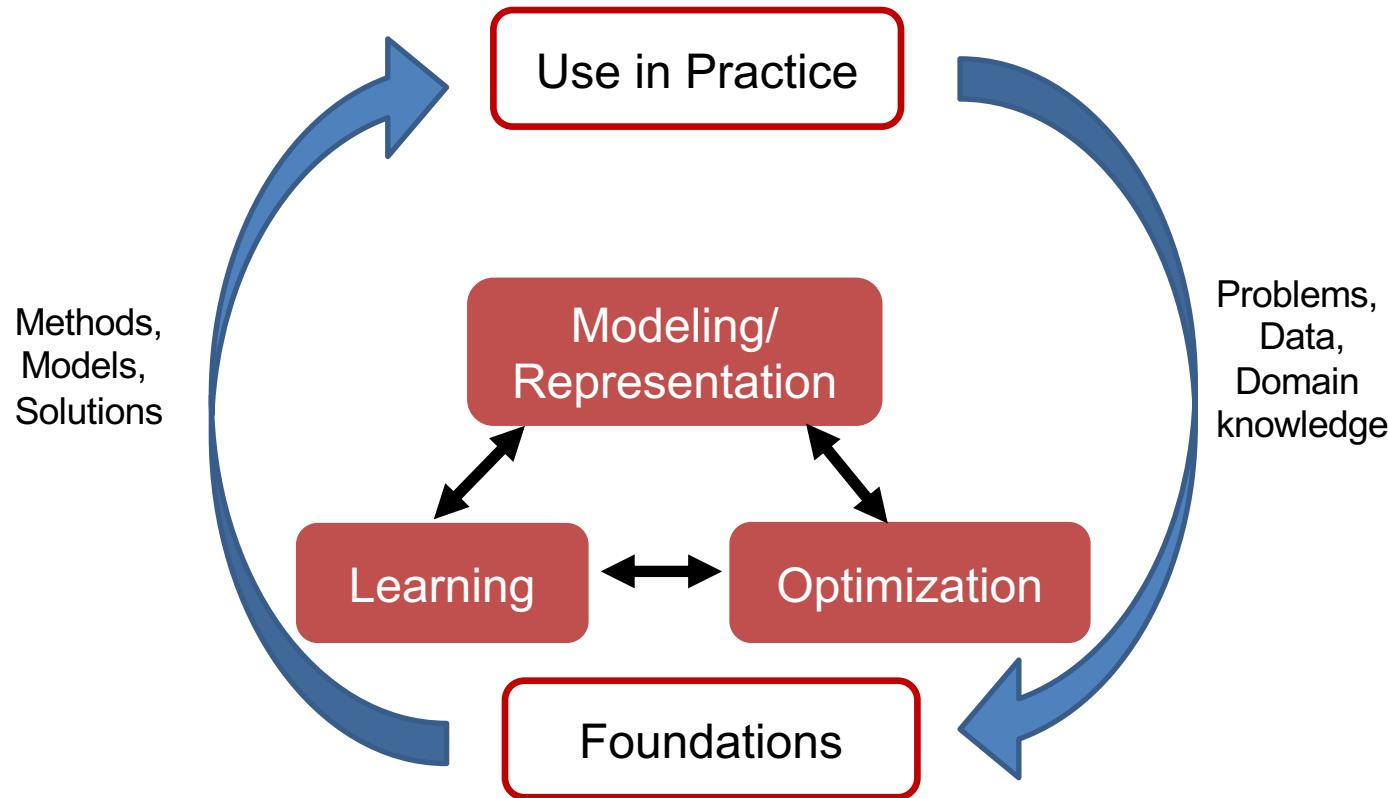
Looking Ahead

- **What are we learning? → Addressed today by majority of EDA works**
 - Base model vs. derived models (for task-specific learning such as congestion, global routing, etc.)
 - Extend to few-shot and transfer learning → largely lacking today
- **Which algorithm / model are we learning? → Largely restricted to hyperparameter tuning**
 - Ensemble, imitation learning that **generalizes** to multiple tasks for production use at industrial scale
- **How are we learning? → Not addressed today in the VLSI/EDA space**
 - Learn the learning rule / objective function best suited for a task
 - Opportunity for algorithm like RL to learn the objective function
 - General vs. specific model

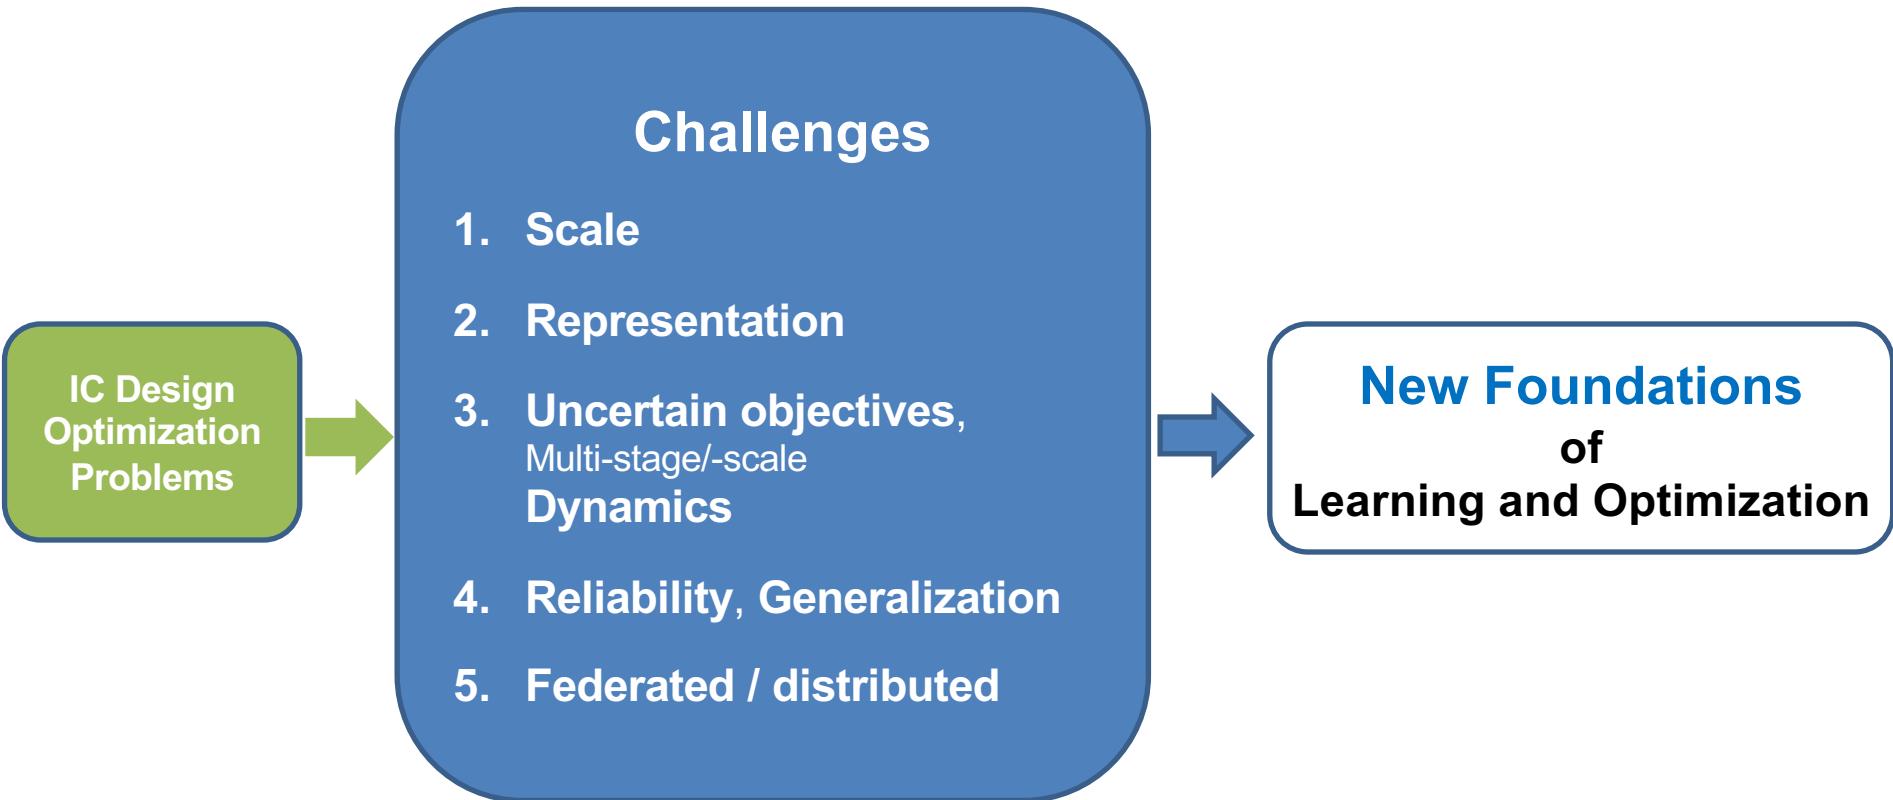
Learning + Optimization

- **How Learning helps Optimization**
 - “Modeling and prediction”: prune doomed runs early
 - “New cost-accuracy tradeoffs in analysis”: less guardbanding
 - Value: more optimization within the available box of resources
 - Center of gravity for most of “ML in/around EDA” so far
- **How Optimization helps Learning**
 - Stochastic gradient descent, nonconvex optimization
 - Distributed/parallel
 - Meta-level: optimization of HW on which learning takes place
- **Virtuous cycle that brings Scaling (of CAD/EDA and IC design)**

Nexus of Theory and Practice



(This Will Keep Us Busy for ... 5-10 Years?)



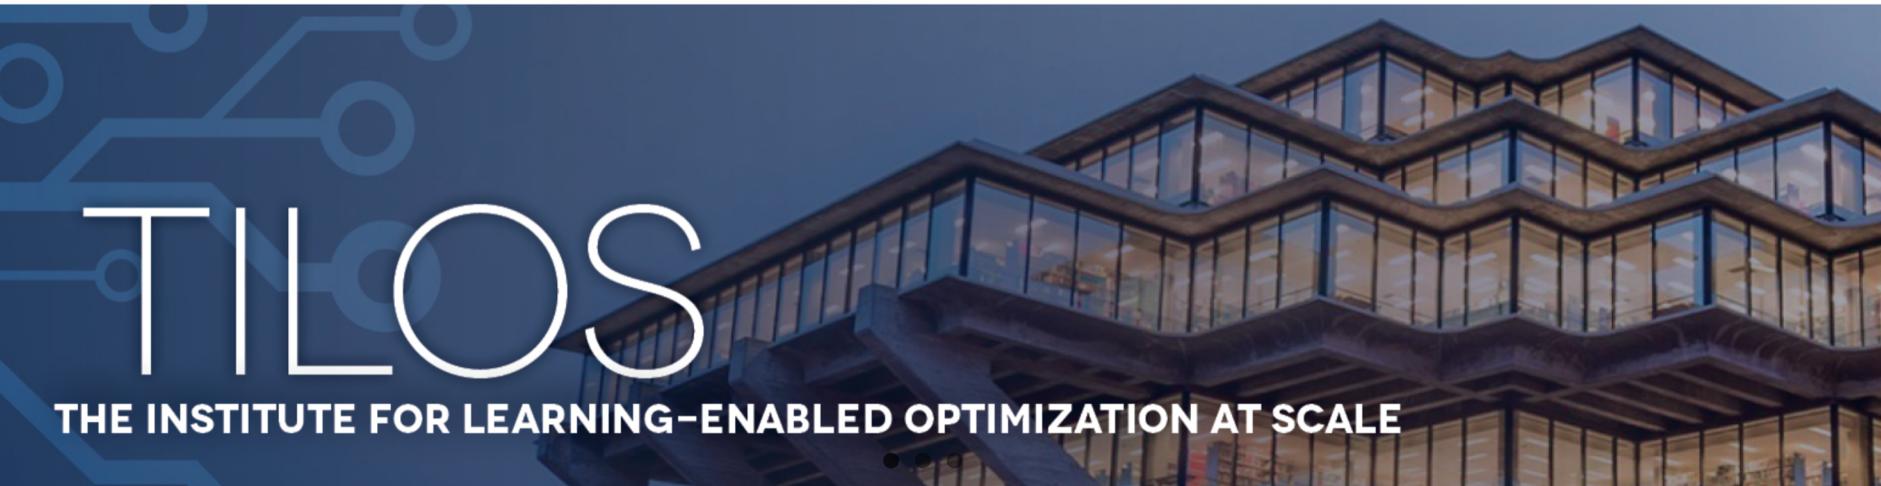
AI and Optimization: Key Directions to Watch

AI advances pose new challenges, provide new tools for optimization

- Bridging Discrete and Continuous
- Distributed, Parallel, and Federated
- Optimization on Manifolds
- Dynamic Decisions under Uncertainty
- Nonconvex Optimization in Deep Learning

New perspectives on classic problems → watch this space!

TILOS: An NSF AI Research Institute for Advances in Optimization

[Team ▾](#)[Advisory](#)[Research ▾](#)[Outreach](#)[News](#)[Events ▾](#)[Publications](#)[Industry ▾](#)[Contact](#)

TILOS

THE INSTITUTE FOR LEARNING-ENABLED OPTIMIZATION AT SCALE

- Partially supported by Intel Corporation
- Launched on November 1, 2021

tilos.ai

THANK YOU !

- Research at UCSD is supported by NSF, DARPA, Qualcomm, Samsung, NXP, Mentor Graphics and the C-DEN center.
- Questions/Feedback: abk@eng.ucsd.edu

Summary of Tutorial 1

- We have presented IEEE CEDA DATC's efforts toward standard platform for ML-enabled EDA and IC design

• Introduction	-----	Part 1
• Challenges in ML-enabled EDA and IC design	-----	Part 2
• DATC's Efforts on standard platform for ML-enabled EDA and IC design	-----	
• Platform for EDA and IC Design	-----	Part 3
• DATC Robust Design Flow	-----	
• Platform for “ML-Enabled” EDA and IC Design	-----	Part 4
• Large-scale data generation in cloud-native way	-----	
• METRICS 2.1: Toward standard metrics collection system	-----	Part 5
• AutoTuner: Autonomous flow tuning framework	-----	Part 6
• Summary and future work	-----	Part 7